

INTRODUCTION TO SPARK WITH SCALA

Introduction to Spark



Agenda

2

- Introduction to Spark
- Core Concepts and Architecture

Introduction to Apache Spark

3

“Big data” is moving beyond one-pass batch jobs to low latency applications that need data sharing

- Matei Zaharia

Introduction to Apache Spark

4

The leading candidate for “successor to MapReduce” is Apache Spark

- Mike Olson

Cloudera Chief Strategy Officer

12/30/2013

Introduction to Apache Spark

5

In-memory compute and machine-learning logic are the key to unlocking the value of Big Data

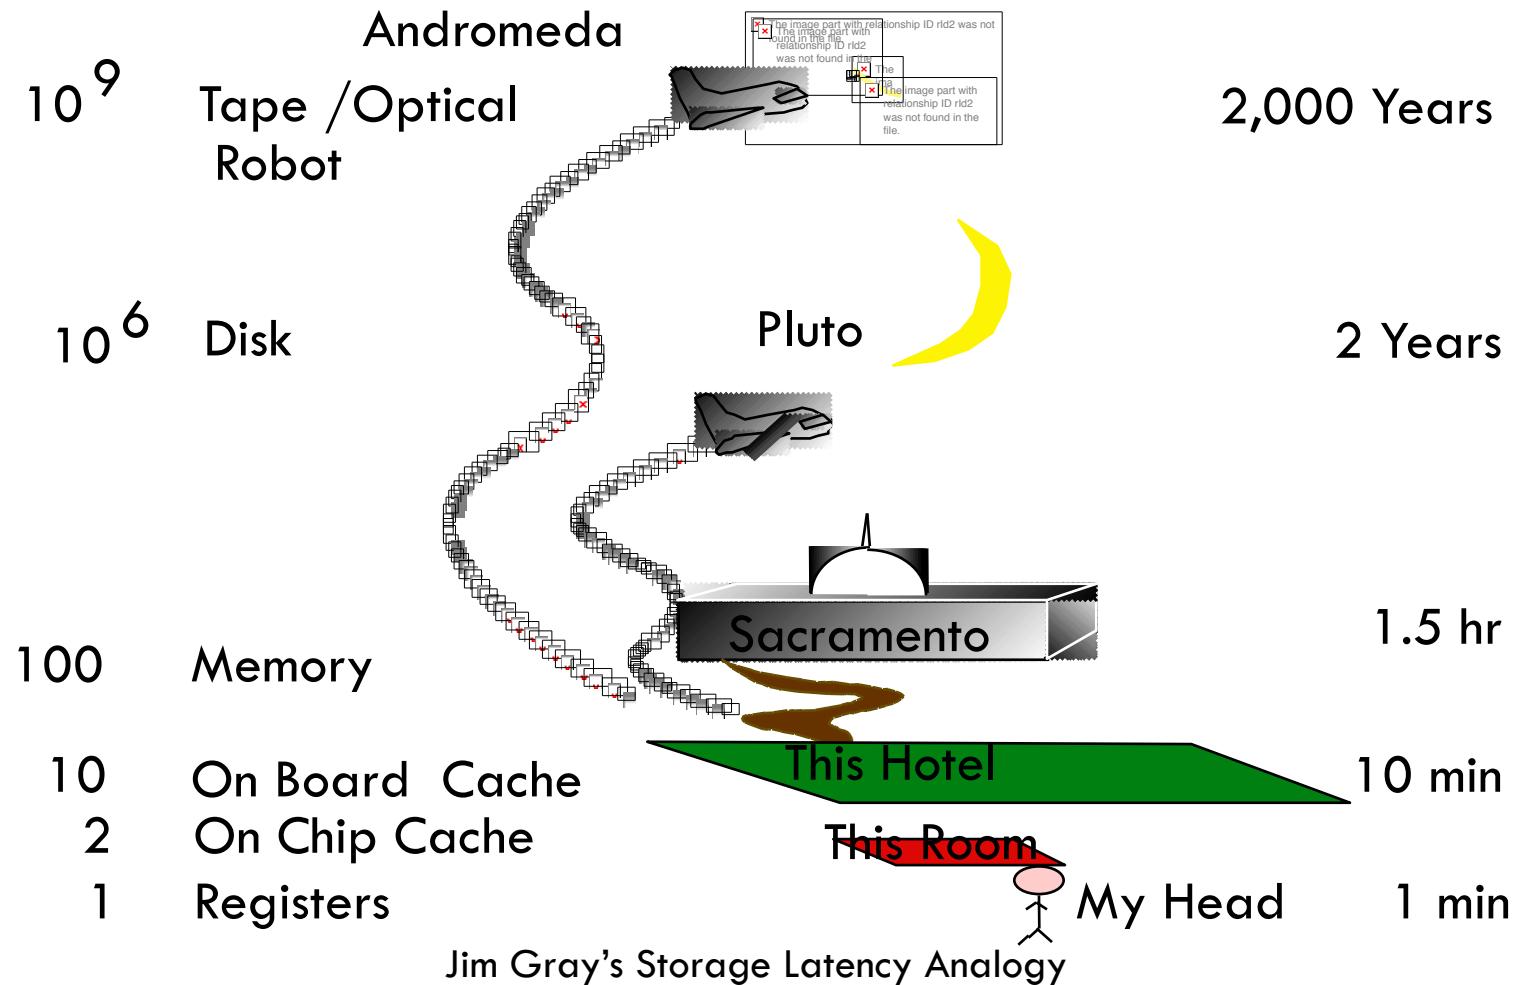
- Christopher Nguyen

CEO Adatao

Introduction to Apache Spark

6

How Far Away is the Data?

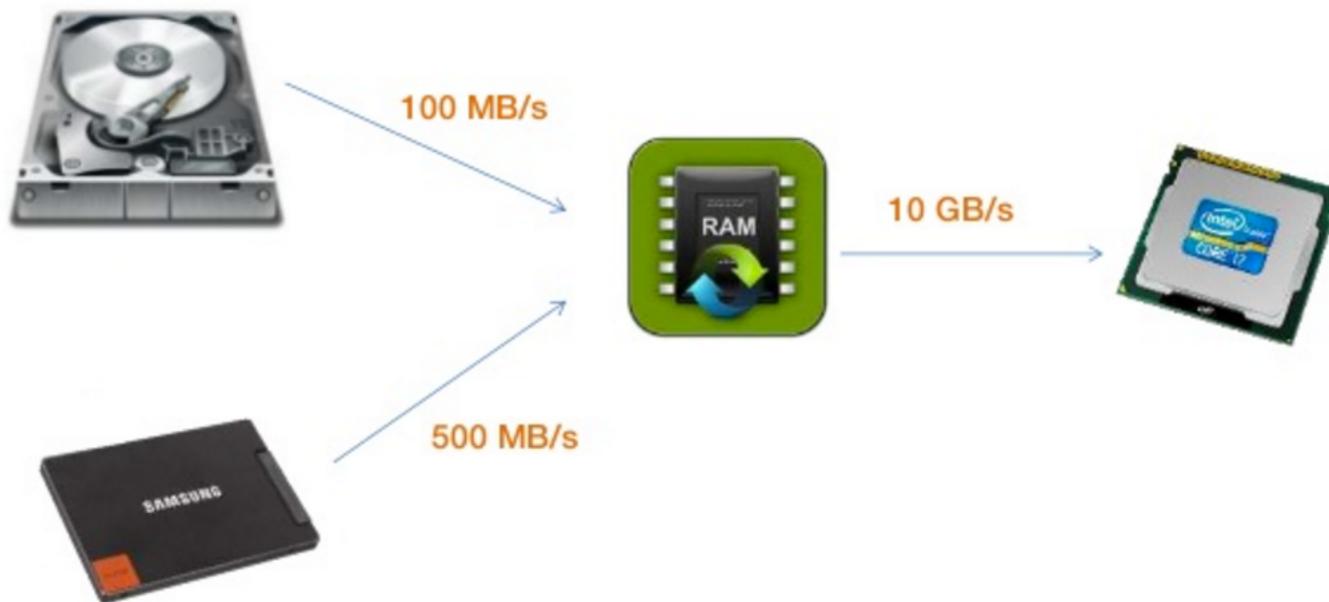


Introduction to Apache Spark

7

Why In-memory Computation Matters

 Clip slide

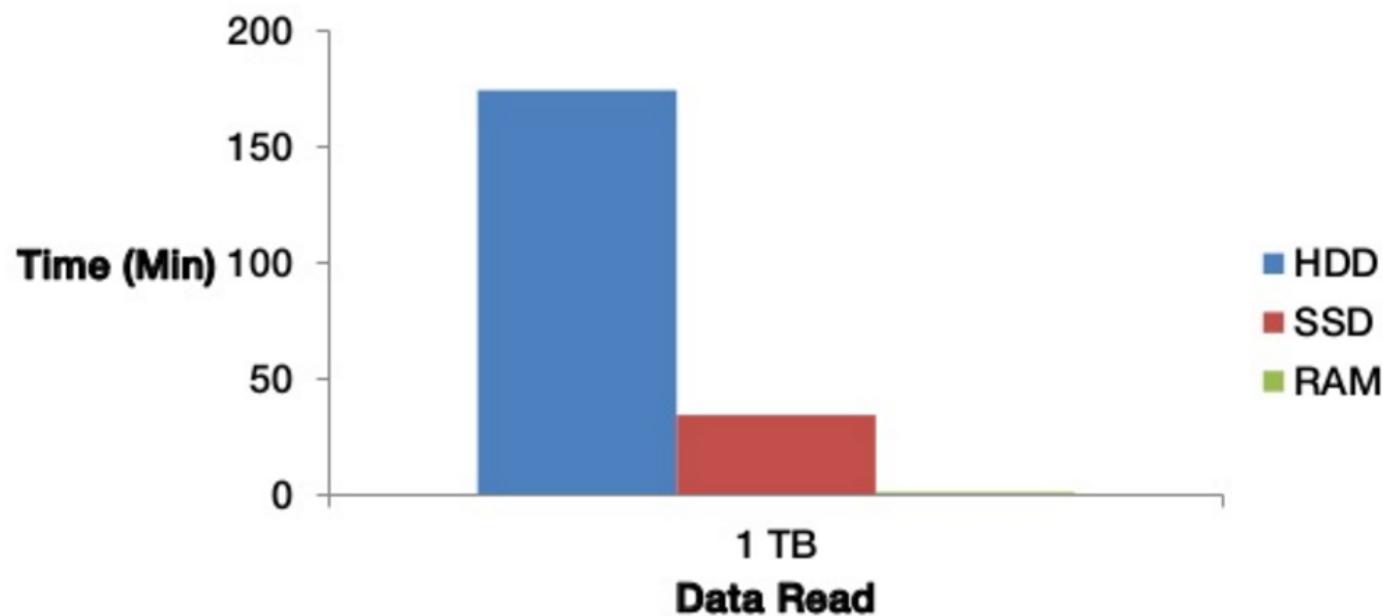


Introduction to Apache Spark

8

Read Time Comparison

 Clip slide



Introduction to Apache Spark

9

Batch Processing

Interactive, iterative,
streaming, graph

Unified Engine

Dremel
Impala
Drill Girah
 Impala
 Presto Storm
 Samza Tez



2004-2013

2007-2015



2014-?

Introduction to Apache Spark

In recent days..... 2010



Spark: Cluster Computing with Working Sets

people.csail.mit.edu/matei/papers/2010/hotcloud_spark.pdf

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing

usenix.org/system/files/conference/nsdi12/nsdi12-final138.pdf

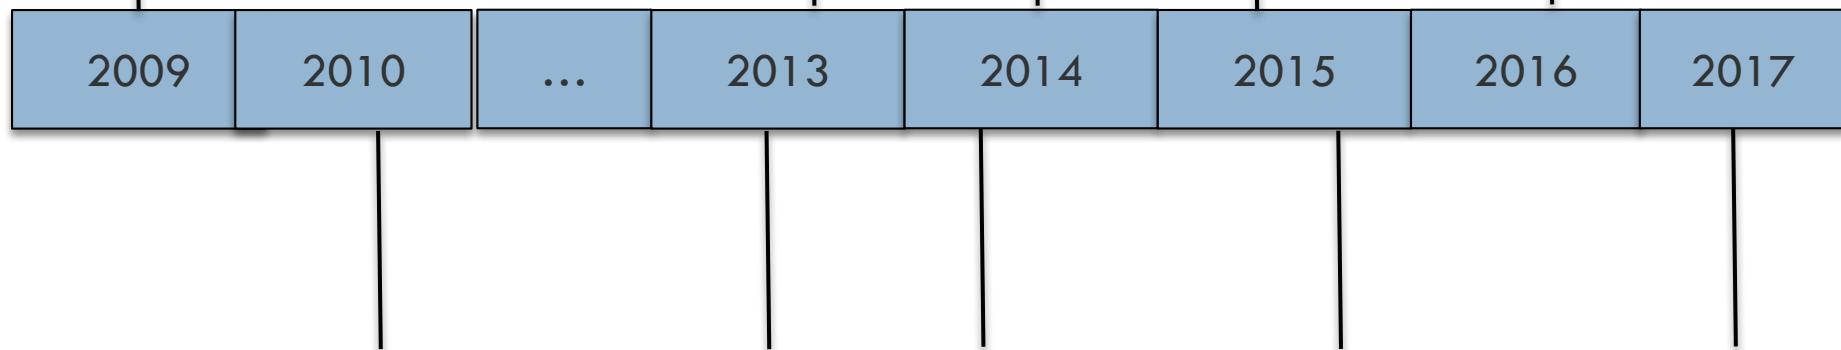
Lightning-fast cluster computing

Introduction to Apache Spark

11

Research project
at Berkeley
Matei Zaharia

Databricks
Founded \$43M
1.0 05/2014 1.2 02/2015 2.0 11/2016



Open sourced
on GitHub

Apache
incubator

Top level
project

1.5
09/2015

2.2
12/2017

Introduction to Apache Spark

12

Clip slide

On-Disk Sort Record:

Time to sort 100TB

2013 Record:
Hadoop

2100 machines



72 minutes



2014 Record:
Spark

207 machines



23 minutes



Also sorted 1PB in 4 hours

Source: Daytona GraySort benchmark, sortbenchmark.org

databricks

Introduction to Apache Spark

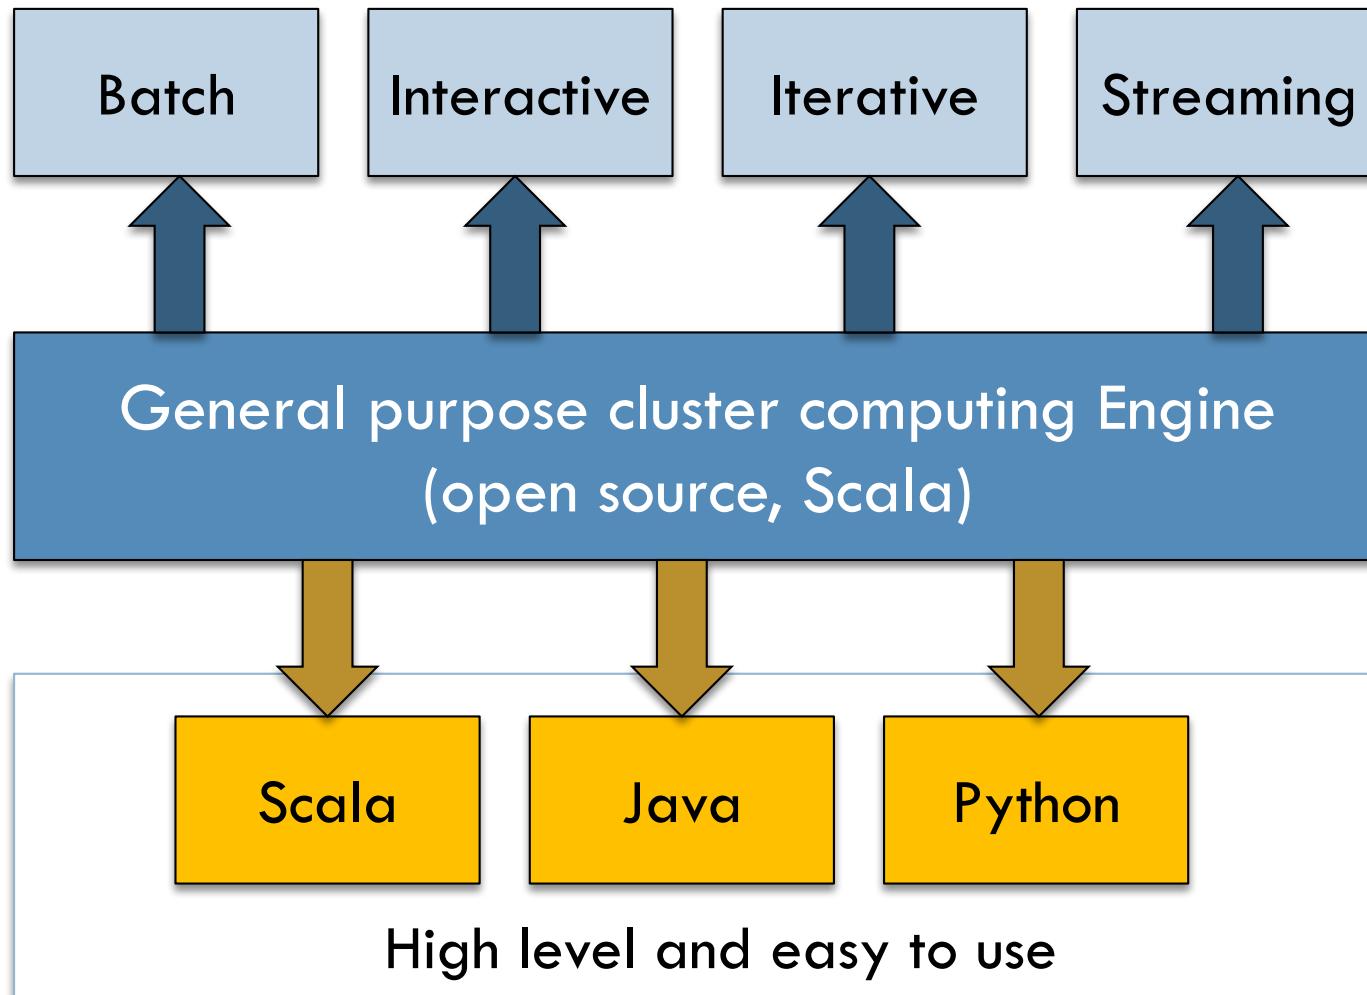
2014 Daytona GraySort Contest

	Hadoop MR Record	Spark Record
Data Size	100TB	100TB
Elapsed Time	72 mins	23 mins
# Nodes	2100	206
# Cores	50400	6592
Sort Rate	1.42 TB/min	4.27 TB/min
Sort Rate/node	0.67 GB/min	20.7 GB/min

3X faster and 10X less resources – Not using in-memory cache

Introduction to Apache Spark

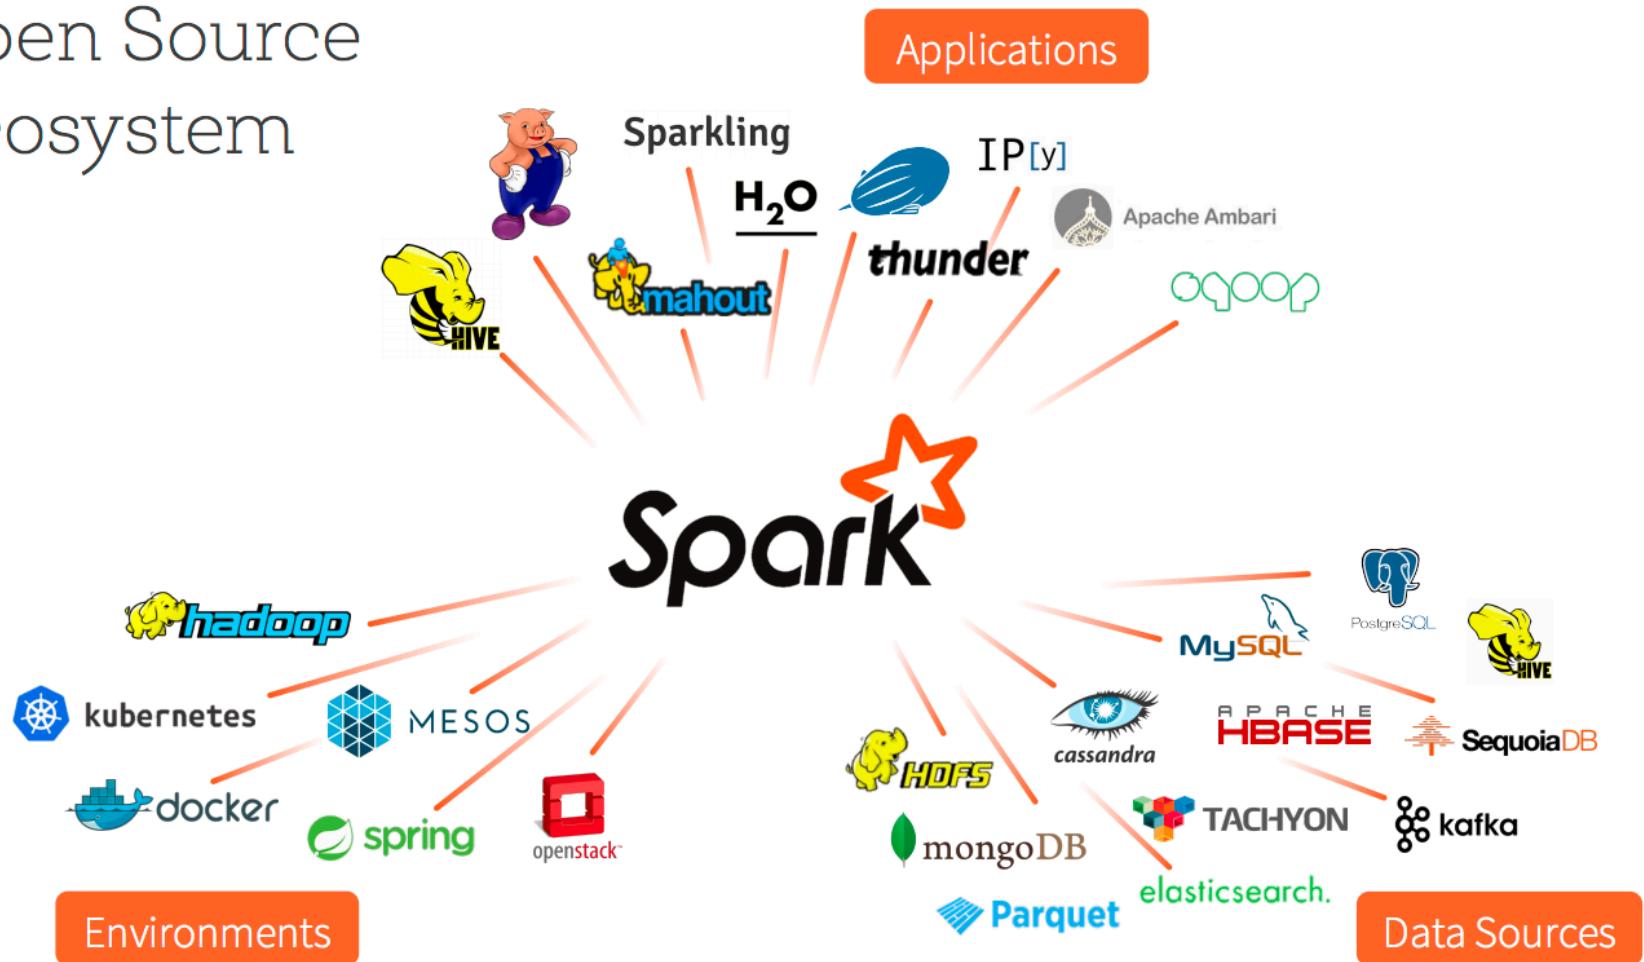
14



Introduction to Apache Spark

15

Open Source
Ecosystem



Introduction to Apache Spark

16

Adoptions



Introduction to Apache Spark

17

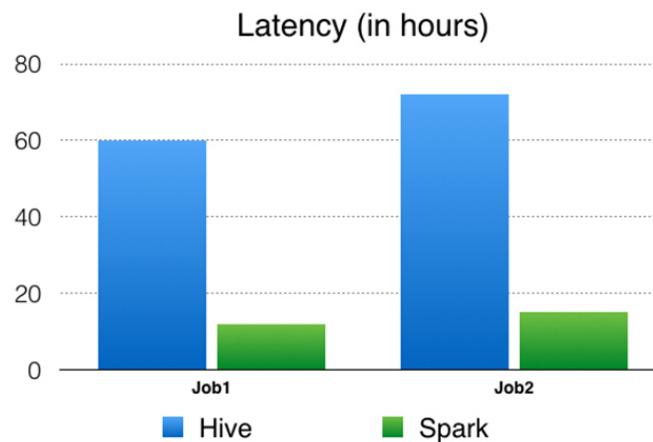
Notable Use Cases



Introduction to Apache Spark

18

- FB
 - Large-scale language model training
 - Converted Hive pipeline to Spark
 - 2.6x faster
 - A 60 TB+ production use case
 - Feature preparation for entity ranking



Introduction to Apache Spark

19

- Salesforce
 - Power insights on Sales Emails
 - Recommending hot sales lead to follow up
 - Recommending next actions based on email interactions
 - Power Einstein platform

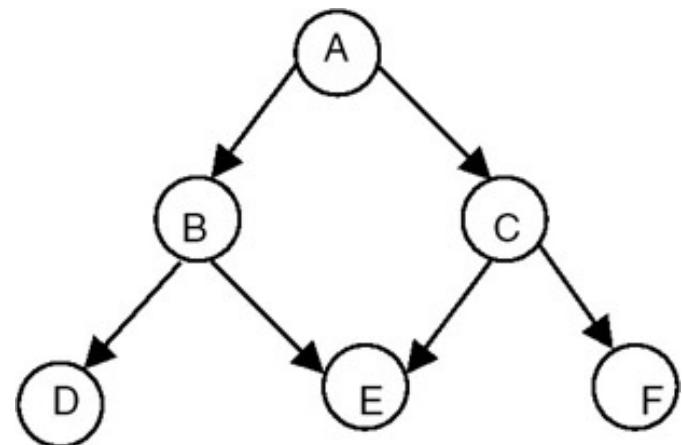
Introduction to Apache Spark

20

Resilient Distributed Dataset



In-memory computing primitives



DAG execution engine

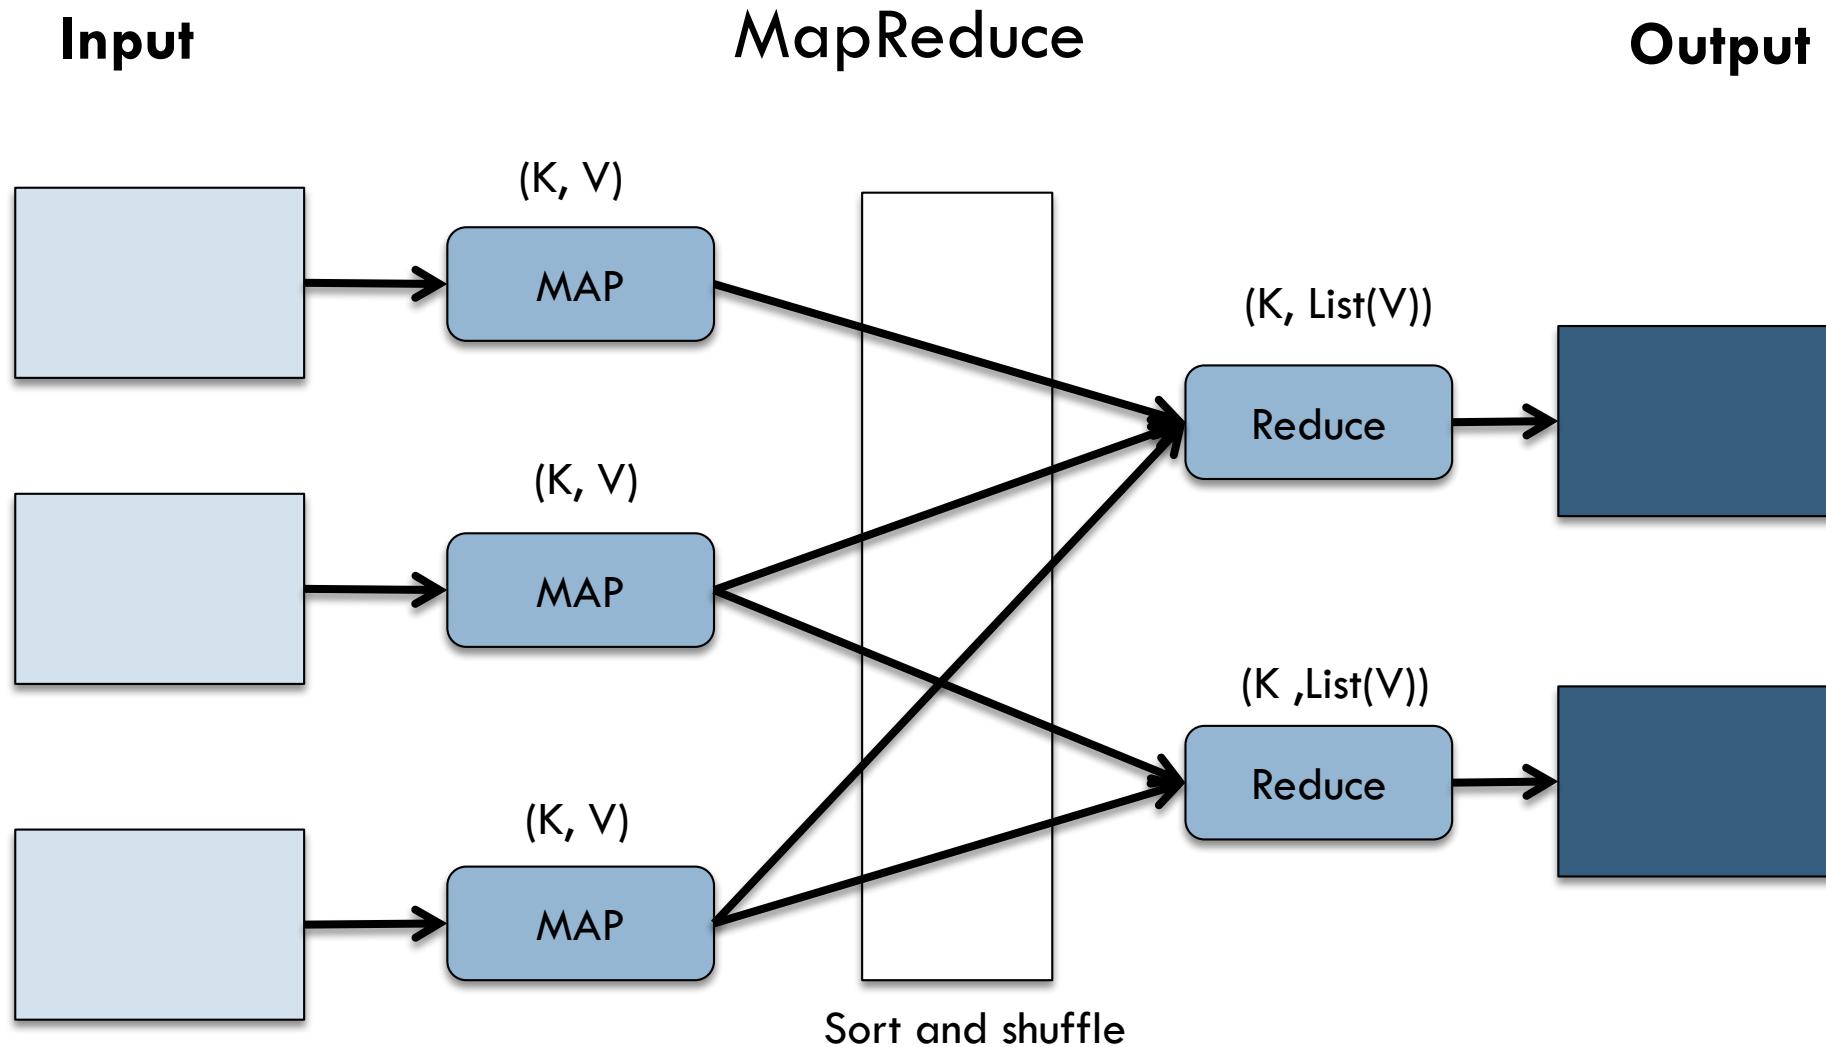
Introduction to Apache Spark

21

- MapReduce limitations
 - Difficult and tedious to program
 - Low level
 - Assembly language for data processing
 - Performance bottlenecks
 - Not easily used for non-batch use cases
 - Iterative, interactive, graph processing, streaming
 - Need to string together a series of MapReduce jobs

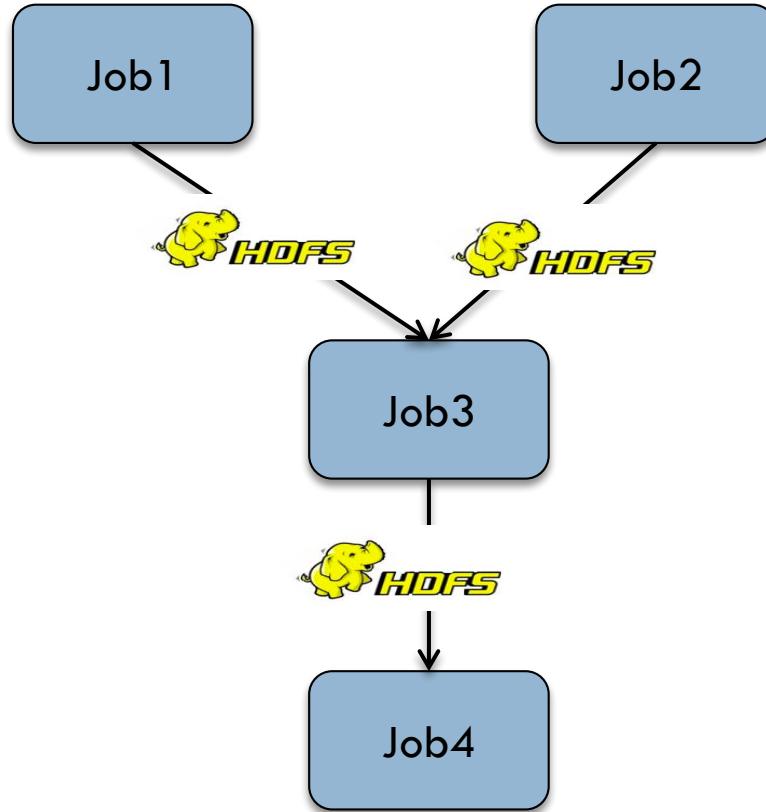
Introduction to Apache Spark

22



Introduction to Apache Spark

23



MapReduce doesn't support multi-stage jobs

Introduction to Apache Spark



24

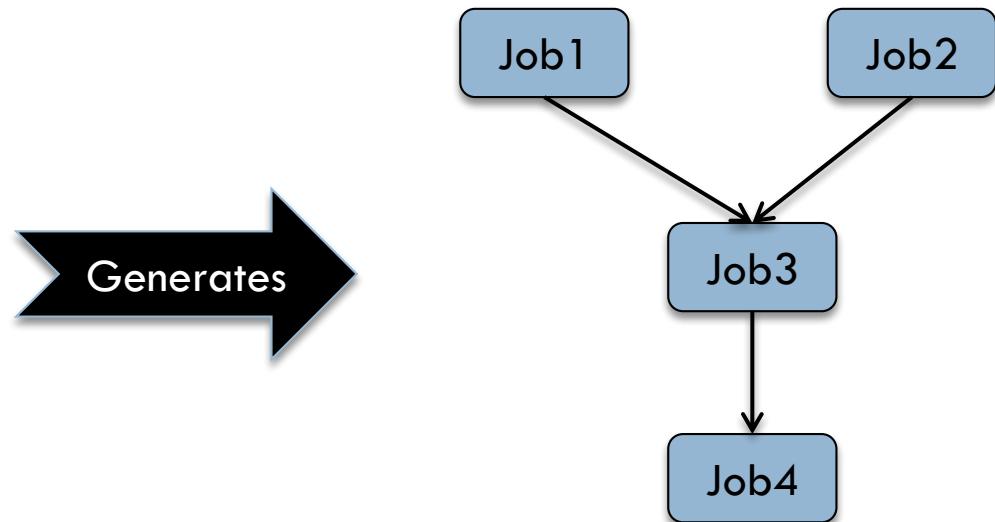
- Main abstraction is the relation
- Provide operators to manipulate the relations

```
// loading data file
input_lines = LOAD '/dataset' AS (line:chararray);

// transformation
- filter, group, join, order, union
```

```
member_account = LOAD
'/data/databases/MEMBE
R2/ MEMBER_ACCOUNT/#LAT
EST' using
LiAvroStorage();

member_account_filtere
d = FILTER
member_account BY
COUNTRY_CODE == 'us'
AND POSTAL_CODE is not
null;
```



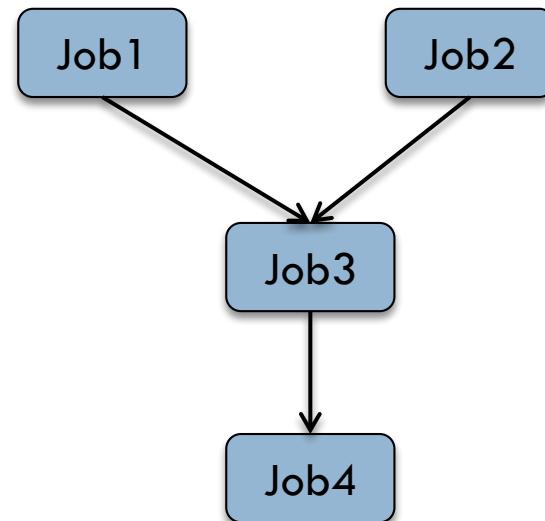
Introduction to Apache Spark



25

- Main abstraction is the table
- Provide SQL to manipulate the tables

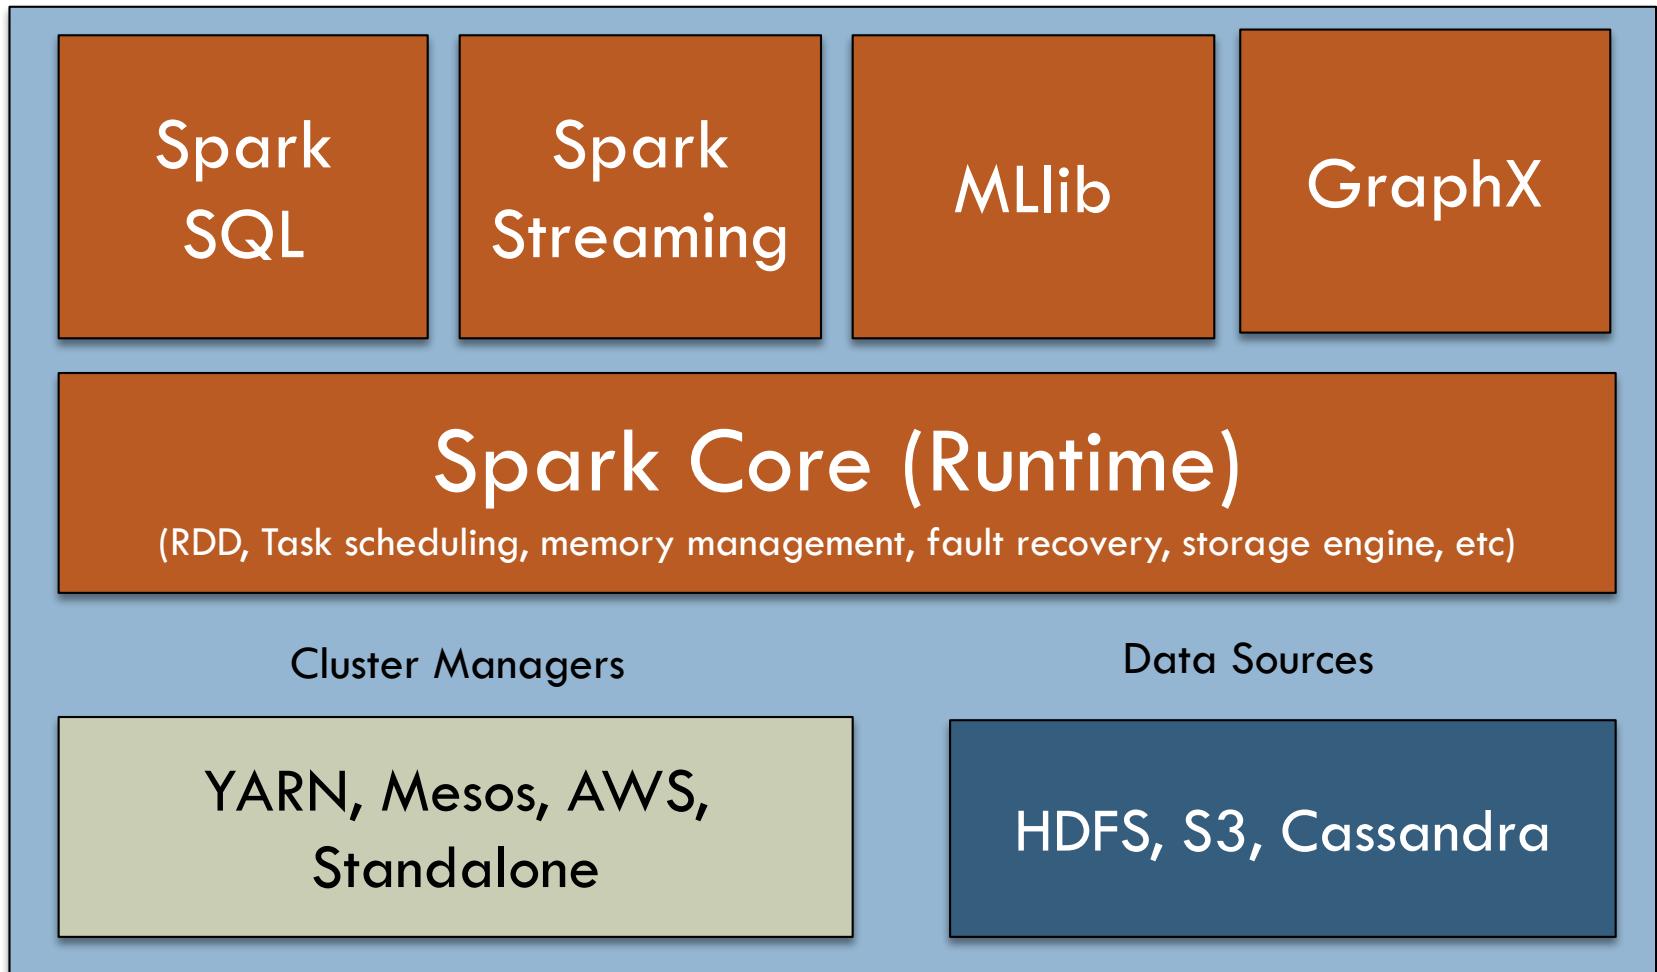
```
select country,  
       count(*) from  
       employee group  
       by country  
       having size >  
       100
```



Core Concepts and Architecture

26

Apache Spark Unified Stack



Core Concepts and Architecture

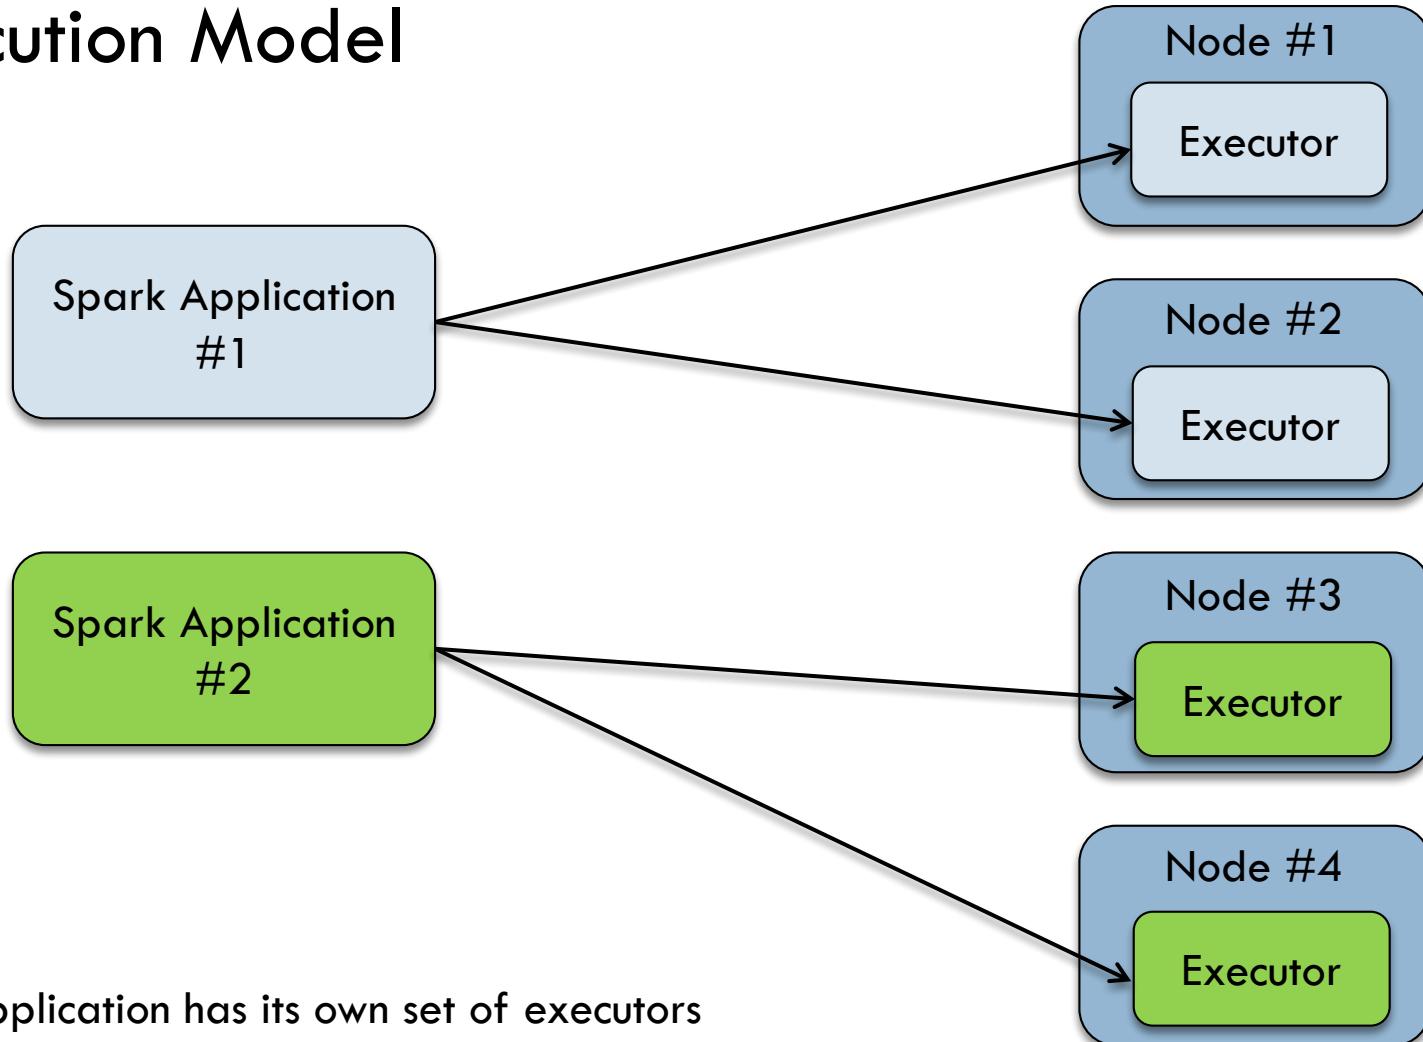
27

- Spark doesn't come with a storage layer
 - Unlike Hadoop
- Can easily integrate with HDFS and others
- Apache Alluxio (fka Tachyon)
 - Memory-centric distributed storage system
 - Reliable data-sharing at memory speed

Apache Spark Overview

28

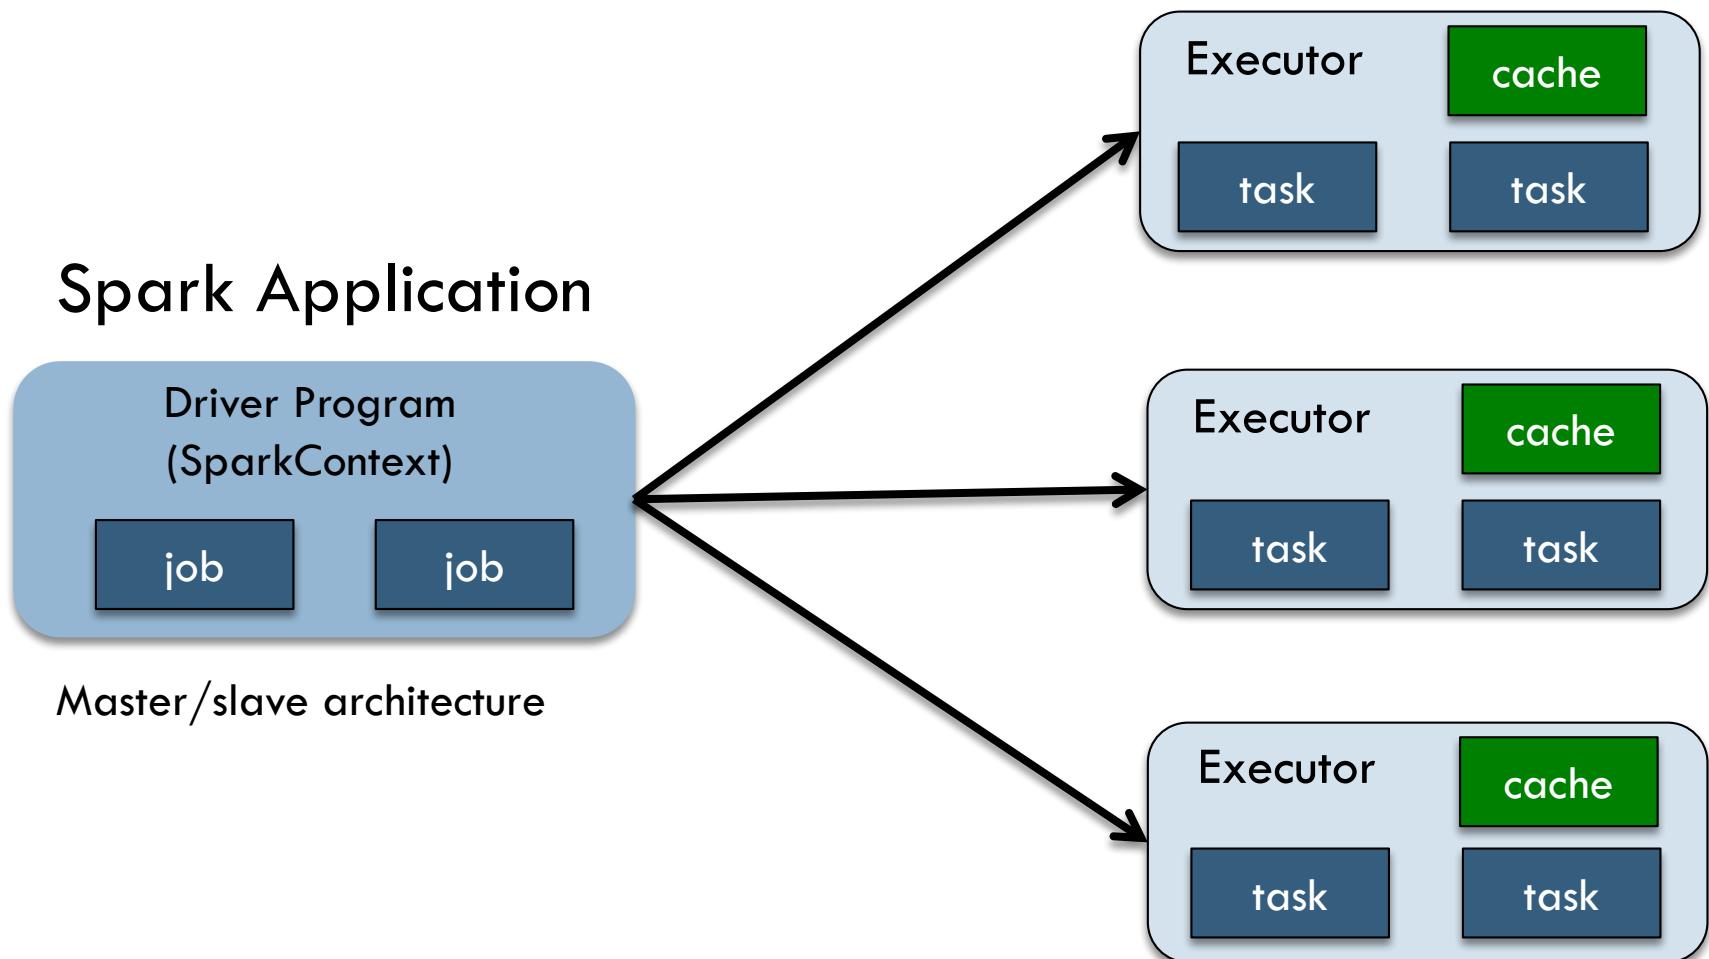
Execution Model



Core Concepts and Architecture

29

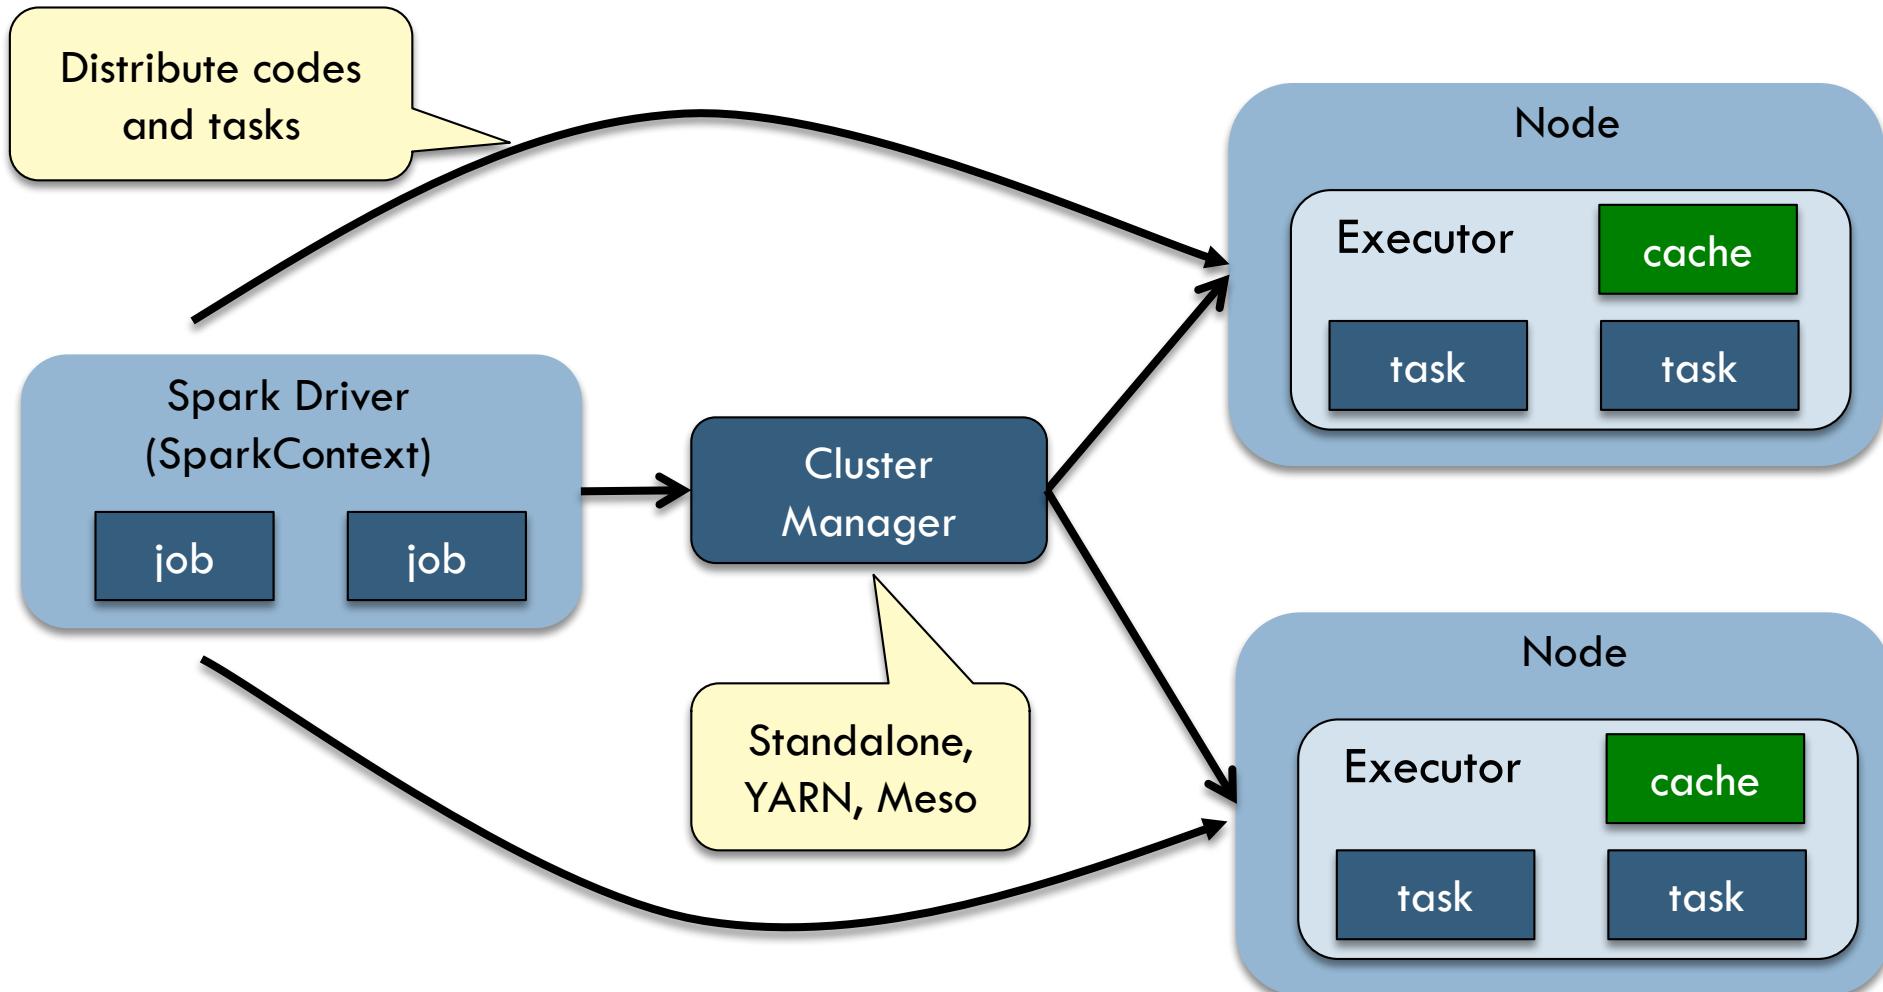
Spark Application Execution Model



Core Concepts and Architecture

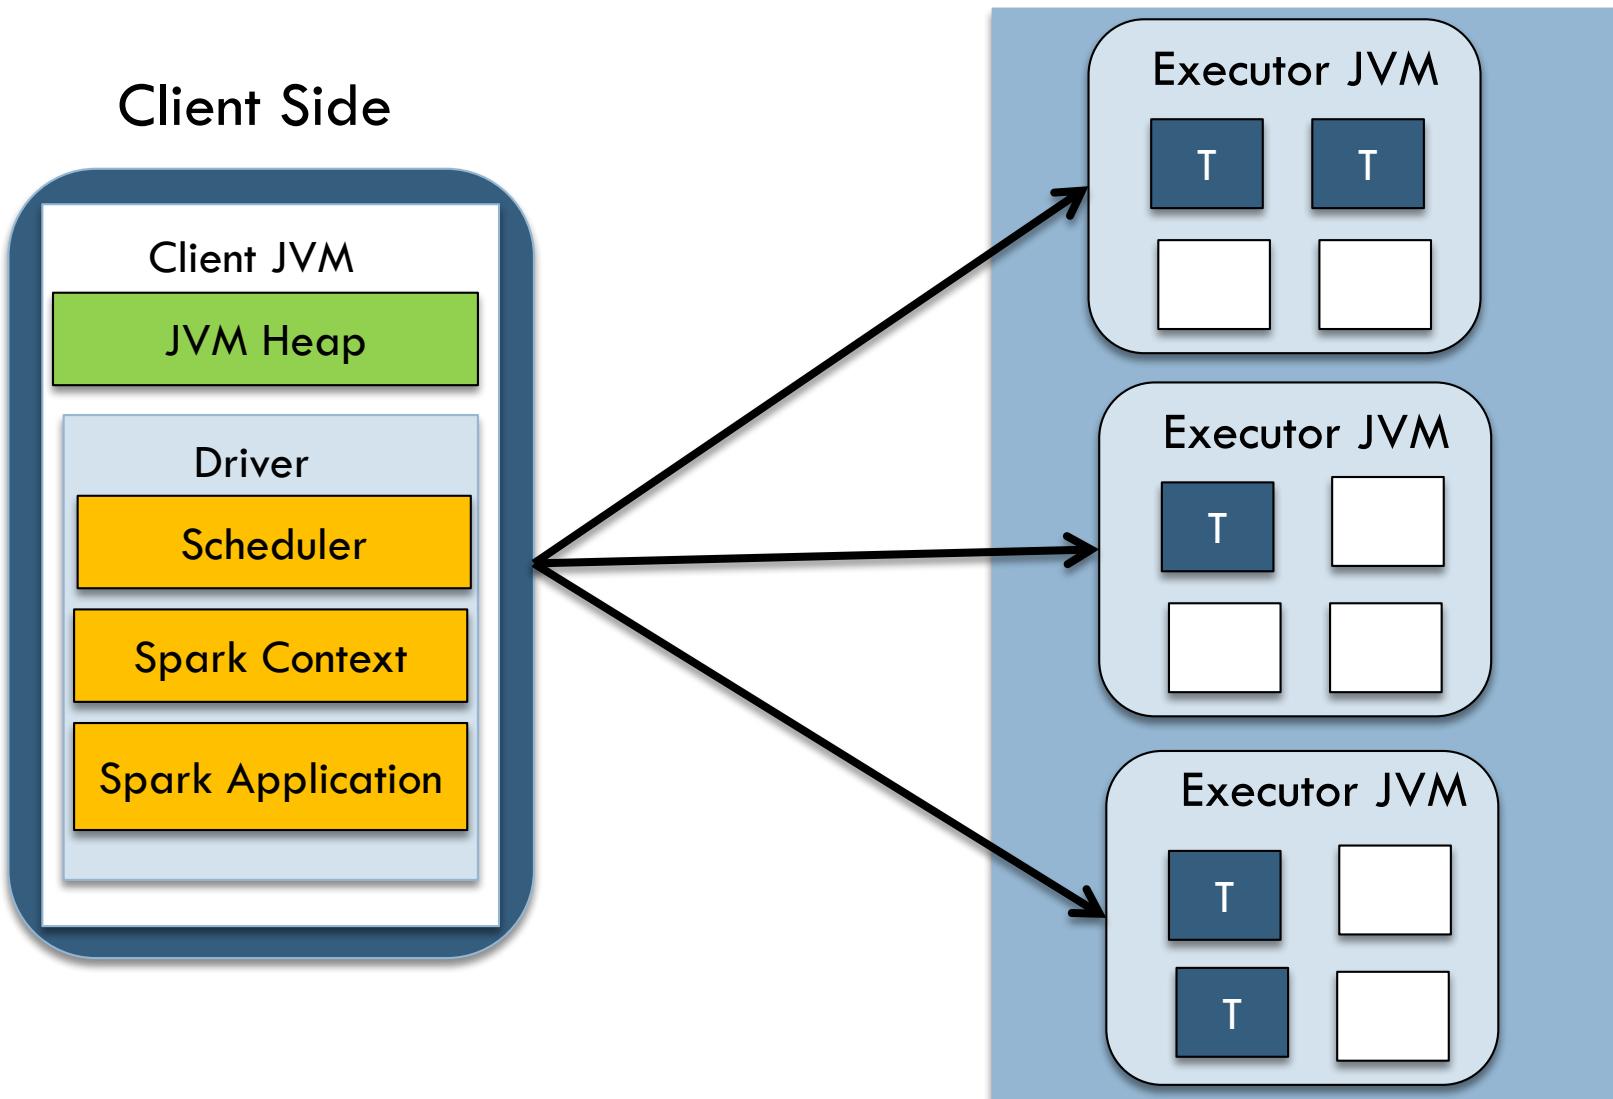
30

Spark Application Execution Model



Core Concepts and Architecture

31



Core Concepts and Architecture

32

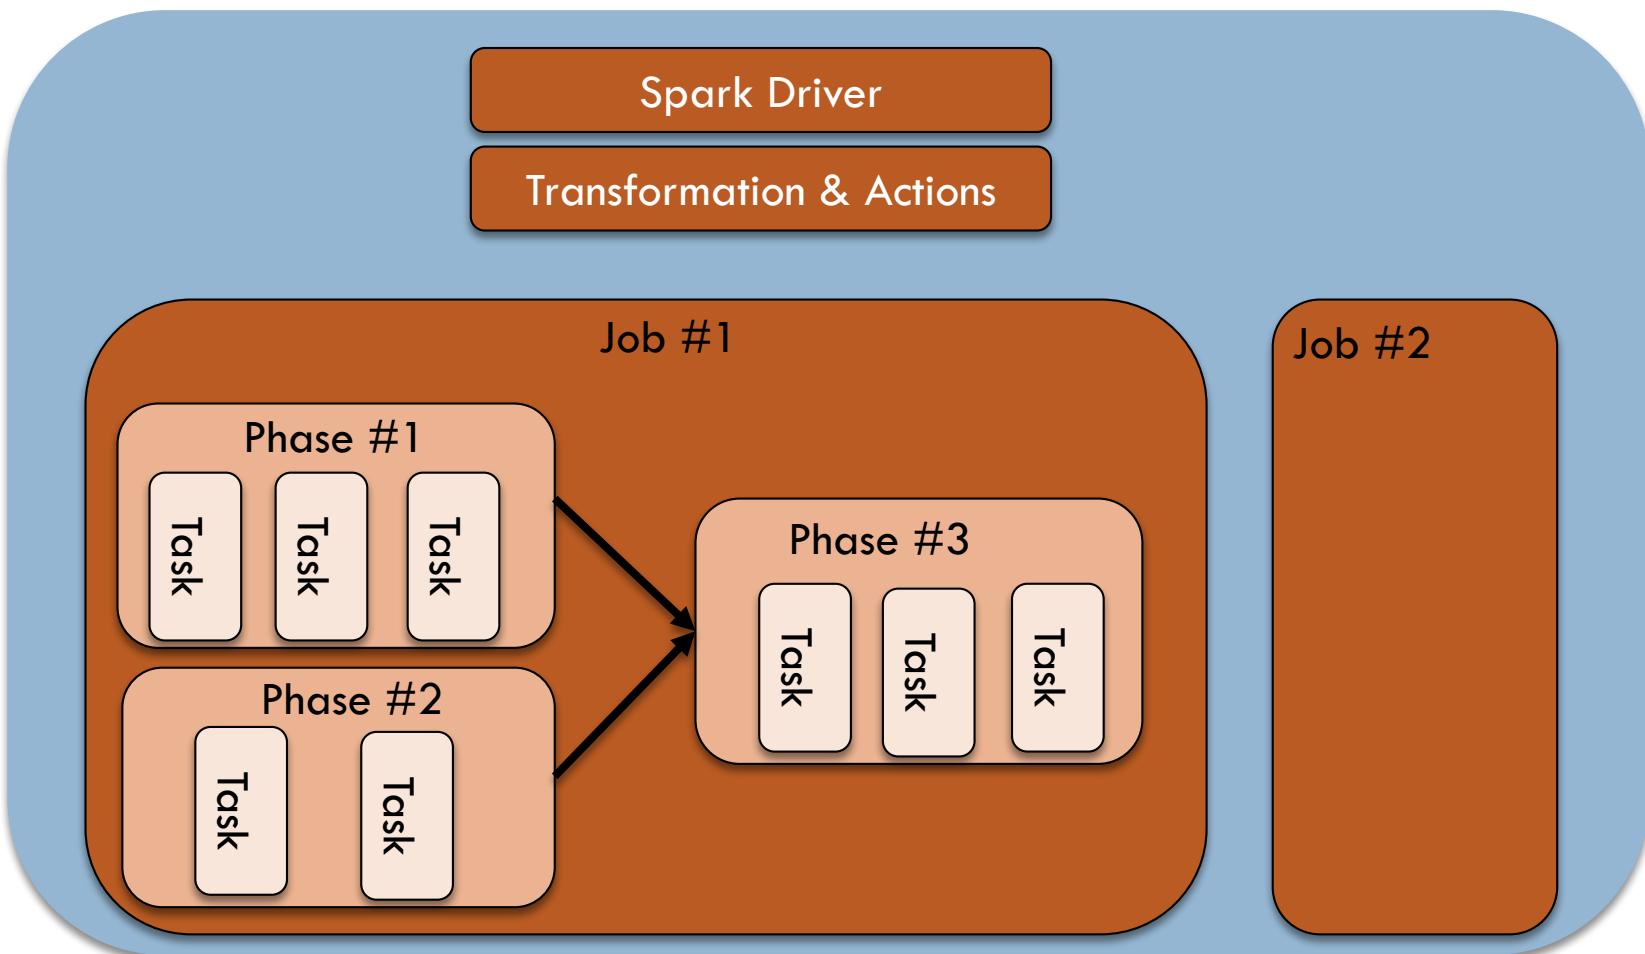
- Driver
 - Central coordinator
 - Optimization
 - Pipelining
 - Convert logic into set of stages
 - Each stage consists of multiple tasks
 - Tasks are smallest unit of work
 - Tasks are packaged and prepared to send to executors
 - Tasks are coordinated and scheduled to run on executors



Core Concepts and Architecture

33

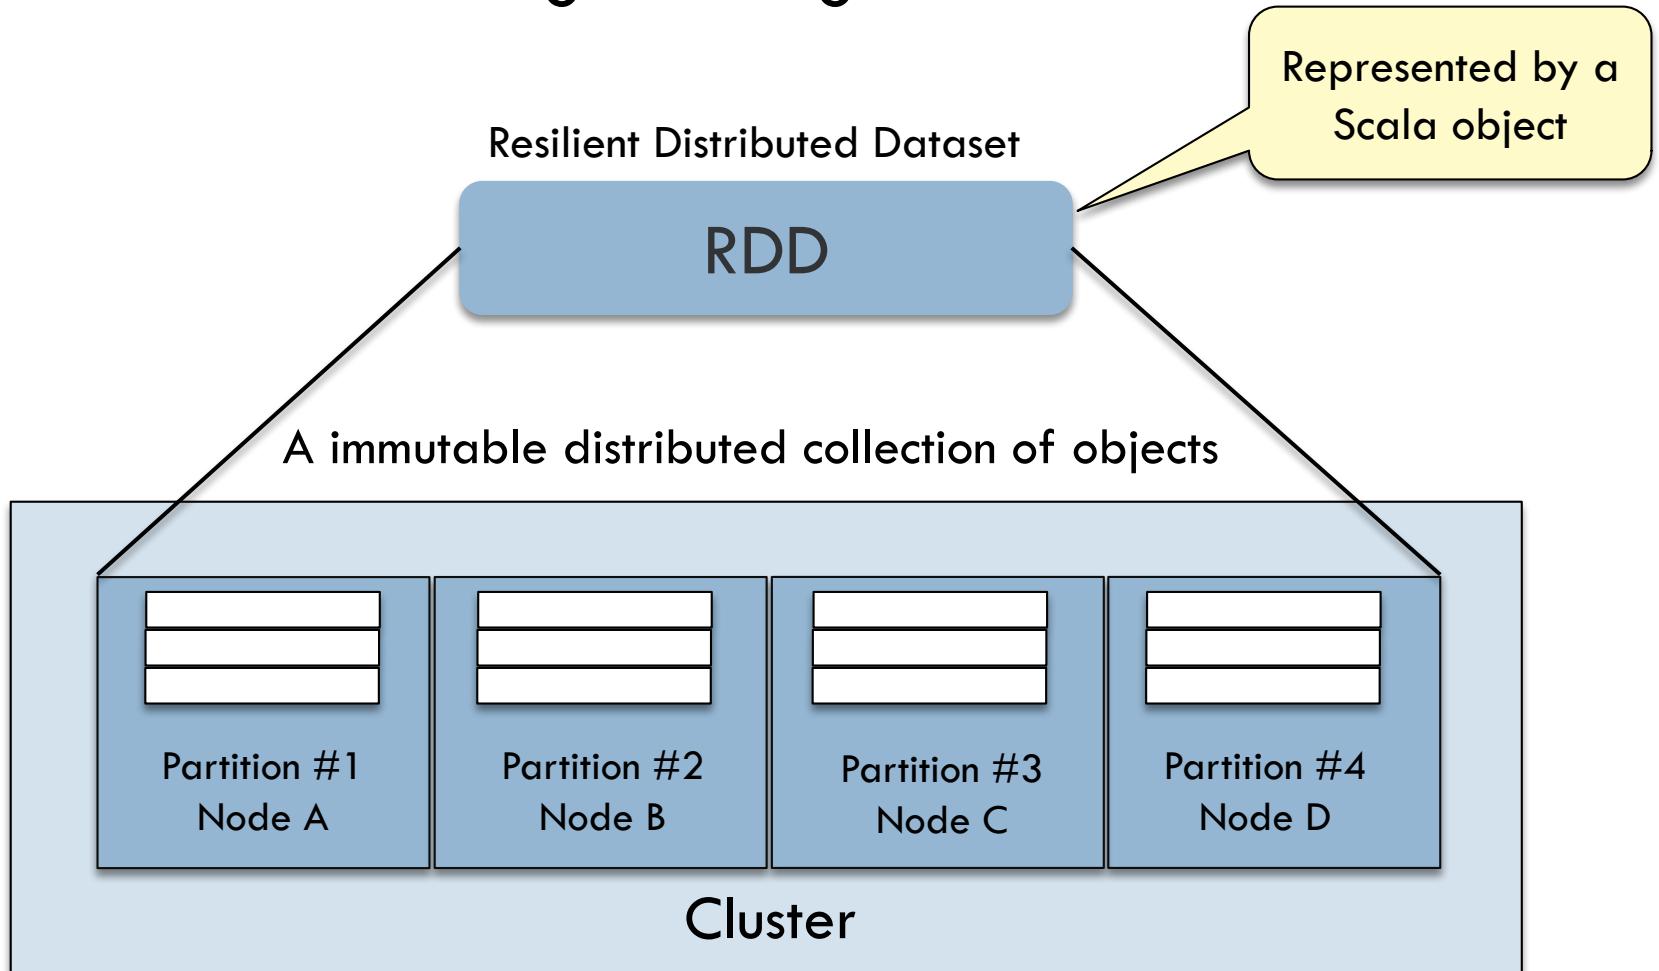
Spark Application



Apache Spark Overview

34

Programming Model



Core Concepts and Architecture

35

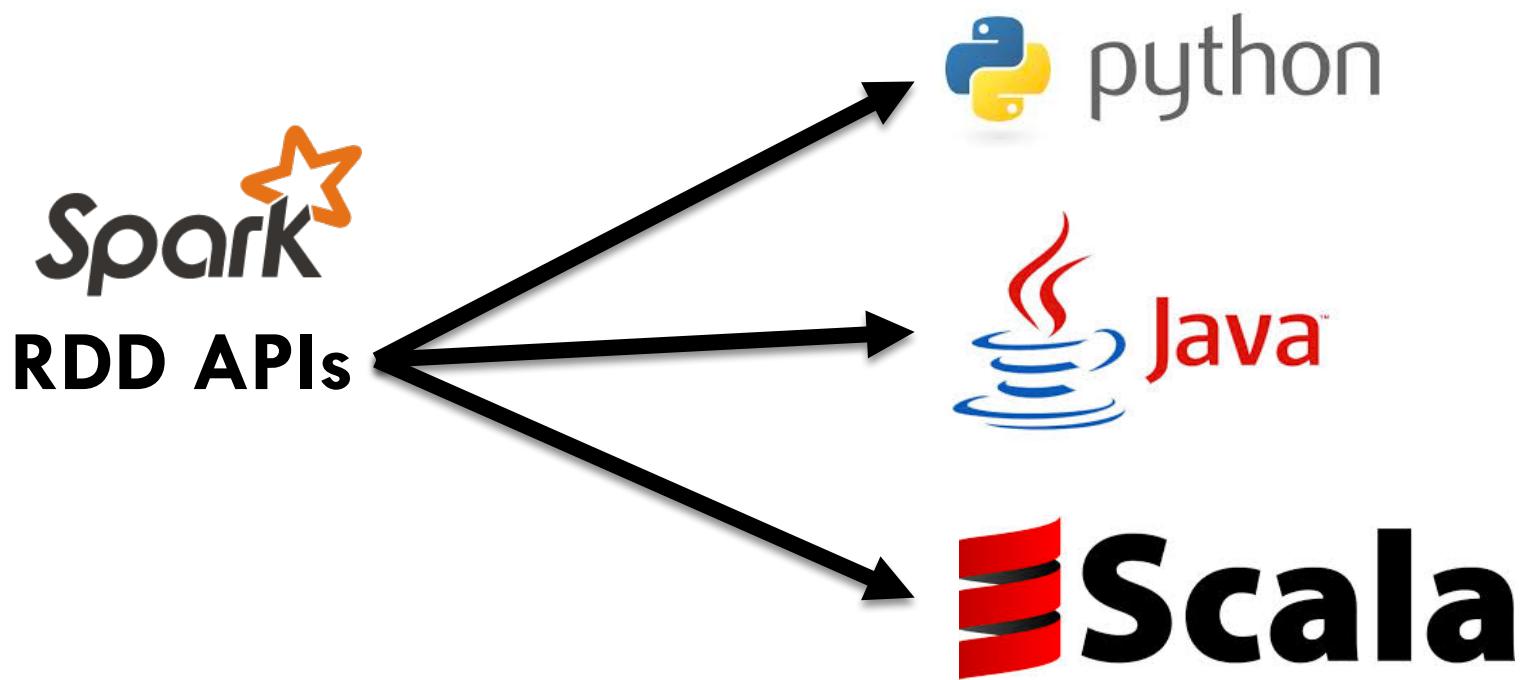
- RDD (Resilient Distributed Dataset)
 - Dataset
 - Initial data can come from a file or created programmatically
 - Distributed
 - Store across nodes in the cluster
 - Resilient
 - Automatic rebuilt on failure (through lineage)
- Able to persist intermediate results in memory
 - Spill to disk if not enough RAM
 - Controllable persistence

Core Concepts and Architecture

36

Power of Spark

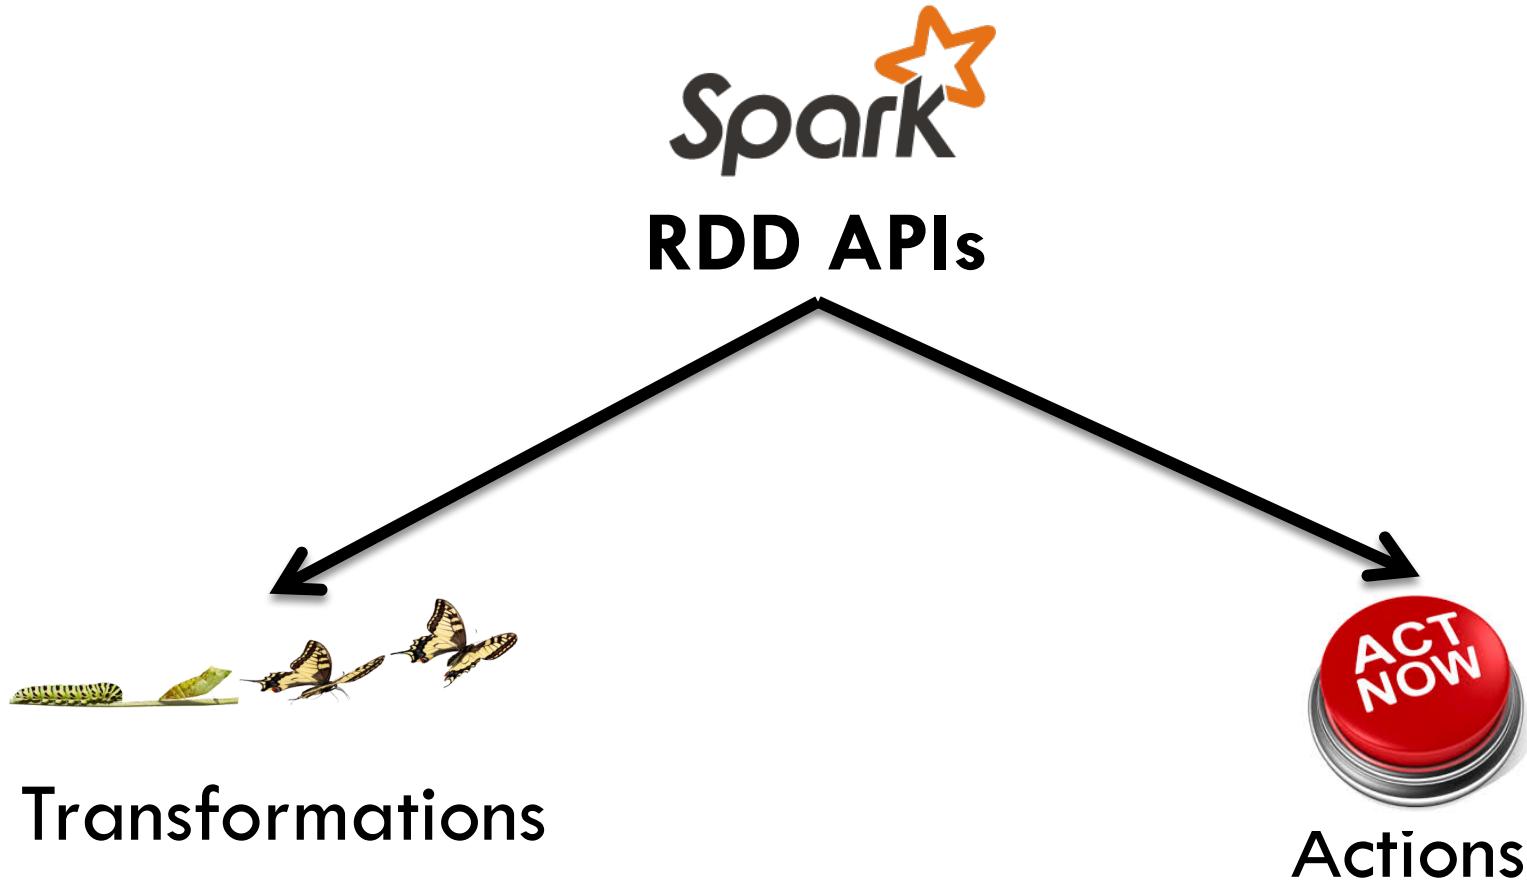
(Rich APIs and interactive shell)



Similar to Scala collection APIs, but for distributed data computation

Core Concepts and Architecture

37



Core Concepts and Architecture

38

- RDD programming interface
 - Transformations (e.g map, filter, groupBy, join)
 - Operations on RDD
 - **Create** new RDD from previous one
 - **Lazy** evaluation (when use an action), why?
 - Actions (e.g count, collect, save)
 - Return a result or write to storage
 - **Force** evaluation of the transformations
 - Restricted share variables
 - Broadcast and accumulators

Core Concepts and Architecture

39

- RDD programming interface
 - Broadcast variables
 - For a large read-only piece of data (lookup table)
 - Distribute and cache on each executor (machine)
 - Rather than tasks
 - Reduce communication costs
 - Accumulators
 - Can only “add” to using an associative operation
 - Only driver can read
 - Can be used to implement like counters in MapReduce

Core Concepts and Architecture

40

- RDD Persistence
 - One of the most important capabilities
 - Reading data from memory is extremely fast
 - Good interactive use cases and iterative algorithms
 - Support different storage level
 - Memory
 - Disk
 - Serialized or unserialized
 - LRU for eviction policy

Core Concepts and Architecture

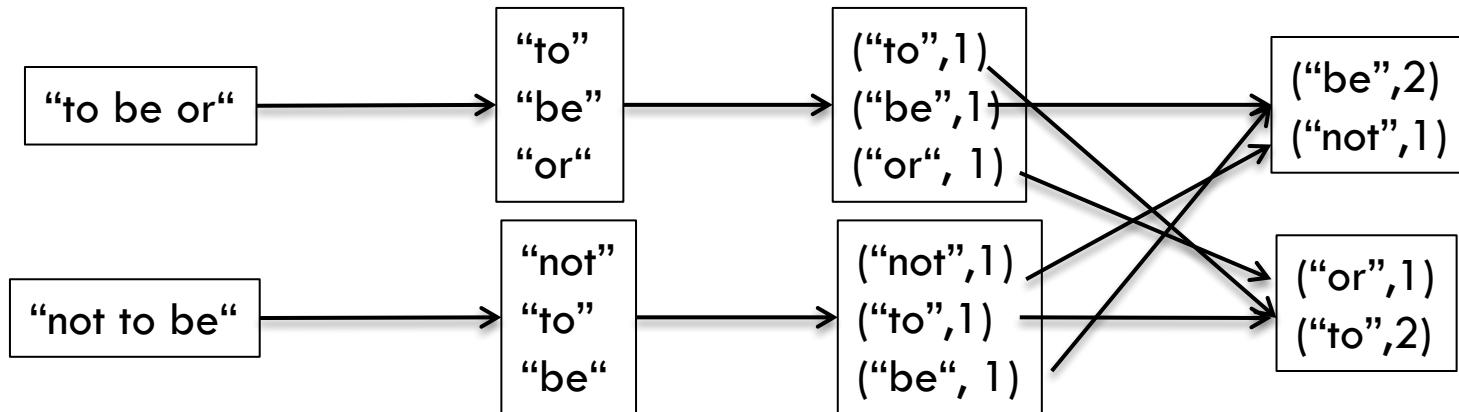
41

Word count example

```
// loading data file
val file = sc.textFile("README.md")

val words = file.flatMap(l => l.split(" "))
val wordCountPairs = words.map(word => (word, 1))
val wordCountSum = wordCountPairs.reduceByKey(_ + _)

// store output
wordCountSum.saveAsTextFile("/tmp/output")
```



Core Concepts and Architecture

42

```
// loading data file
val file = sc.textFile("README.md")

val words = file.flatMap(l => l.split(" "))
val wordCountPairs = words.map(word => (word, 1))
val wordCountSum = wordCountPairs.reduceByKey(_ + _)
// store output
wordCountSum.saveAsTextFile("/output")
```

