

INTRODUCTION TO SPARK WITH SCALA

Apache Spark Programing Model



Agenda

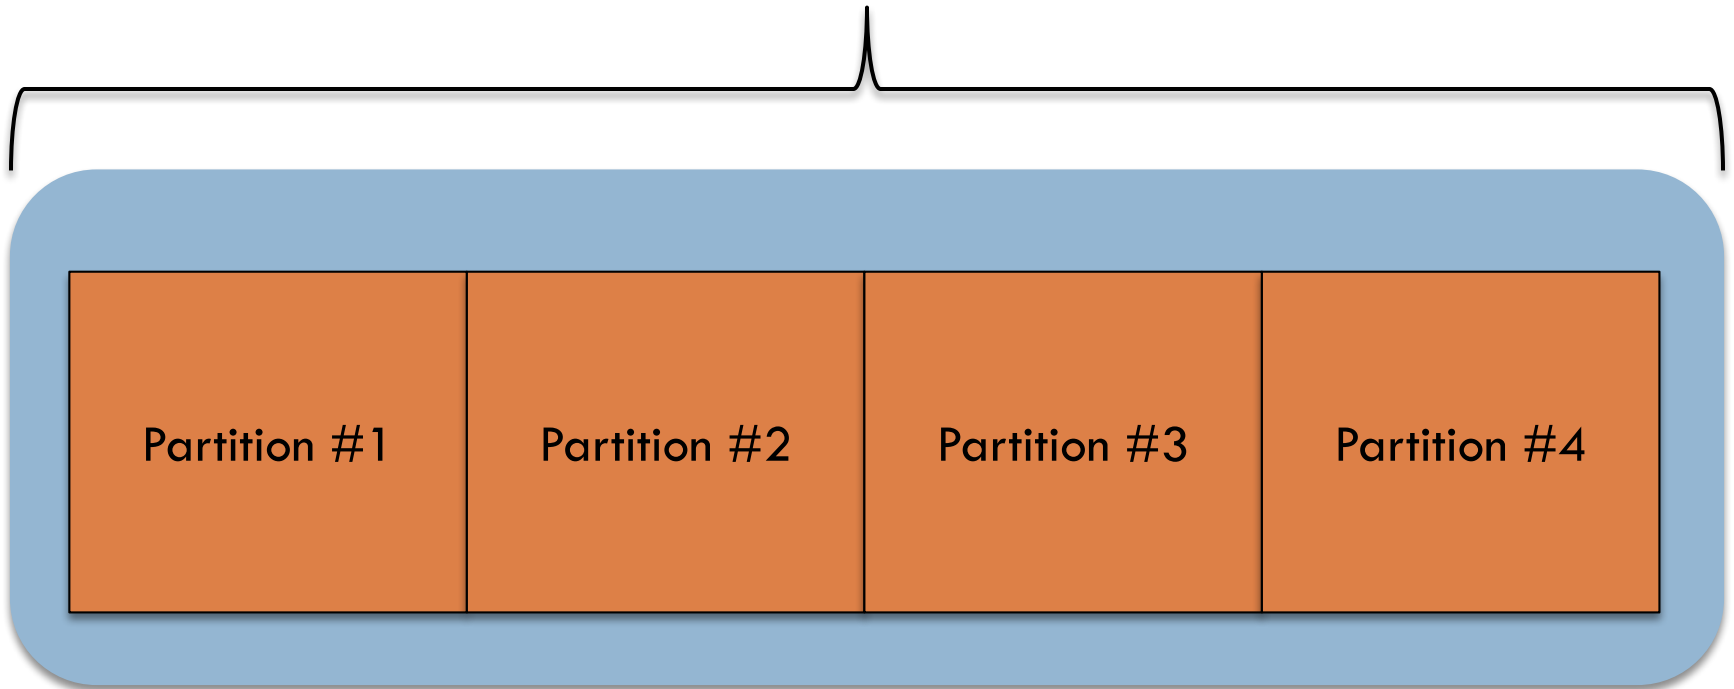
2

- RDD Dependency
- RDD Data Partitioning
- RDD Persistence
- RDD Loading & Saving
- RDD Shared Variables
- PageRank in Spark

RDD Dependency

3

RDD




Immutable partitioned data set distributed across multiple nodes


RDD Dependency

4


(map, filter)



(union)



Narrow Dependency




Join with co-partitioned

RDD Dependency


5

Wide Dependency

(groupByKey)



Multiple child RDD partitions may depend on a single parent RDD partition



Join without co-partitioned

RDD Data Partitioning


6

- Partitioning
 - ▣ Control the degree of parallelism
 - ▣ Group elements based on a function of each key
 - ▣ Ensure a set of keys appear together on some node
- Minimize data movement for performance
- Applicable when RDD is reused multiple times

RDD Data Partitioning

7

Shuffling while reduceByKey



Expensive Operation

- Disk I/O
- Data serialization
- Network I/O

Physical movement of data between partitions

RDD Data Partitioning

8

□ Use cases

- ▣ In key/value pair operations to ensure a set of keys reside on some node
- ▣ `join`, `reduceByKey`, `groupByKey`, `cogroup`, `sortByKey`

□ Partition options

- ▣ Hash partitioning
 - `Hash(key) & numPartitions`
- ▣ Range partitioning
 - Partition data elements into roughly equal range

RDD Data Partitioning

9

Data Partition APIs

| Transformation | Description |
|---|---|
| <code>partitionBy(partitioner)</code> | Return new RDD with specified number of partitions partition by given partitioner |
| <code>coalesce(numPartitions)</code> | Decrease number of partitions |
| <code>repartition(numPartitions)</code> | Reshuffle data randomly to create either more or fewer partitions and balance across them |

RDD Data Partitioning

10

partitionBy(partitioner) –

- New RDD with specified # of partitions and partition data elements based on given partitioner

```
val rdd1 = sc.textFile("<path>")
Val keyValueRdd = rdd1 .. Convert to key value pair RDD

val partitionedRDD = keyValueRdd.partitionBy(new
HashPartitioner(20))

// to find out # of partitions

partitionedRDD.partitions.length // 20
partitionedRDD.partitioner        // HashPartitioner
```

RDD Data Partitioning

11

coalesce(numPartitions) –

- ❑ Reduce number of partitions down to given partitions w/o shuffling
- ❑ Useful after filtering down a large dataset

```
import scala.util.Random
val randNums = Seq.fill(100000)(Random.nextInt(80000))
val numberCounts = sc.parallelize(randNums, 10)

val oddNumbers = numberCounts.filter(n => n % 2 == 1)

val rdd2 = oddNumbers.coalesce(3)
rdd2.partitions.length
```

RDD Data Partitioning

12

repartition(numPartitions) –

- Reshuffle data randomly to create more or fewer partitions

```
val rdd1 = sc.parallelize(1 to 15 by 3, 3)
rdd1.glom().collect


// Array(Array(1), Array(4, 7), Array(10, 13))

val rdd2 = rdd1.repartition(2)
rdd2.glom().collect

//Array(Array(1, 4, 10), Array(7, 13))
```

RDD Data Partitioning

13




```
val memberData = sc.textFile("<path>")
val pageViewData = sc.textFile("<path>")

memberData.join(pageViewData)
```

RDD Data Partitioning

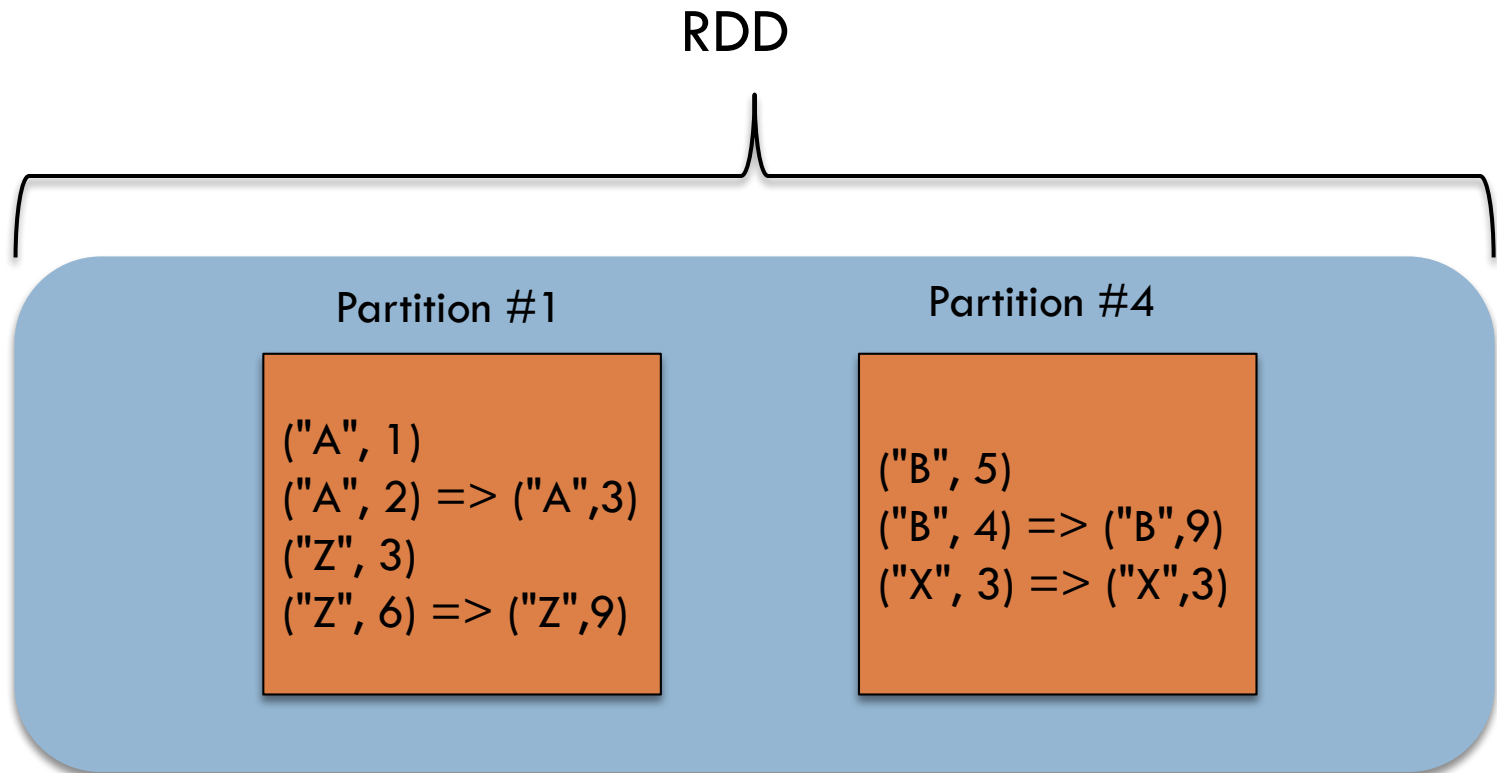
14



```
val memberData =  
sc.textFile("<path>").partitionBy(..).persist()  
val pageViewData = sc.textFile("<path>")  
memberData.join(pageViewData)
```

RDD Data Partitioning

15



```
rdd.reduceByKey(_ + _).collect()  
  
// ("A", 3), ("Z", 9), ("B", 9), ("X", 3)
```

RDD Persistence

16

- RDD Persistence/caching
 - ▣ Important and unique feature in Spark
 - ▣ Persisting dataset in memory across operations
 - ▣ Intermediate result of a computation – reuse
 - ▣ Great for interactive use case or iterative algorithms
 - ▣ Fault-tolerant caching
 - ▣ Caching storage options
 - Tradeoff between memory usage and CPU usage
 - Memory, disk
 - Serialized vs deserialized objects
 - ▣ LRU eviction policy

RDD Persistence

17

| Caching Option | Description |
|---------------------------------|---|
| MEMORY_ONLY (default) | Store as Java deserialized Java objects. Recompute partitions that don't fit in memory on the fly |
| MEMORY_AND_DISK | Store as Java deserialized Java objects. Spill onto disk for don't fit partitions |
| MEMORY_ONLY_SER | Store as serialized Java objects. Space efficient but more CPU-intensive |
| MEMORY_AND_DISK_SER | Similar to MEMORY_ONLY_SER, but spill out to disk for don't fit partitions |
| DISK_ONLY | Store only on disk |
| MEMORY_ONLY2 MEMORY_AND_DISK | Same ask above, but replicate each partition on two nodes |
| OFF_HEAP | Store in serialized format on Tachyon |

RDD Persistence

18

Caching Option Tradeoffs

| Caching Option | Space Used | CPU Time | In Memory | On Disk |
|-----------------------|------------|----------|-----------|---------|
| MEMORY_ONLY (default) | High | Low | Y | N |
| MEMORY_AND_DISK | High | Medium | Some | Some |
| MEMORY_ONLY_SER | Low | High | Y | N |
| MEMORY_AND_DISK_SER | Low | High | Some | Some |
| DISK_ONLY | Low | High | N | Y |

RDD Persistence

19

`persist(storageLevel)` – persist in cache

`unpersist()` – remove from cache

```
// loading data file
val logs = sc.textFile("hdfs://service-log/")

val errors = logs.filter(_.contains("ERROR"))
errors.persist()

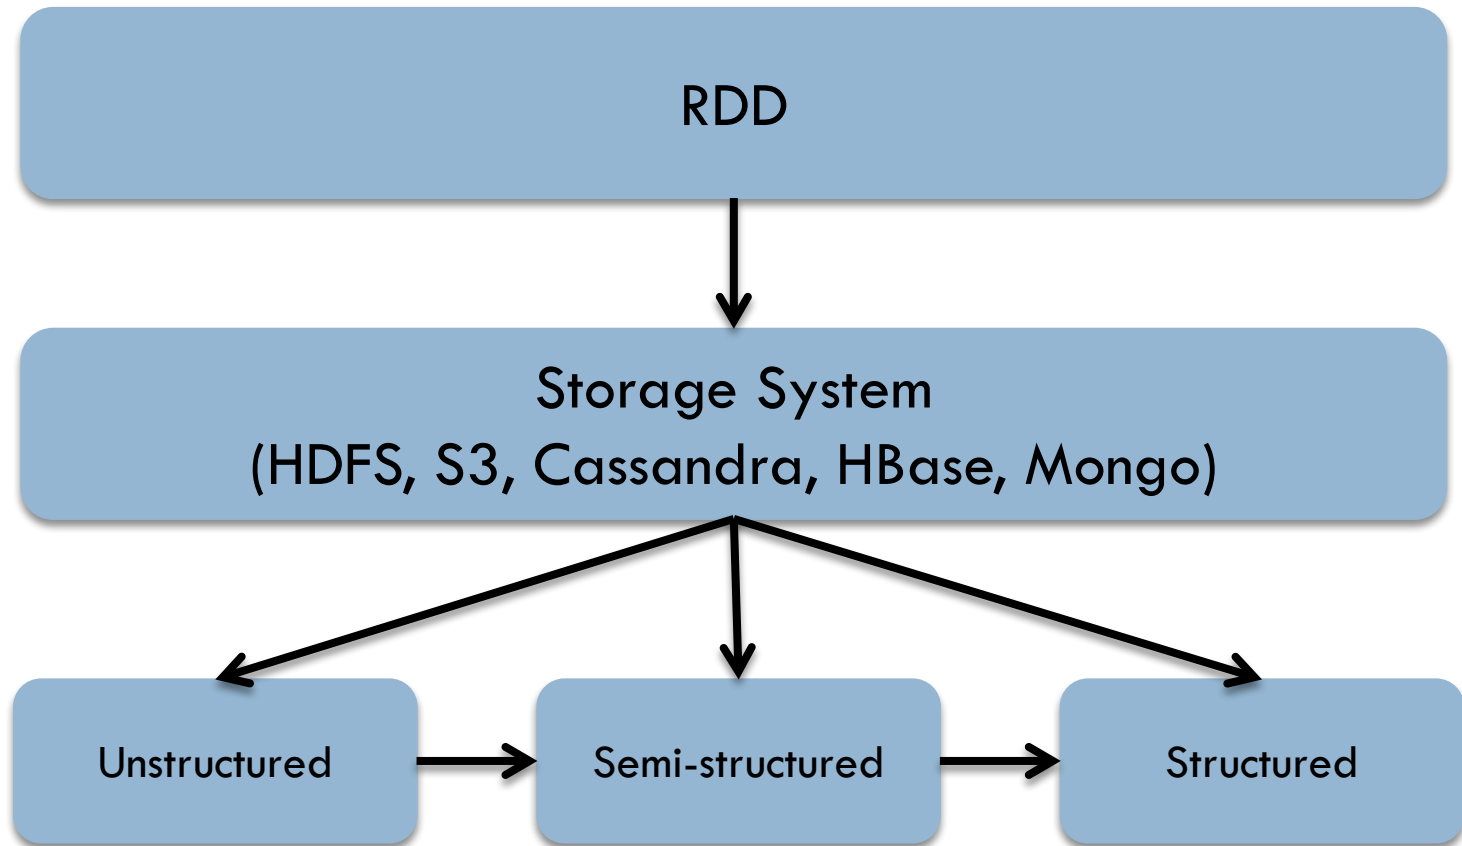
// figure out what kind of errors
errors.filter(_.contains("SQLException")).count

// should be faster at this point
errors.filter(_.contains("RemoteException")).count

// to evict RDD from cache
// errors.unpersist()
```

RDD Loading & Saving

20



RDD Loading & Saving

21

File Formats

| File Format | Structured | Description |
|------------------|------------|--|
| Text | No | Records are assumed to be one per line |
| JSON | Semi | Common text-based format |
| CSV | Yes | Popular text-based format |
| Protocol Buffers | Yes | A fast space efficient cross language format |
| Avro | Yes | A fast, space efficient, schema based format |
| Parquet | Yes | A columnar storage format |
| ORC | Yes | A columnar storage format |

RDD Loading & Saving

22

textFile(uri, numPartitions) –

- Lazily loading data set

```
// load a specific file
val rdd1 = sc.textFile("/data/part1.txt")
val rdd1 = sc.textFile("file:///data/part1.txt")

// load all files in /data
val rdd2 = sc.textFile("file:///data")

// load all files end with ".txt" in /data on HDFS
val rdd3 = sc.textFile("hdfs:///data/*.txt")

// load all compressed files end with ".tgz" in /data
val rdd4 = sc.textFile("/data/*.gz")
```

RDD Loading & Saving

23

wholeTextFiles(uri, numPartitions) –

- Read files in a directory
- Return tuples with (file name, content)

```
val rdd1 = sc.wholeTextFiles("/data")

val fileNames = rdd1.map({
  case(fn, content) => (fn, content.split(" ").length)
})

// file name with # of words in each file
fileNames.collect()
```

RDD Loading & Saving

24

saveAsTextFile(path) –

- Save RDD as text file(s) to given path
- Path shouldn't exist
- One part file per partition

```
// load a specific file
val rdd1 = sc.textFile("file:///data/part1.txt")

val rdd2 = rdd1.filter(line => line.contains("ERROR"))

// rdd2.saveAsTextFile("/data-with-error")

rdd2.saveAsTextFile("file:///data-with-error")
```


RDD Loading & Saving

25

- JSON Support
 - ▣ Better in SparkSQL
- JDBC
 - ▣ Through JdbcRDD
- Cassandra
 - ▣ Datastax Cassandra connector
- HBase
 - ▣ `SparkContext.newAPIHadoopRDD()`
 - ▣ `TableInputFormat`
- ElasticSearch
 - ▣ `ElasticSearch InputFormat`

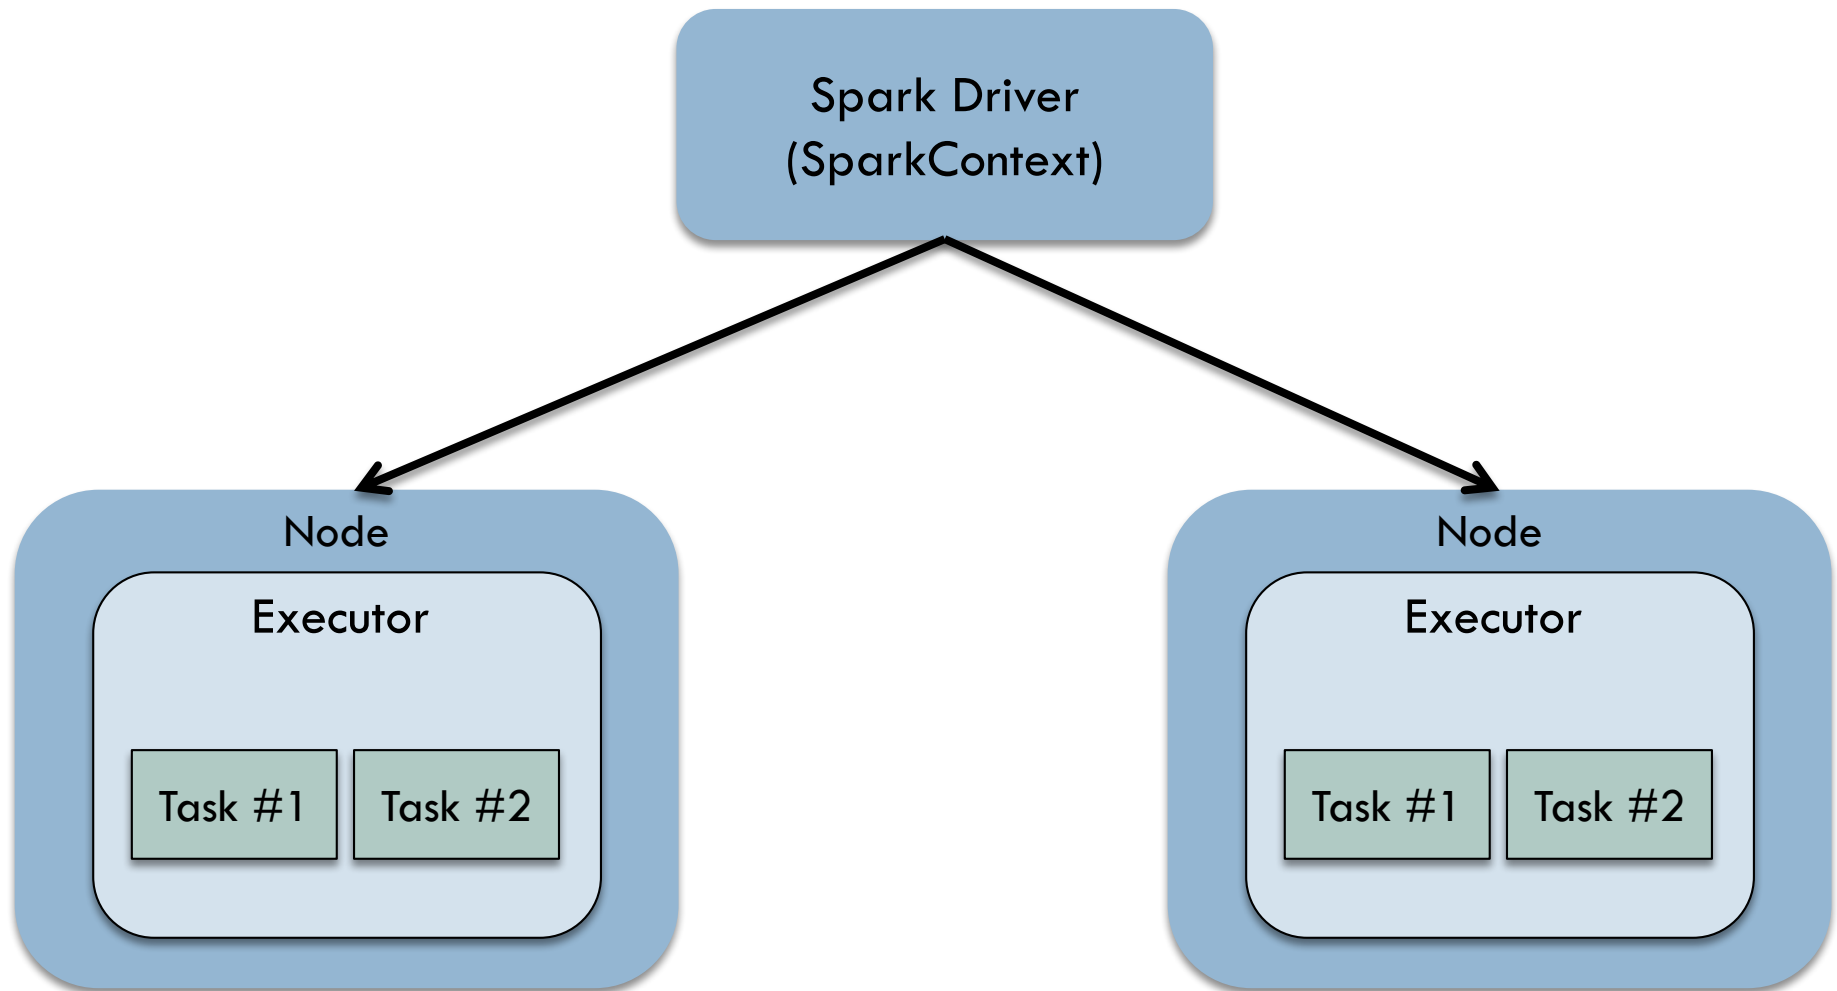
RDD Shared Variables

26

- Accumulators – distributed counters
 - ▣ Aggregate values from workers
 - ▣ Can only “add” to using an associative operation
 - ▣ Only driver can read the values
 - ▣ Can be used to implement counters like in MapReduce
 - ▣ Generally used for debugging purposes only
 - ▣ Out of the box support
 - Int, Double, Long, Float
 - ▣ Custom accumulator is supported

RDD Shared Variables

27



RDD Shared Variables

28

Regular Variable Example

```
// load a specific file
val rdd1 = sc.textFile("file:///data/part1.txt")

// create a variable
var badLines = 0;

// logic below is executed on each of the workers
val rdd2 = rdd1.flatMap(line => {
    if (line == "") {
        badLines += 1
    }
    line.split(" ")
})
rdd2.saveAsText("file:///output")
println("Bad lines: " + badLines)
```

RDD Shared Variables

29

Accumulator Example

```
// load a specific file
val rdd1 = sc.textFile("file:///data/part1.txt")

// create an accumulator
val badLines = sc.accumulator(0)

// logic below is executed on each of the workers
val rdd2 = rdd1.flatMap(line => {
    if (line == "") {
        badLines += 1
    }
    line.split(" ")
})

rdd2.saveAsText("file:///output") // needed?
println("Bad lines: " + badLines.value)
```

RDD Shared Variables

30

□ Broadcast Variables

- ▣ Efficiently share large, read-only value to worker nodes
- ▣ Efficient broadcast algorithm to reduce network costs
- ▣ When multiple tasks across stages need same data
- ▣ Read-only lookup table
 - Large feature vector
- ▣ Efficiently join when one of dataset is small

RDD Shared Variables

31

Broadcast Variables Example

```
// load a specific file
val rdd1 = sc.textFile("file:///data/part1.txt")
val countryMap = loadContryMap(...)

val countryMapBC = sc.broadcast(countryMap)

// countryMapBC will be broadcasted to each nodes
val rdd2 = rdd1.map ({line =>
    val row = line.split(",")
    (row(0), countryMapBC.value.get(row(1)))
})

rdd2.saveAsText("file:///output")
```

PageRank in Spark

32

- Showing off the power of Spark
- Iterative application
 - ▣ Join and shuffling
- Partitioning to reduce shuffling

PageRank works by counting the number and quality of links to a page to determine a rough estimate of how important the website is. The underlying assumption is that more important websites are likely to receive more links from other websites.

PageRank Algorithm

33

PageRank works by counting the number and quality of links to a page to determine a rough estimate of how important the website is. The underlying assumption is that more important websites are likely to receive more links from other websites.


- Wikipedia (<https://en.wikipedia.org/wiki/PageRank>)

PageRank in Spark

34

□ PageRank Algorithm

- Give page ranking scores based on # links to them
- Links from many pages scores higher rank
- Link from a high-rank page boosts rank higher




PageRank in Spark

35

□ Algorithms

- Start each page at a rank of 1
- On each iteration, have page p contribute
 - $\text{rank}_p / \text{neighbors}_p$
- Set each page's rank to $0.15 + 0.85 * \text{contribs}$




PageRank in Spark

36

□ Algorithms

- On each iteration, have page p contribute
 - $\text{contrib}_p = \text{rank}_p / \text{neighbors}_p$
- Set each page's rank to $0.15 + 0.85 * \text{contribs}$




PageRank in Spark

37

□ Algorithms

- On each iteration, have page p contribute
 - $\text{contrib}_p = \text{rank}_p / \text{neighbors}_p$
- Set each page's rank to $0.15 + 0.85 * \text{contribs}$




PageRank in Spark

38

□ Algorithms


- On each iteration, have page p contribute
 - $\text{contrib}_p = \text{rank}_p / \text{neighbors}_p$
- Set each page's rank to $0.15 + 0.85 * \text{contribs}$



PageRank in Spark

39

- Data format
 - ▣ { source, destination }




page1, page3
page2, page1
page3, page1
page3, page4
page4, page1
page4, page2

PageRank in Spark


40

Iterative Algorithm Example

```
val lines = sc.textFile("<path>")
val links = lines.map{ s =>
    val parts = s.split(",")
    (parts(0), parts(1))
}.distinct().groupByKey().cache()
```



page1, page3
page2, page1
page3, page1
page3, page4
page4, page1
page4, page2




(page1, [page3])
(page2, [page1])
(page3, [page1, page4])
(page4, [page2, page1])

PageRank in Spark


41

Iterative Algorithm Example

```
val lines = sc.textFile("<path>")
val links = lines.map( s => {
    val parts = s.split(",")
    (parts(0), parts(1))
}).distinct().groupByKey().cache()
var ranks = links.mapValues(v => 1.0)
```



(page1, [page3])
(page2, [page1])
(page3, [page1, page4])
(page4, [page2, page1])



(page1, 1.0)
(page2, 1.0)
(page3, 1.0)
(page4, 1.0)

PageRank in Spark

42

(page1, [page3])
(page2, [page1])
(page3, [page1, page4])
(page4, [page2, page1])

(page1, 1.0)
(page2, 1.0)
(page3, 1.0)
(page4, 1.0)



(page1, ([page3], 1.0))
(page2, ([page1], 1.0))
(page3, ([page1, page4], 1.0))
(page4, ([page2, page1], 1.0))



(page3, 1.0)
(page1, 1.0)
(page1, 0.5)
(page4, 0.5)
(page2, 0.5)
(page1, 0.5)

```
for (i <- 1 to 10) {  
  val linkWithRank = links.join(ranks)  
  val contribs = linkWithRank.values.flatMap{  
    case (urls, rank) =>  
      val size = urls.size  
      urls.map(url => (url, rank / size))  
  }  
  ranks = contribs.reduceByKey(_ + _)  
    .mapValues(0.15 + 0.85 * _)  
}
```

PageRank in Spark

43

```
for (i <- 1 to 10) {  
  val linkWithRank = links.join(ranks)  
  val contribs = linkWithRank.values.flatMap{  
    case (urls, rank) =>  
      val size = urls.size  
      urls.map(url => (url, rank / size))  
  }  
  ranks = contribs.reduceByKey(_ + _)  
    .mapValues(0.15 + 0.85 * _)  
}
```

(page3, 1.0)
(page1, 1.0)
(page1, 0.5)
(page4, 0.5)
(page2, 0.5)
(page1, 0.5)



(page3, 1.0)
(page1, 2.0)
(page4, 0.5)
(page2, 0.5)

PageRank in Spark

44

```
for (i <- 1 to 10) {  
  val linkWithRank = links.join(ranks)  
  val contribs = linkWithRank.values.flatMap{  
    case (urls, rank) =>  
      val size = urls.size  
      urls.map(url => (url, rank / size))  
  }  
  ranks = contribs.reduceByKey(_ + _)  
    .mapValues(0.15 + 0.85 * _)  
}
```

(page3, 1.0)
(page1, 1.0)
(page1, 0.5)
(page4, 0.5)
(page2, 0.5)
(page1, 0.5)



(page3, 1.0)
(page1, 2.0)
(page4, 0.5)
(page2, 0.5)



(page3, 1.0)
(page1, 1.85)
(page4, 0.58)
(page2, 0.58)

PageRank in Spark

45


Iterative Algorithm Example

```
val lines = sc.textFile("<path>")
val links = lines.map{ s =>
    val parts = s.split("\\s+")
    (parts(0), parts(1))
}.distinct().groupByKey().cache()
var ranks = links.mapValues(v => 1.0) // (url, rank)
for (i <- 1 to 10) {
    val linkWithRank = links.join(ranks)
    val contribs = linkWithRank.values.flatMap{
        case (urls, rank) =>
            val size = urls.size
            urls.map(url => (url, rank / size))
    }
    ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
}
val output = ranks.collect()
output.foreach(rank => println(ranks._1 + " has rank " + ranks._2))
```

PageRank in Spark

46

- Each iteration requires a join - links and ranks
- Each join requires shuffling data over the network



PageRank in Spark

47

```
val lines = sc.textFile("<path>")  
val links = lines.map(...).partitionBy(new HashPartitioner(8))
```

