

# INTRODUCTION TO SPARK WITH SCALA

## Spark Structured Streaming



# Agenda

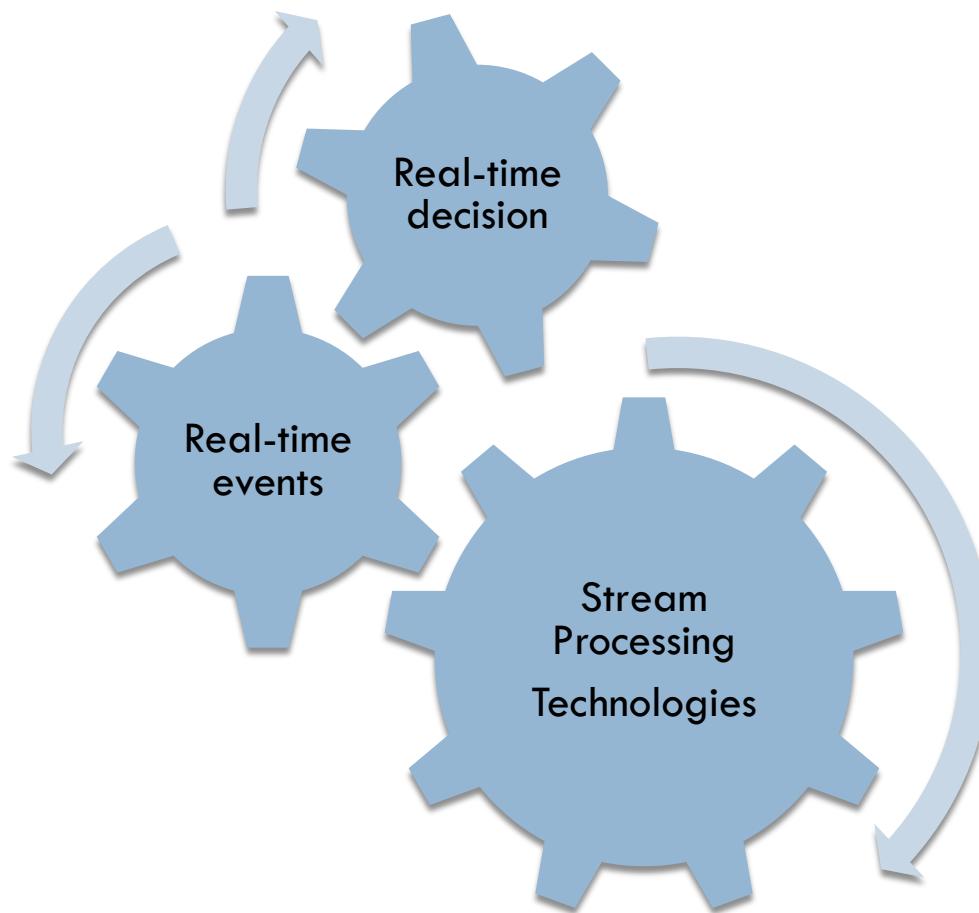
2

- Stream Processing
- Overview of Structured Streaming
- Structured Streaming Concepts
- Structured Streaming API
  - Datasets & DataFrames
- Working with Streaming Queries
- Execution Model
- Fault Tolerance
- Resources

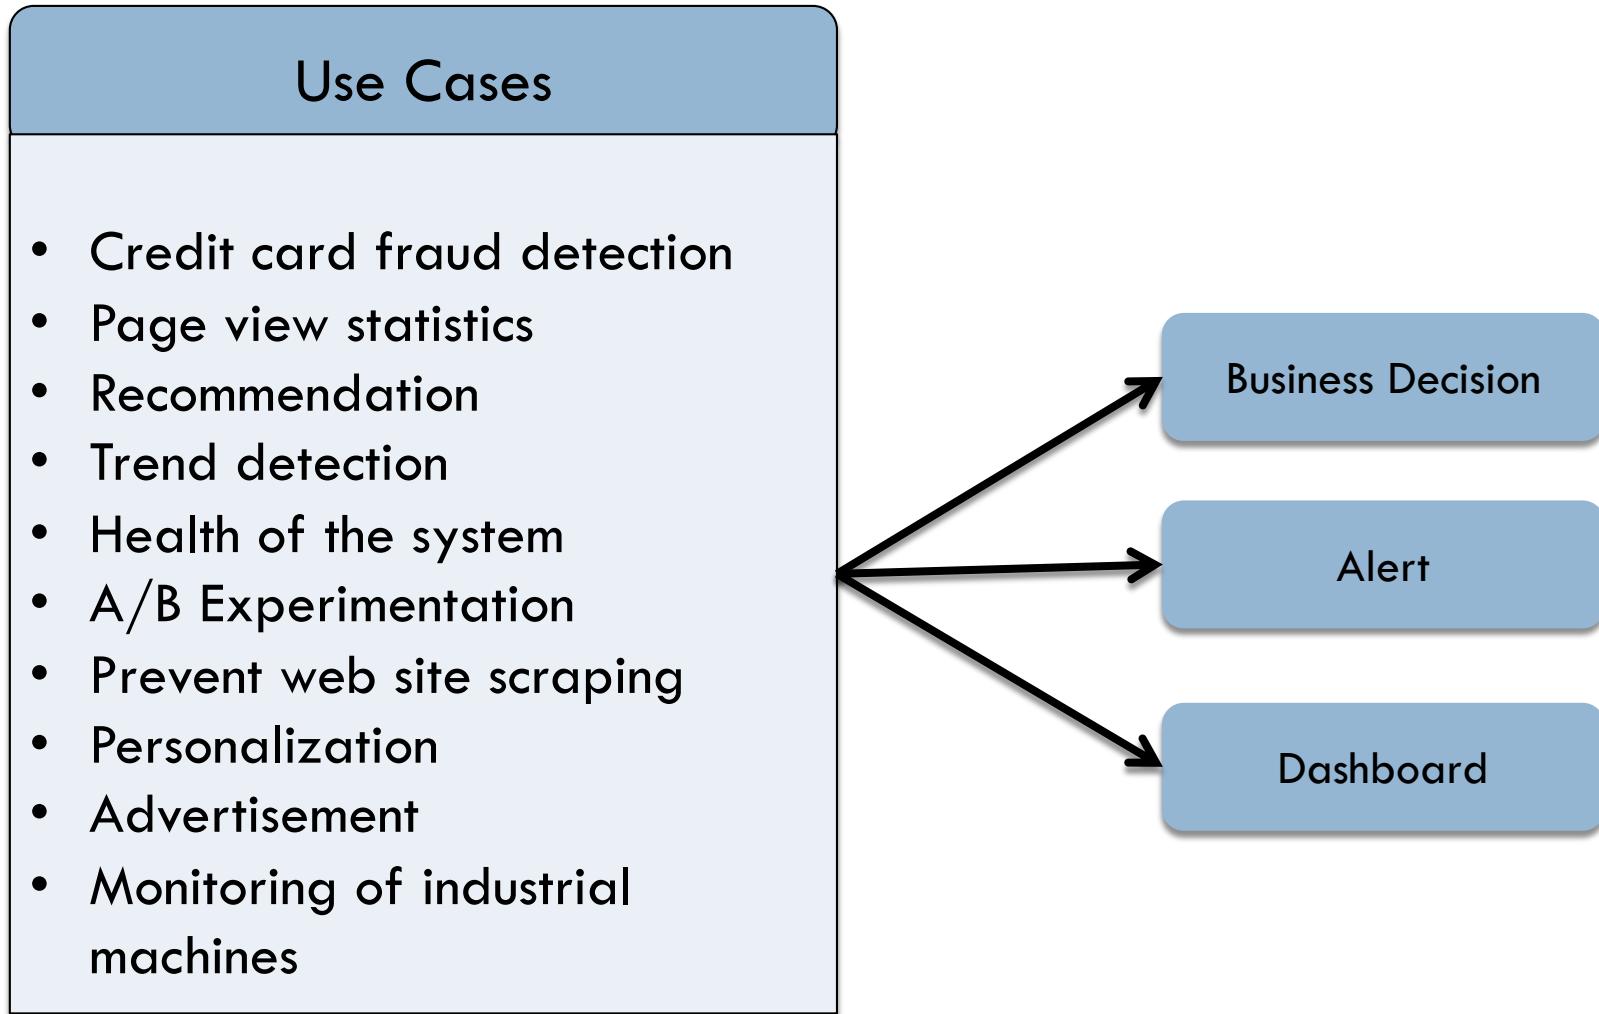
# Stream Processing

3

## The Rise of Real Time



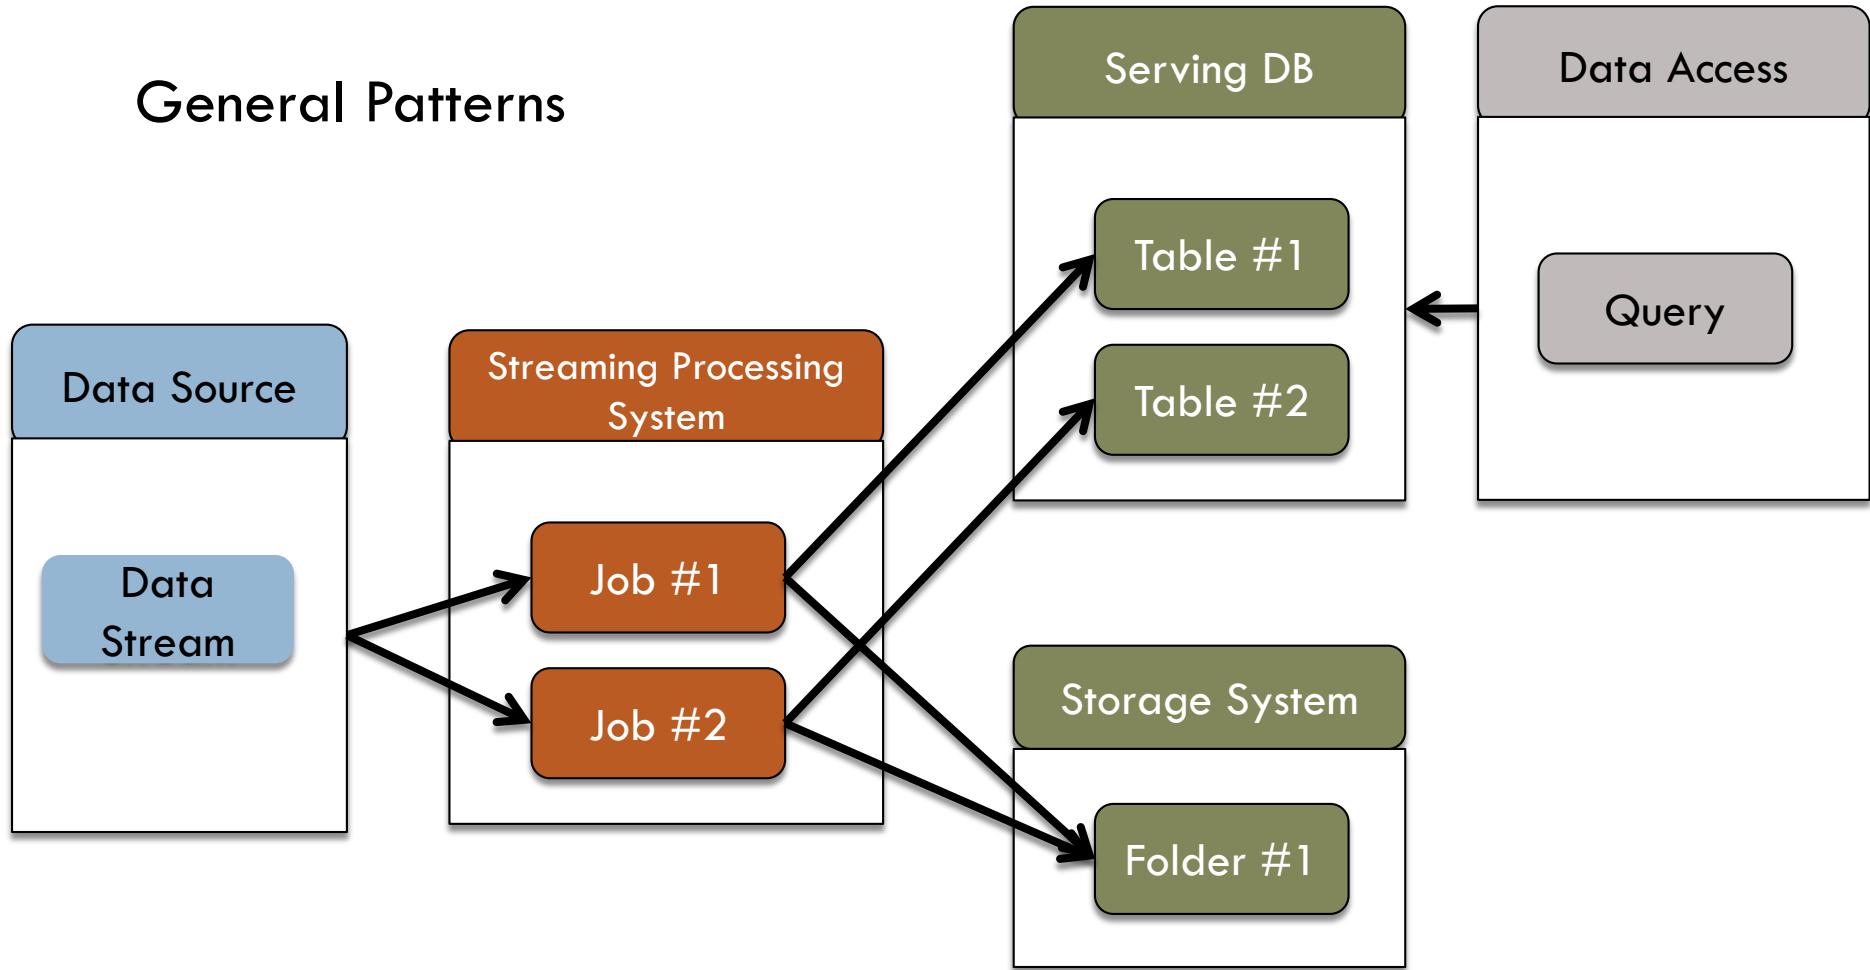
# Stream Processing



# Stream Processing

5

## General Patterns

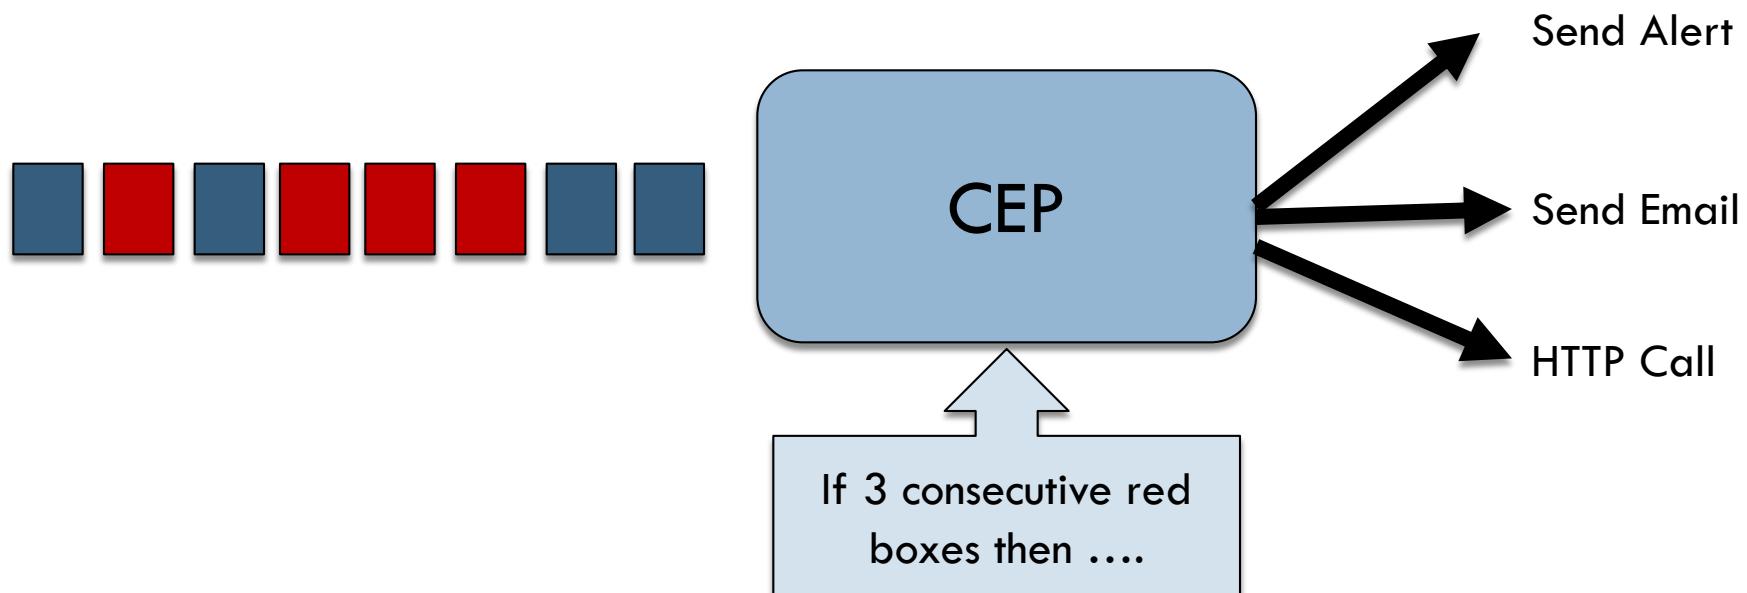


# Stream Processing

6

## □ Complex Event Processing

- Detect complex event patterns in streams of data



# Stream Processing

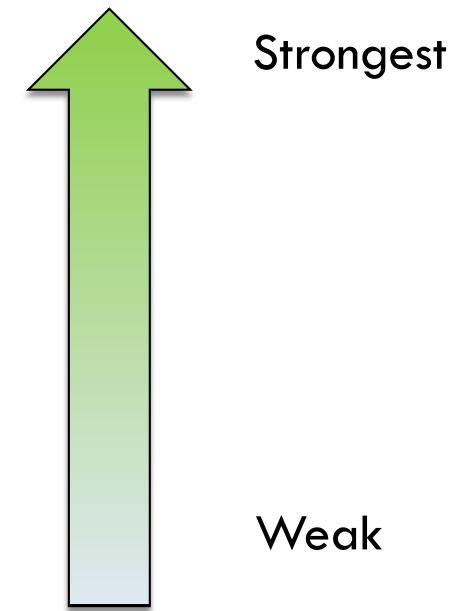
7

## Messaging Semantics

Exactly once => 1 only

At least once => 1 or n

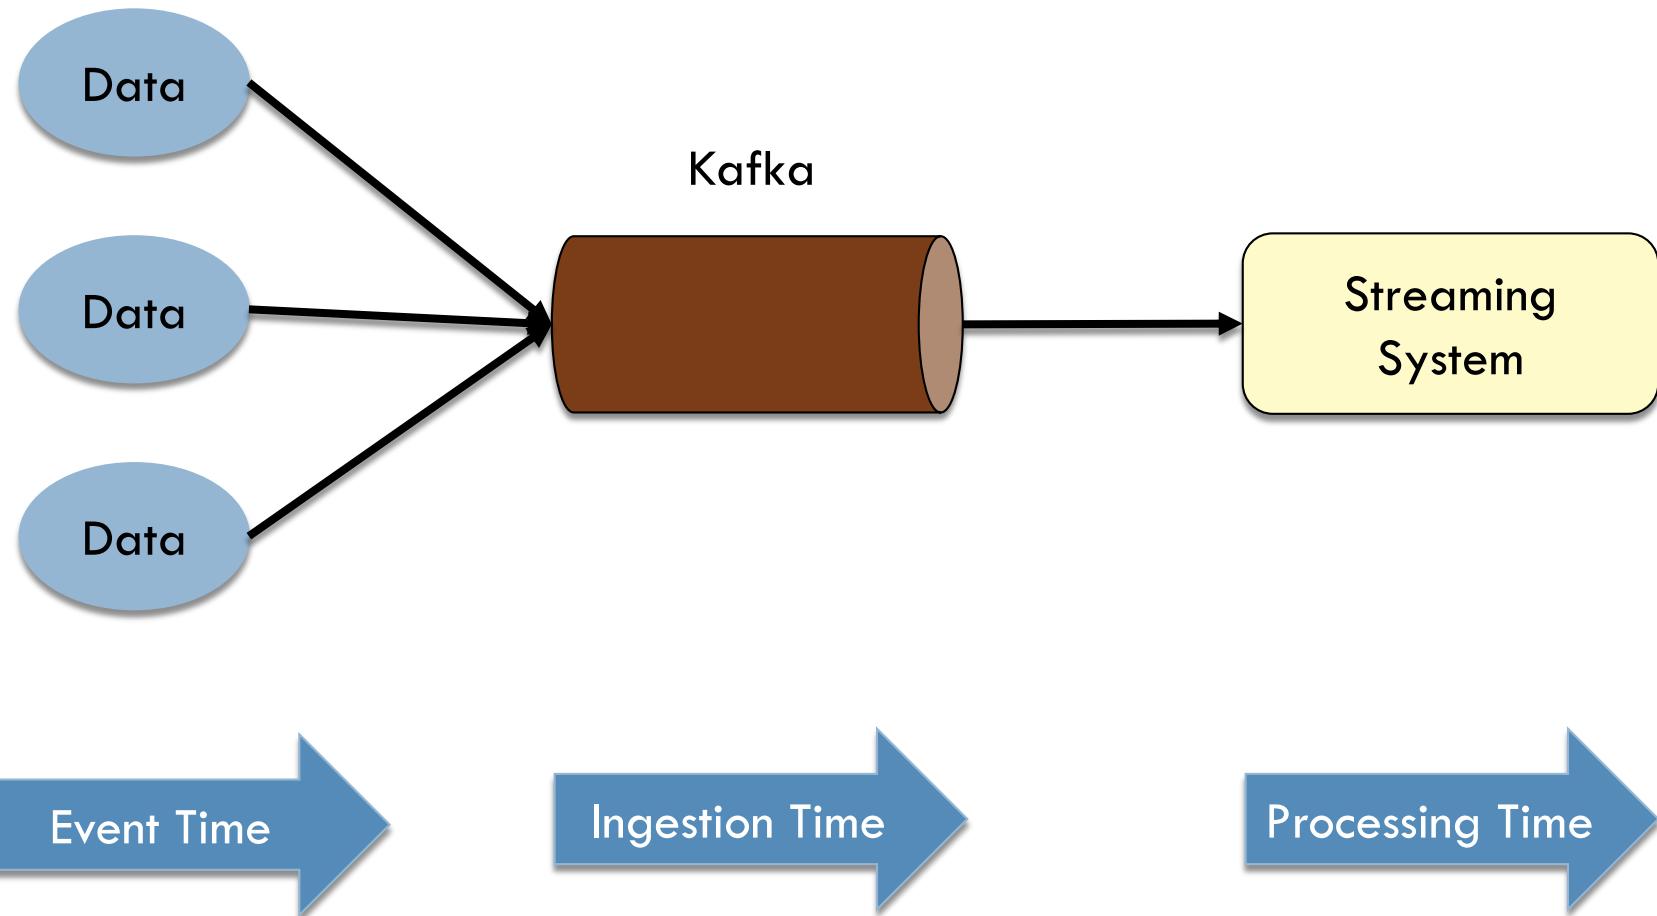
At most once => 1 or 0



# Stream Processing

8

## Notion of Event Time



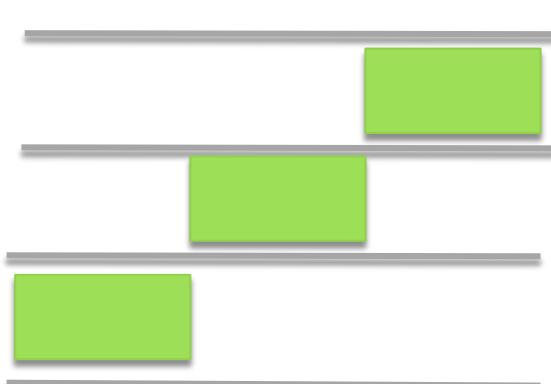
# Stream Processing

9

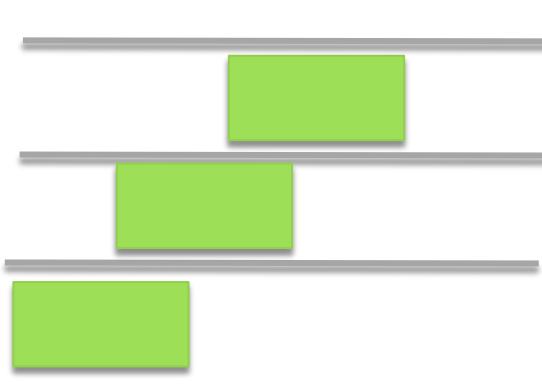
## Windowing

- Slicing up incoming data into chunks for processing
- Delineate finite boundaries for grouping

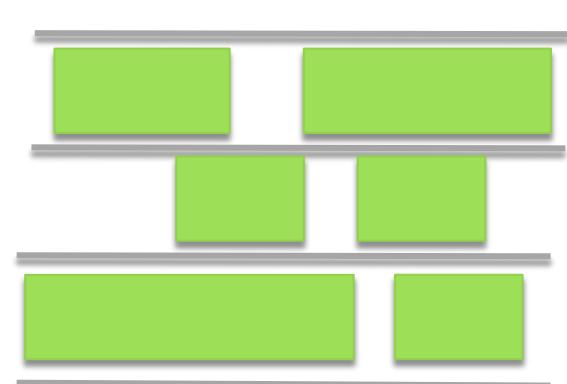
Fixed/Tumbling Window



Sliding Window



Session Window



Time

Time

Time

# Stream Processing

10

## Data Streaming Processing Systems



(Heron)



dataArtisans

confluent



Apache Beam

# Stream Processing

11

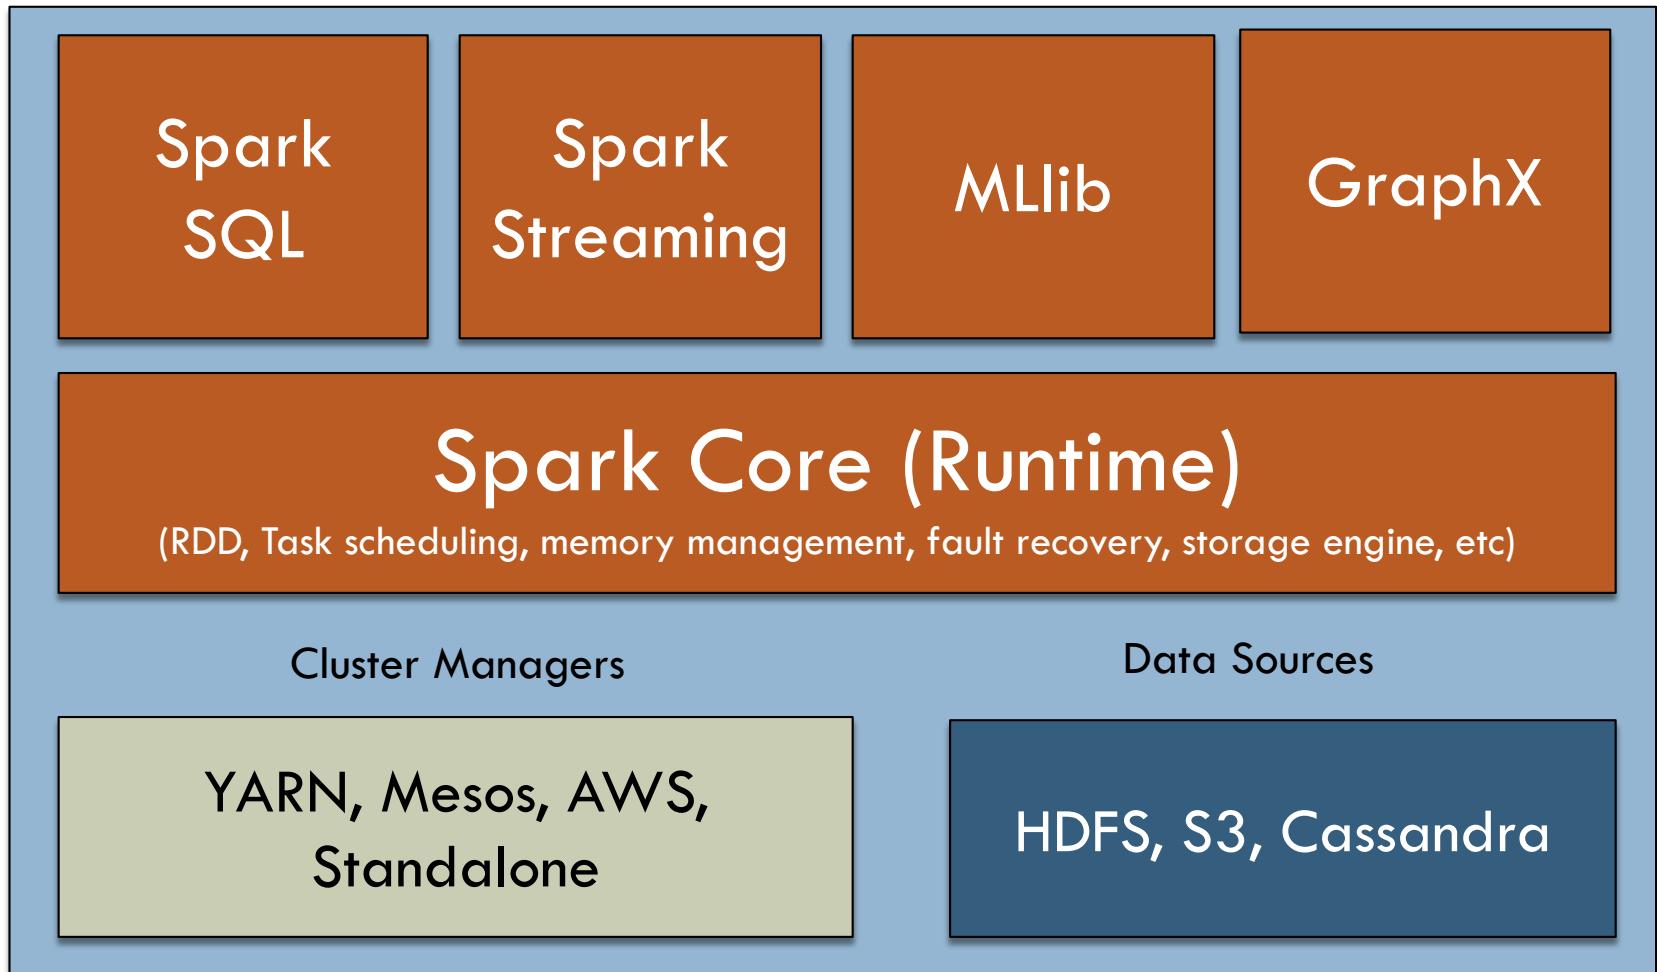
## Requirements for Streaming Processing Engine

- Keep the data moving
  - Streaming architecture
- Declarative access
  - E.g. StreamSQL, CQL
- Handle imperfections
  - Late, missing, unordered items
- Predictable outcomes
  - Consistency, event time
- Integrate stored and streaming data
  - Hybrid stream and batch
- Data safety and availability
  - Fault tolerance, durable state
- Automatic partitioning and scaling
  - Distributed processing
- Instantaneous processing and response

# Spark Streaming Concepts

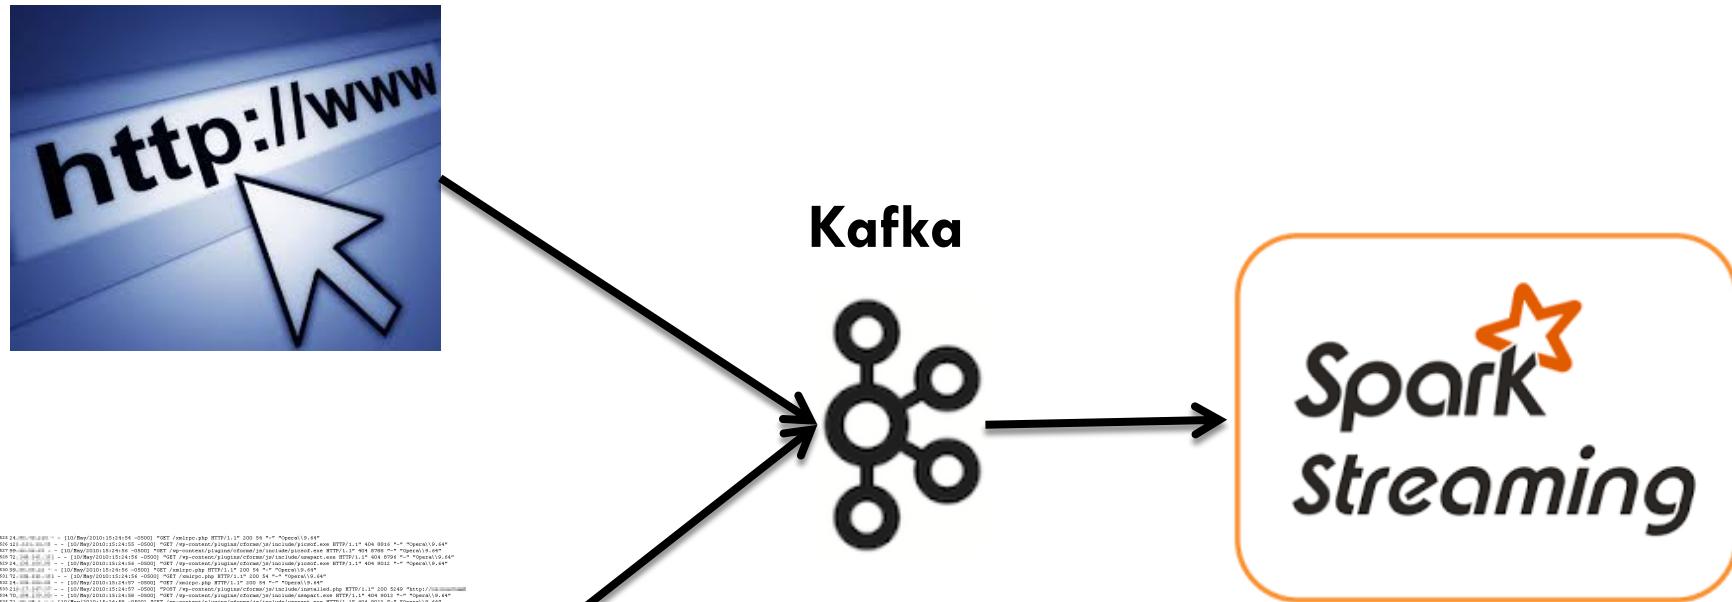
12

## Apache Spark Unified Stack



# Spark Streaming Concepts

13



```
44842 24 - [0.0] GET / HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44843 99 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44844 99 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44845 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44846 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44847 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44848 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44849 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44850 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44851 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44852 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44853 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44854 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44855 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44856 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44857 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44858 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44859 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44860 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44861 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44862 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44863 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44864 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44865 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44866 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44867 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44868 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44869 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44870 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44871 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44872 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44873 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44874 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44875 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44876 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44877 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44878 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44879 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44880 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44881 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44882 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44883 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44884 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44885 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44886 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44887 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44888 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44889 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44890 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44891 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44892 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44893 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44894 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44895 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44896 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44897 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44898 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44899 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44900 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"  
44901 24 - [0.0] GET /index.php HTTP/1.1 200 54 "-" "OpenSUSE/1.44"
```

Event data

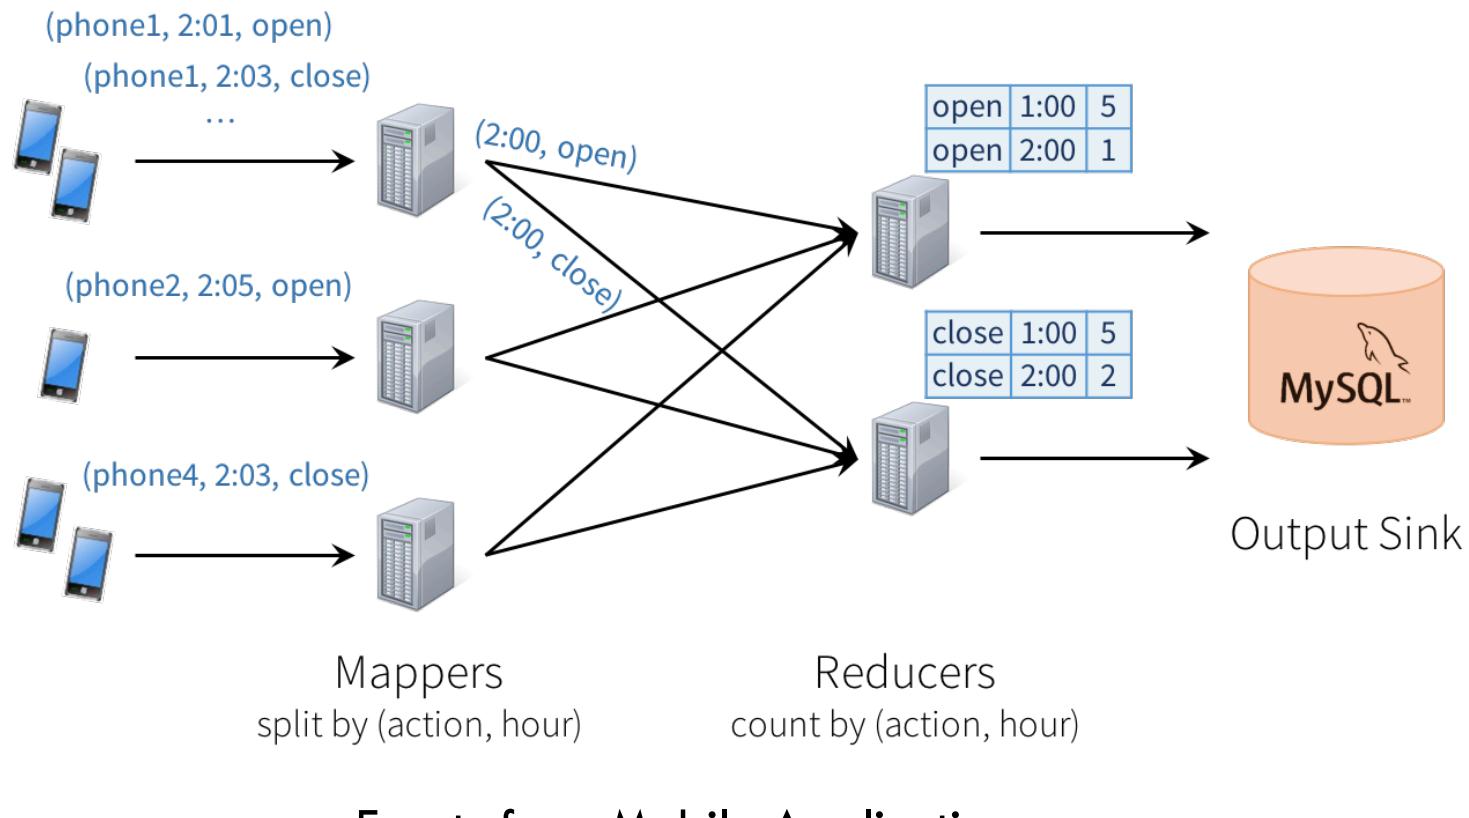
Buffering

Processing

# Overview of Structured Streaming

14

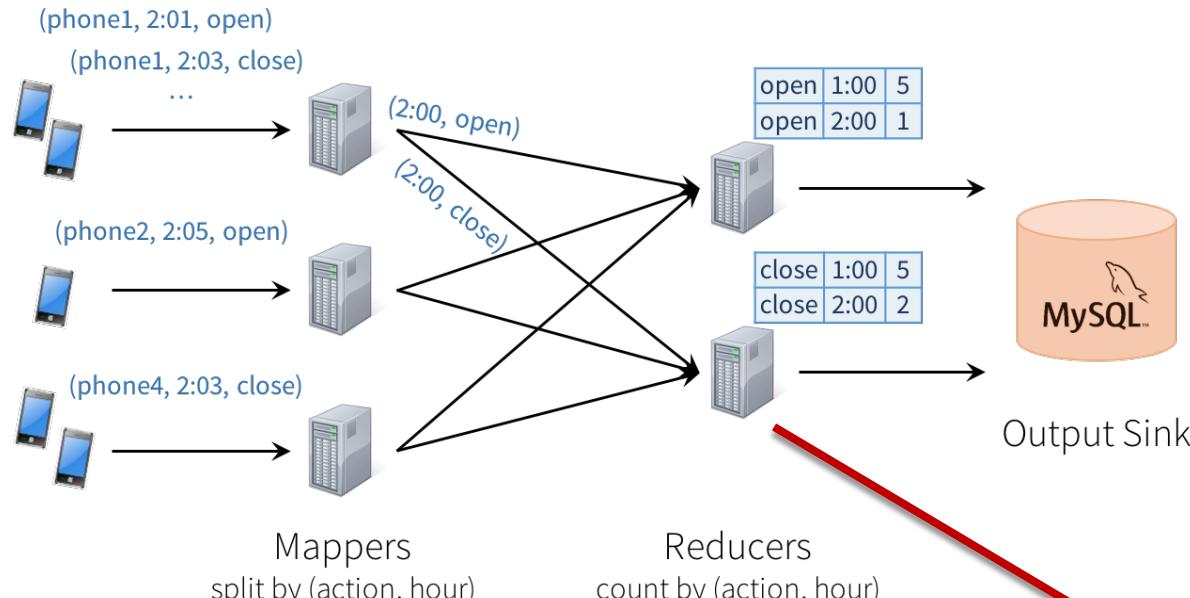
## Real-time Aggregation of Actions Per Hour



# Overview of Structured Streaming

15

## Real-time Aggregation of Actions/Hour



### Potential Issues:

- Consistency
- Fault tolerance
- Out-of-order data

Slow or  
failed

# Overview of Structured Streaming

16

- Streaming applications are complex
  - Integrating with other systems
    - Interactive, batch, RDBMS, machine learning
  - Dealing with late event or out-of-order data
  - Fault tolerance & consistency
- DStream
  - No event time processing support
  - Limited interoperability between batch and interactive
  - Difficult to reasoning about end-to-end guarantees

# Overview of Structured Streaming

17

*The simplest way to perform streaming analytics is  
not having to **reason** about streaming*

- Reynold Xin, chief architect @ Databricks

# Overview of Structured Streaming

18

*Structured Streaming provides fast, scalable, fault-tolerant, end-to-end exactly once streaming processing without user having to reason about streaming*

# Overview of Structured Streaming

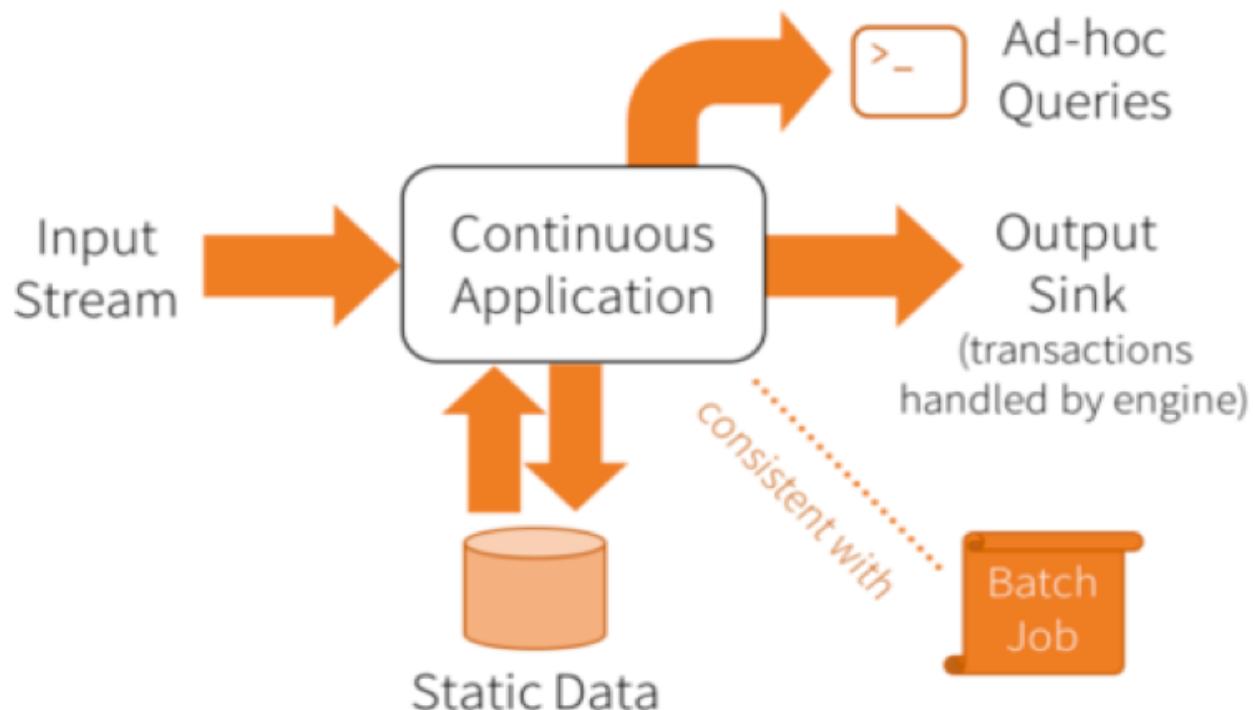
19

- High-level API built on top of Spark SQL engine
  - Streaming computation similar to batch computation
- Strong guarantees about consistency with batch jobs
- Ensure end-to-end exactly-once guarantees
  - Transactional integration with storage system
- Handling late or out-of-order data
- Tight integration with the rest of Spark
  - Serving interactive queries on streaming state
  - Integrate with MLlib

# Overview of Structured Streaming

20

*Continuous application – an end-to-end application that reacts to data in real-time*



# Structured Streaming Concepts

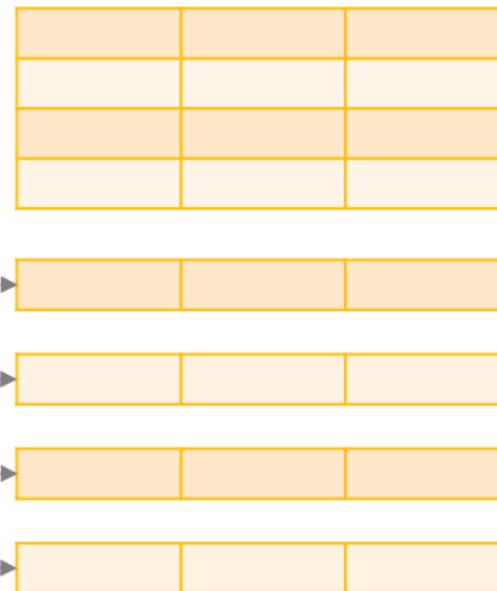
21

## Data Stream as An Unbounded Table

Data stream



Unbounded Table



Developers specify:

- Query
- Trigger point

new data in the  
data stream

=

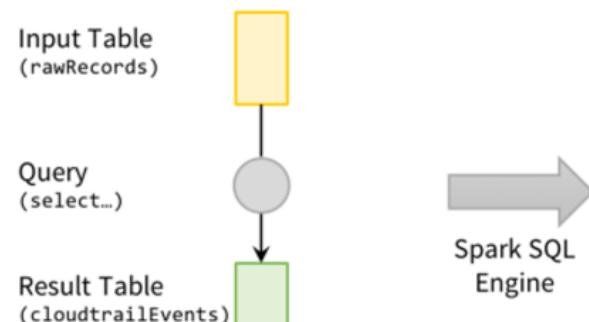
new rows appended  
to a unbounded table

# Structured Streaming Concepts

22

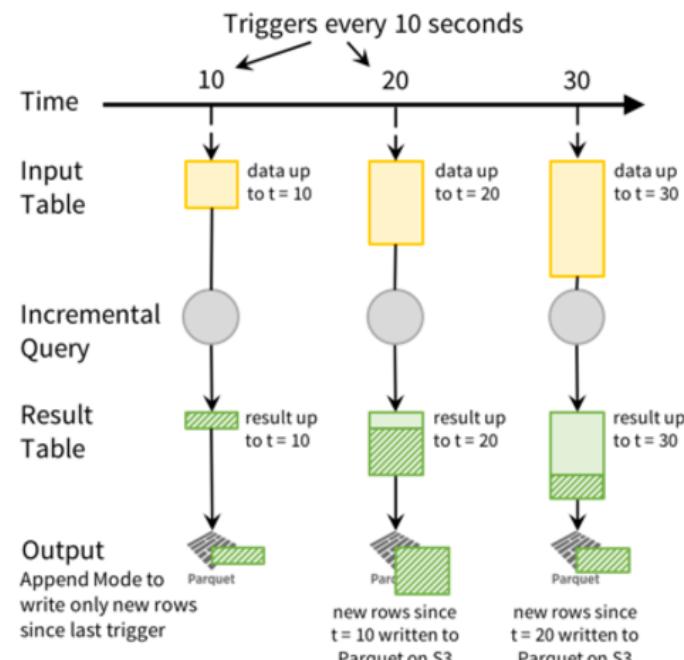
## Structured Streaming Processing Model

### Batch Processing



user's query on Input Table using batch API

### Stream Processing

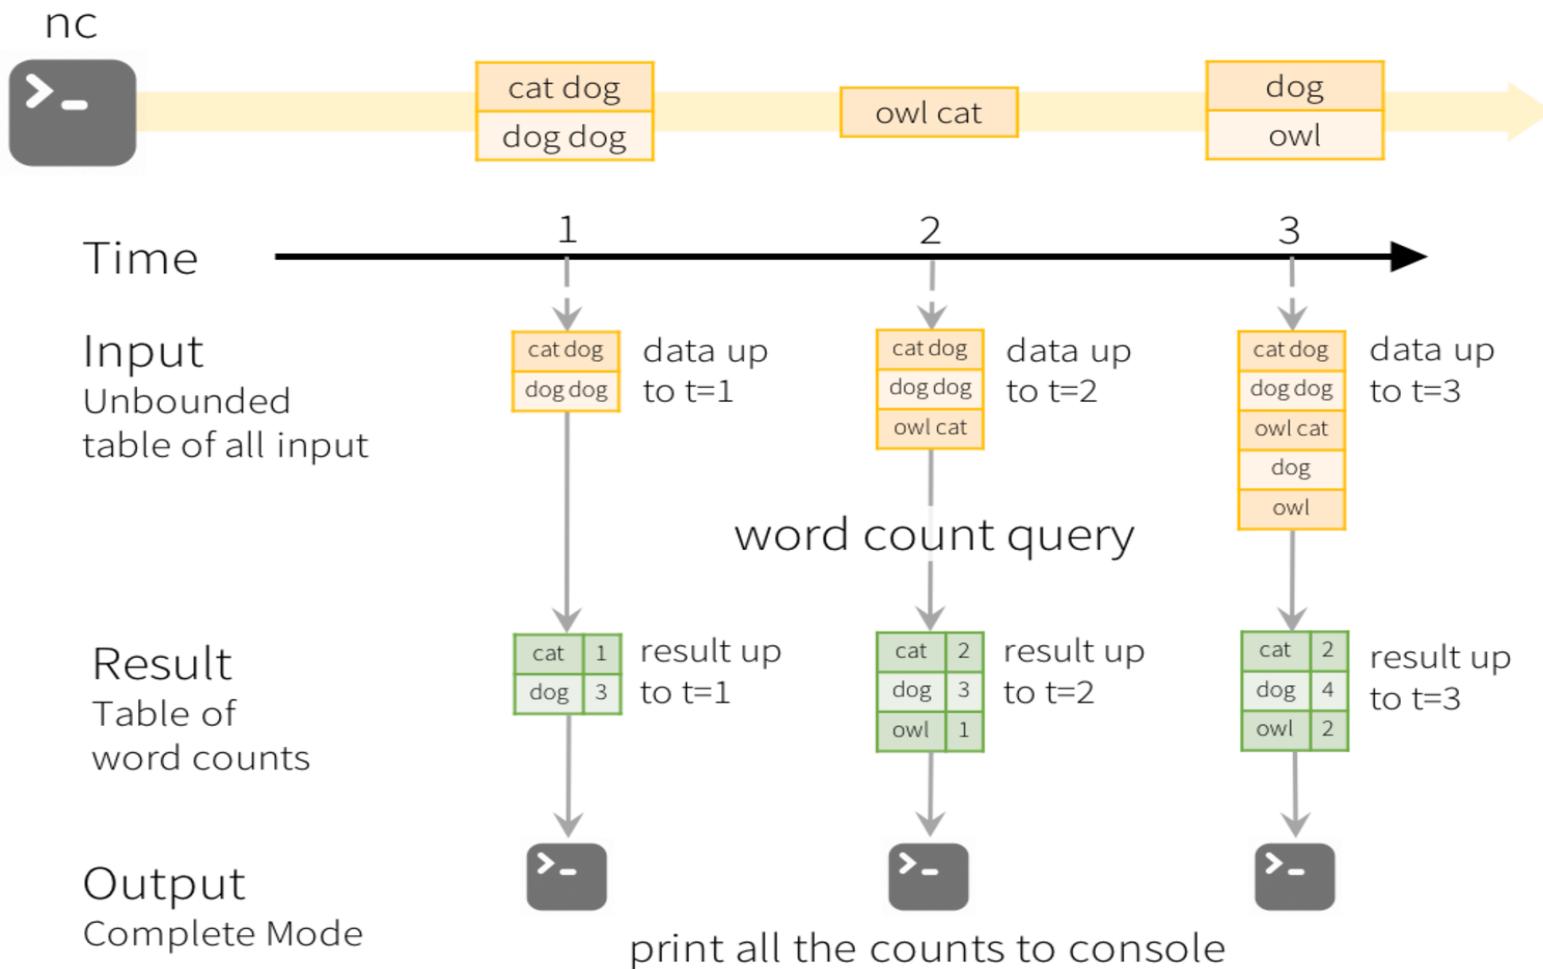


incremental execution on streaming data

# Structured Streaming Concepts

23

## Structured Streaming Word Count Example

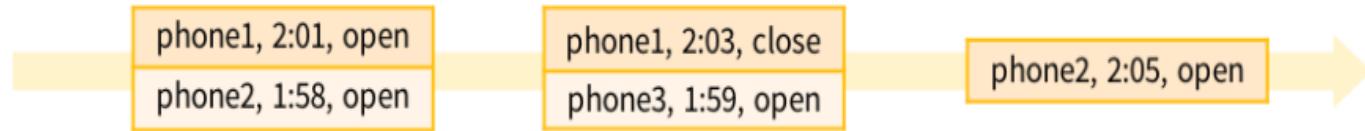


# Structured Streaming Concepts

24

## Events from Mobile Application

Arriving  
Records



```
// schema (phone id, time, action)

val inputDF = spark.readStream.json("<path>")

val countDF = inputDF.groupBy($"action",
    window($"time", "1 hour")).count()
```

# Structured Streaming Concepts

25

Arriving  
Records

phone1, 2:01, open
phone2, 1:58, open

phone1, 2:03, close
phone3, 1:59, open

phone1, 2:05, open
--------------------

System time

2:02

2:04

2:06

Input table

phone1, 2:01, open
phone2, 1:58, open

phone1, 2:01, open
phone2, 1:58, open
phone1, 2:03, close
phone3, 1:59, open

phone1, 2:01, open
phone2, 1:58, open
phone1, 2:03, close
phone3, 1:59, open
phone1, 2:05, open

Query

open	2:00	1
open	1:00	1

open	2:00	1
open	1:00	2
close	2:00	1

open	2:00	2
open	1:00	2
close	2:00	1

Result table

open	2:00	1
open	1:00	1

open	1:00	2
close	2:00	1

open	2:00	2
------	------	---

Output  
(update)

# Structured Streaming Concepts

26

- Fault Recovery & Storage System Requirements
  - Results must still be valid even if machine failed
  - Input sources must be replayable
    - Amazon Kinesis, Apache Kafka, File system
  - Output sinks must support transactional updates
    - Idempotent
  - Spark will manage its internal storage
    - On a reliable storage system
    - S3 or HDFS

# Structured Streaming Concepts

27

## Output Modes

Mode	Description
Complete	The entire updated result table will be written to external storage. (Monitoring)
Append	Only the new rows appended in the result table since the last trigger. Rows will be output only once. (ETL)
Update	Only updated rows in the result table since the last trigger.. Applicable for output sinks that can be updated in place, such as RDBMS table. (Alerting)

# Structured Streaming API

28

API built on top of Spark SQL engine

## Batch ETL

```
val input = spark.read.json("<path>")

val result = input.select("phoneId", "action").where($"action" === "open")

result.write.format("parquet").save("<path>")
```

## Streaming

```
val input = spark.readStream.json("<path>")

val result = input.select("phoneId", "action").where($"action" === "open")

result.writeStream.format("parquet").option("path", "<path>")
```

# Structured Streaming API

29

## Word Count Example

```
import org.apache.spark.sql.functions._  
import org.apache.spark.sql.SparkSession  
import spark.implicits._  
  
val lines = spark.readStream.format("socket")  
    .option("host", "localhost").option("port", 9999).load()  
  
// Split the lines into words  
val words = lines.as[String].flatMap(_.split(" "))  
  
// Generate running word count  
val wordCounts = words.groupBy("value").count()  
  
val query = wordCounts.writeStream.outputMode("complete")  
    .format("console").start()
```

# Structured Streaming API

30

- Data Sources
  - File – txt, json, csv, parquet
  - Kafka
  - Socket source – for testing purpose
- Output Sinks
  - File – store output to a directory
  - Console – for debugging
    - Print output every time there is a trigger
  - Memory – for debugging
    - In memory table for low volume data

# Structured Streaming API

31

```
// socket
val dataDF = spark.readStream.format("socket").option("host",
"localhost").option("port", 8888).load()

val dataSchema = new StructType().add("id", StringType).add("time",
TimestampType).add("action", StringType)

// file
val dataDF = spark.readStream.schema(dataSchema).json("<path>")

val dataDF = spark.readStream.schema(dataSchema )
    .option("maxFilesPerTrigger", 1).json("<path>")

// is DF streaming?
dataDF.isStreaming

// File sink
writeStream.format("parquet").start()
// Console sink
writeStream.format("console").start()
```

# Structured Streaming API

32

- Selection, Project, Aggregation
  - Most common operations are supported for streaming
  - `select`, `map`, `filter`, `flatMap`
- Window Operations on Event Time

```
// schema {timestamp:Timestamp, word:String}
val wordsDF = ..

// aggregation over sliding window
val windowCount = dataDF.groupBy(
    window($"timestamp", "10 minutes", "5 minutes"),
    $"word").count()
```

# Structured Streaming API

33

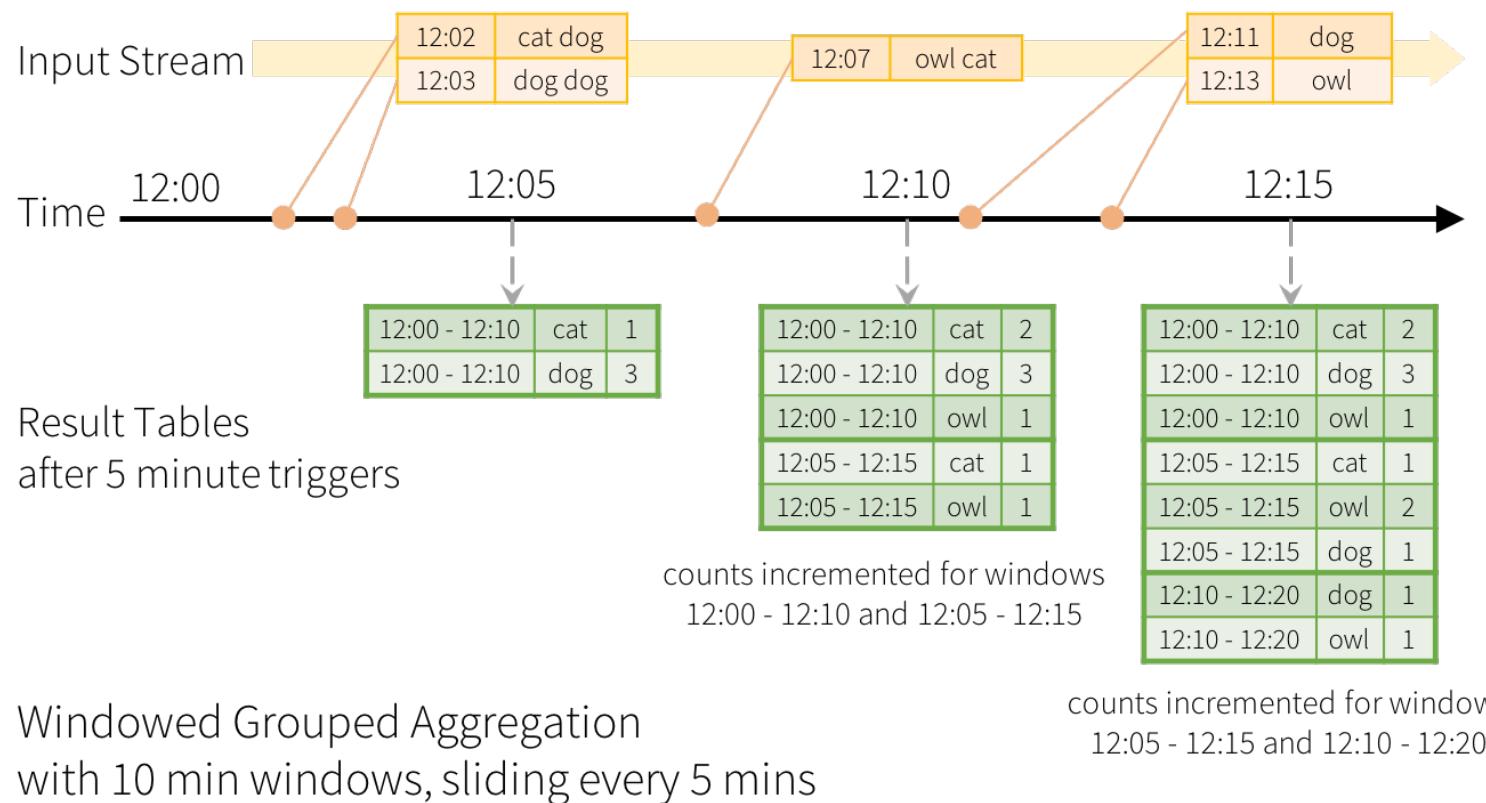
- Joining streams with static data
  - Very easy and without having to think

```
val pageViewEvents = spark.readStream.format("kafka").load()  
  
val userInfo = spark.read.parquet("<path>")  
  
val pageViewWithInfo = pageViewEvents.join(userInfo, "userId")
```

# Structured Streaming API

34

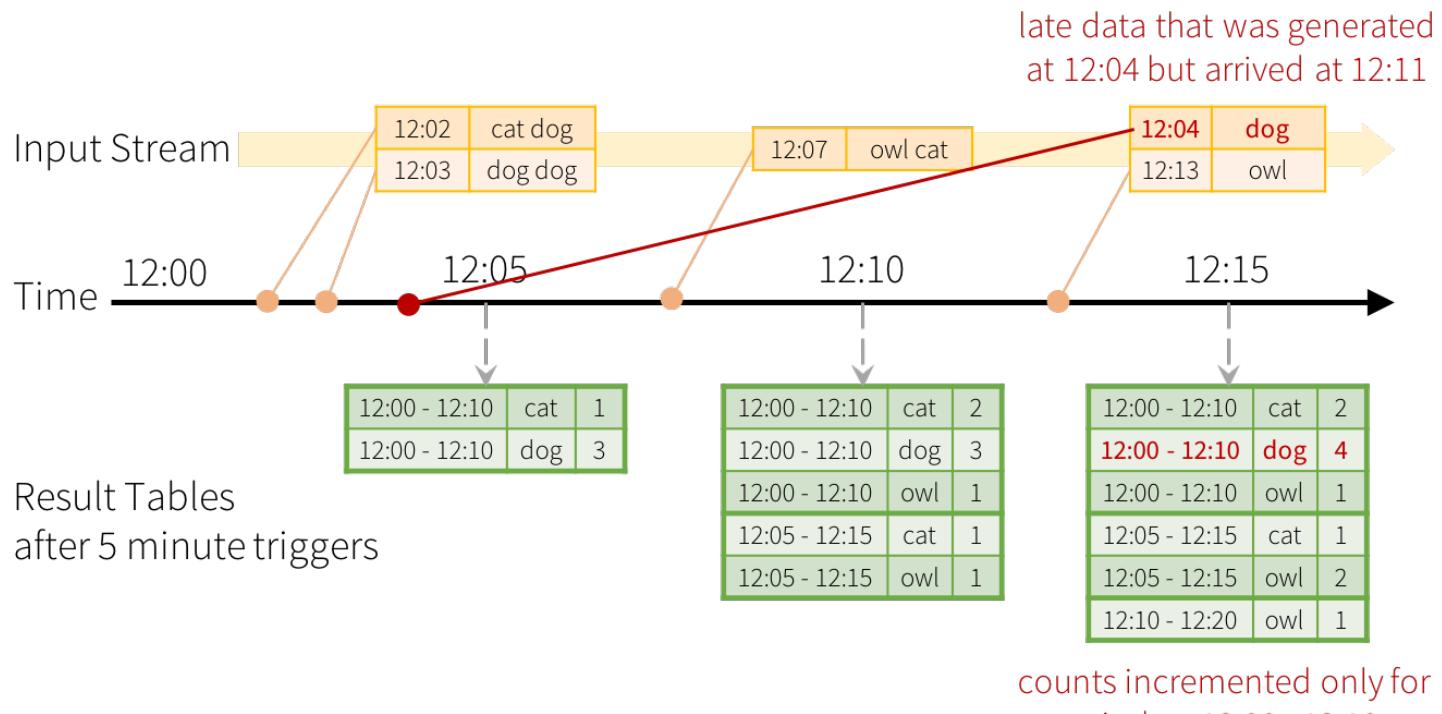
## Aggregation Over Sliding Window



# Structured Streaming API

35

## Handling Late Data & Watermarking



Late data handling in  
Windowed Grouped Aggregation

# Structured Streaming API

36

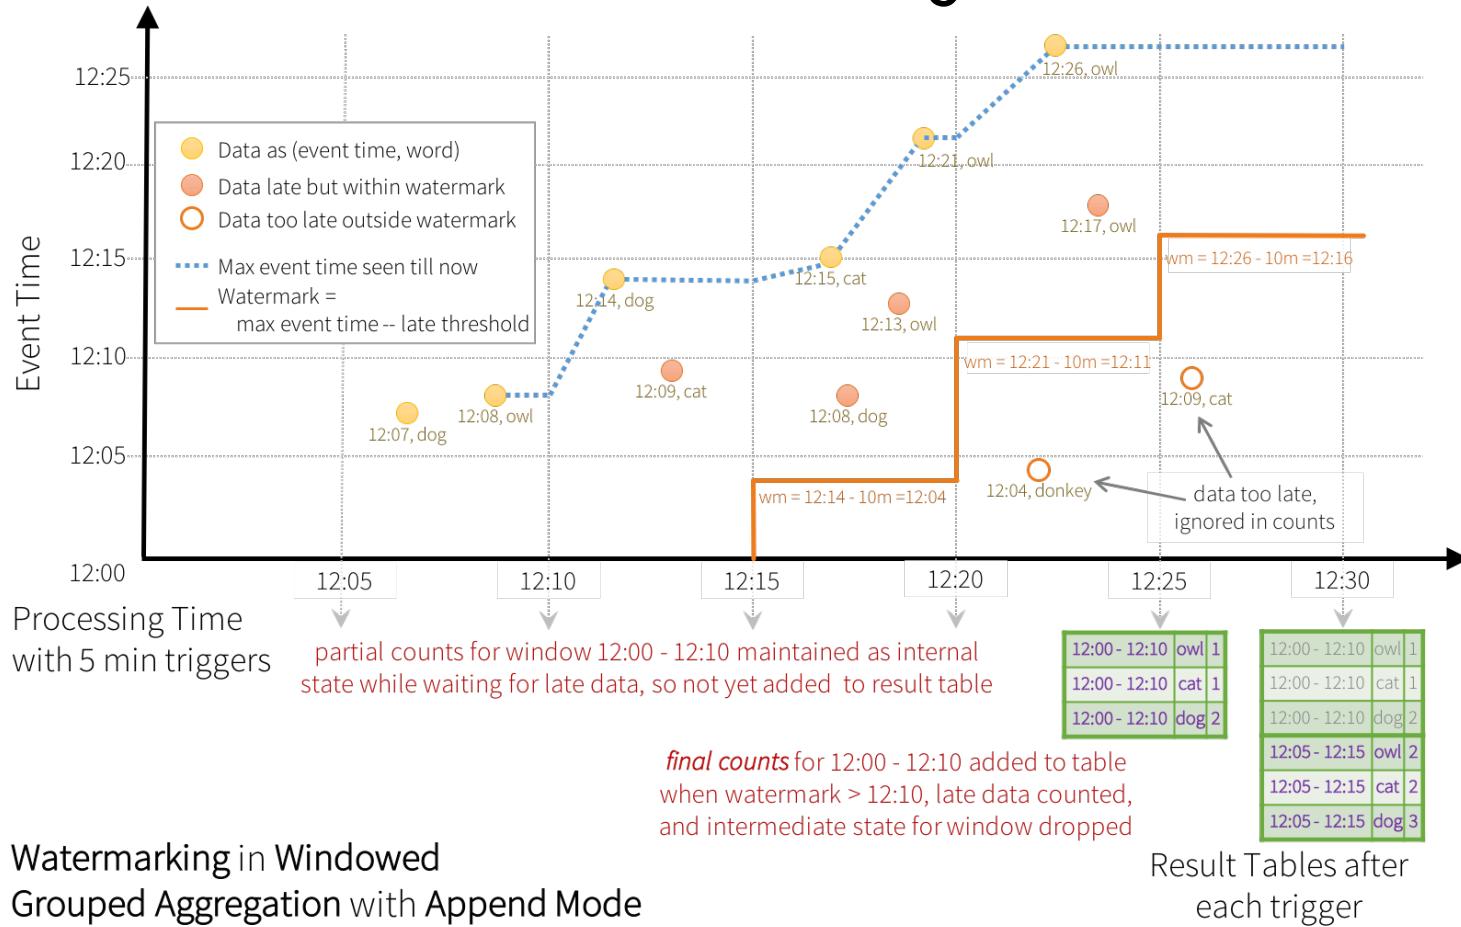
## Watermarking

- A moving threshold trailing behind the max seen event time
- A gap defines how late the data is expected to be
- Data older than the watermark will be ignored
- State older than watermark is removed to reduce intermediate state size
- Of course useful only for stateful operation

# Structured Streaming API

37

## Watermarking



Watermarking in Windowed  
Grouped Aggregation with Append Mode

Result Tables after  
each trigger

# Structured Streaming API

38

- Watermarking example
  - Add call to watermark API before calling groupBy
  - Watermark event-time column must be the same as the one groupBy

```
// schema {timestamp:Timestamp, action:String}  
val actionLogDF = ..  
  
// aggregation over sliding window  
val countDF = actionLogDF  
    .watermark("timestamp", "15 minutes")  
    .groupBy(  
        window($"timestamp", "10 minutes", "5 minutes"),  
        $"word").count()
```

# Structured Streaming API

39

## □ Arbitrary Stateful Processing

- User defined processing logic on each group of data
- Call once per group in batch
- `mapGroupsWithState`
- `flatMapGroupsWithState`
- Support in Scala and Java only
- Need to set time out in order to reduce amount of state to maintain

# Working with Streaming Queries

40

- Expose the result table
  - To interactive queries via Spark's JDBC server
  - Spark 2.0 – just rudimentary “memory” output sink

```
// schema {timestamp:Timestamp, action:String}
val actionLogDF = ...

// aggregation over sliding window
val countDF = actionLogDF.groupBy(
    window($"timestamp", "10 minutes", "5 minutes"),
    $"word").count()

countDF.writeStream.format("memory").queryName("count")
    .outputMode("complete").start()

sql("select sum(count) from count where action='login'")
```

# Working with Streaming Queries

41

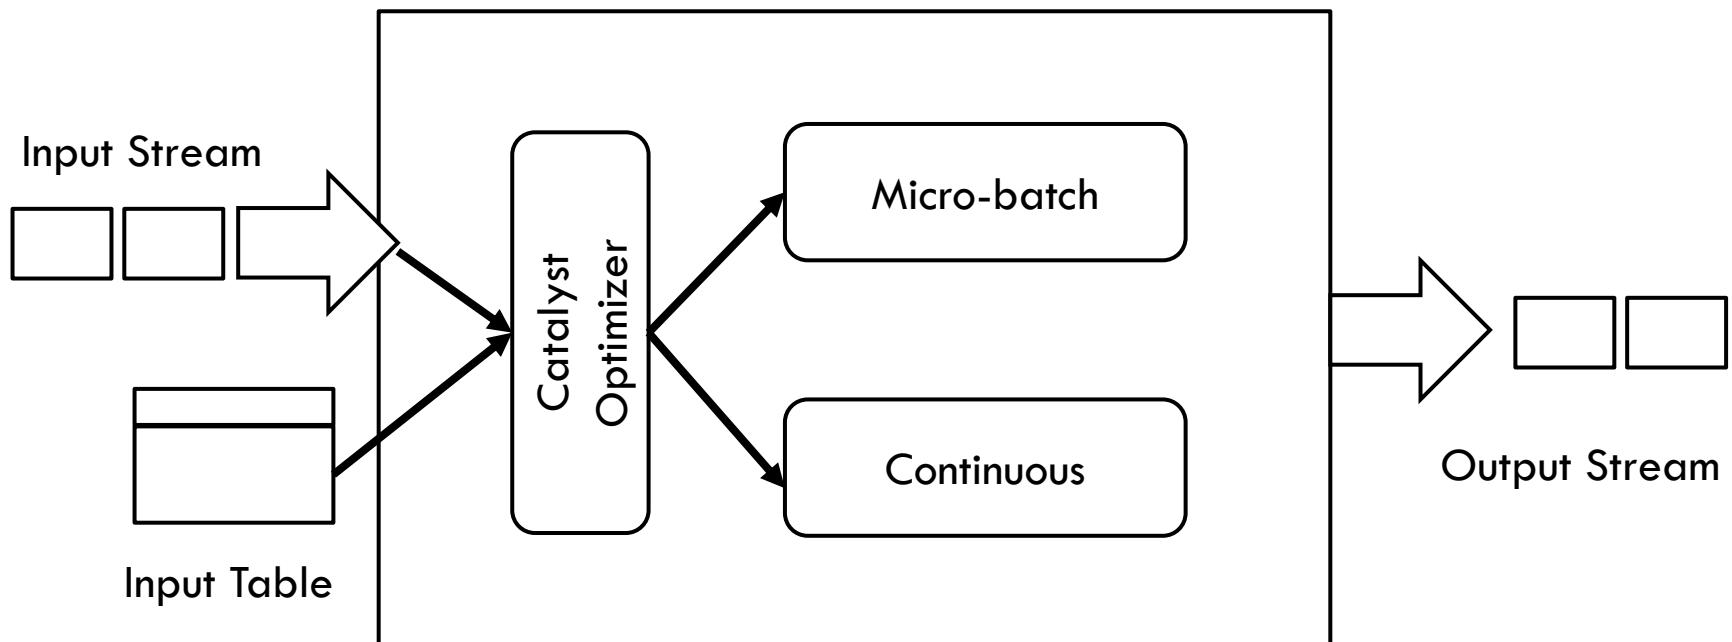
## Triggers

Type	Description
unspecified (default)	Micro-batch mode. Trigger as soon as the next batch has completed
Fixed-interval	Micro-batch mode. Trigger based on the specified interval. For example – every 5 seconds.
One-time	Trigger only once, and then stop. Good for low volume data stream that need to process a few times a day
Continuous	(Experimental 2.3) Continuous processing of new incoming data. Low latency.

# Working with Streaming Queries

42

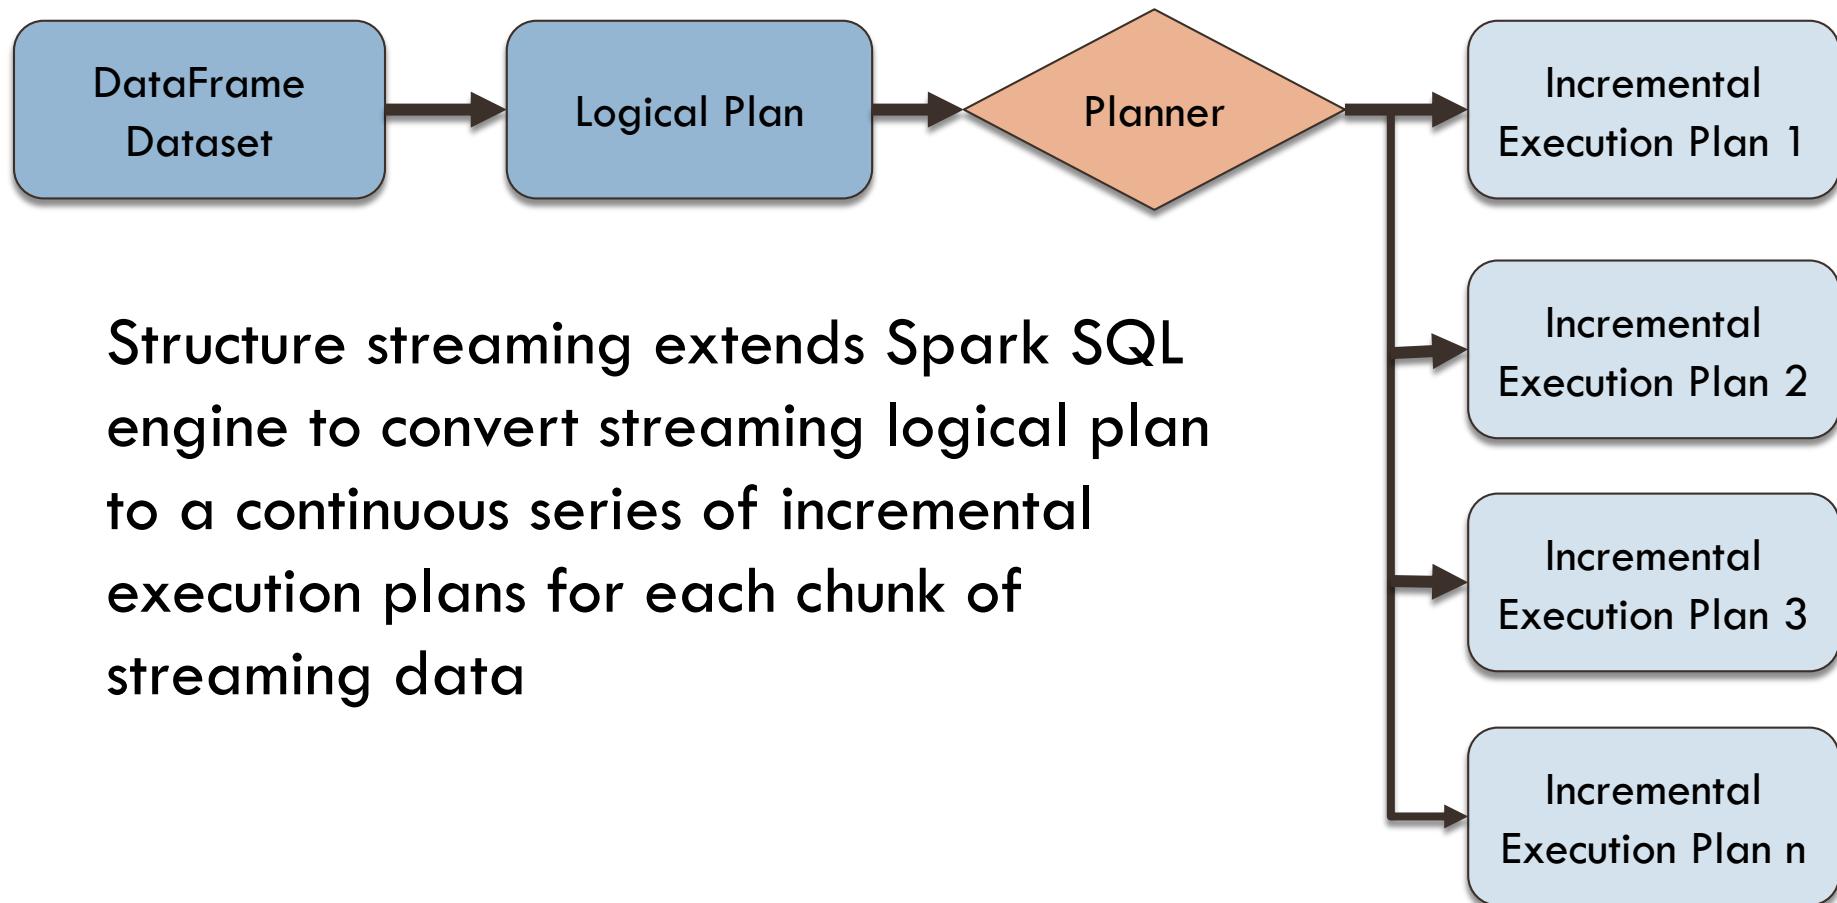
- Continuous processing mode (Spark 2.3)
  - For low latency use cases - < 100ms
  - Real-time fraud detection
  - Long running tasks that process events as they come in



# Execution Model

43

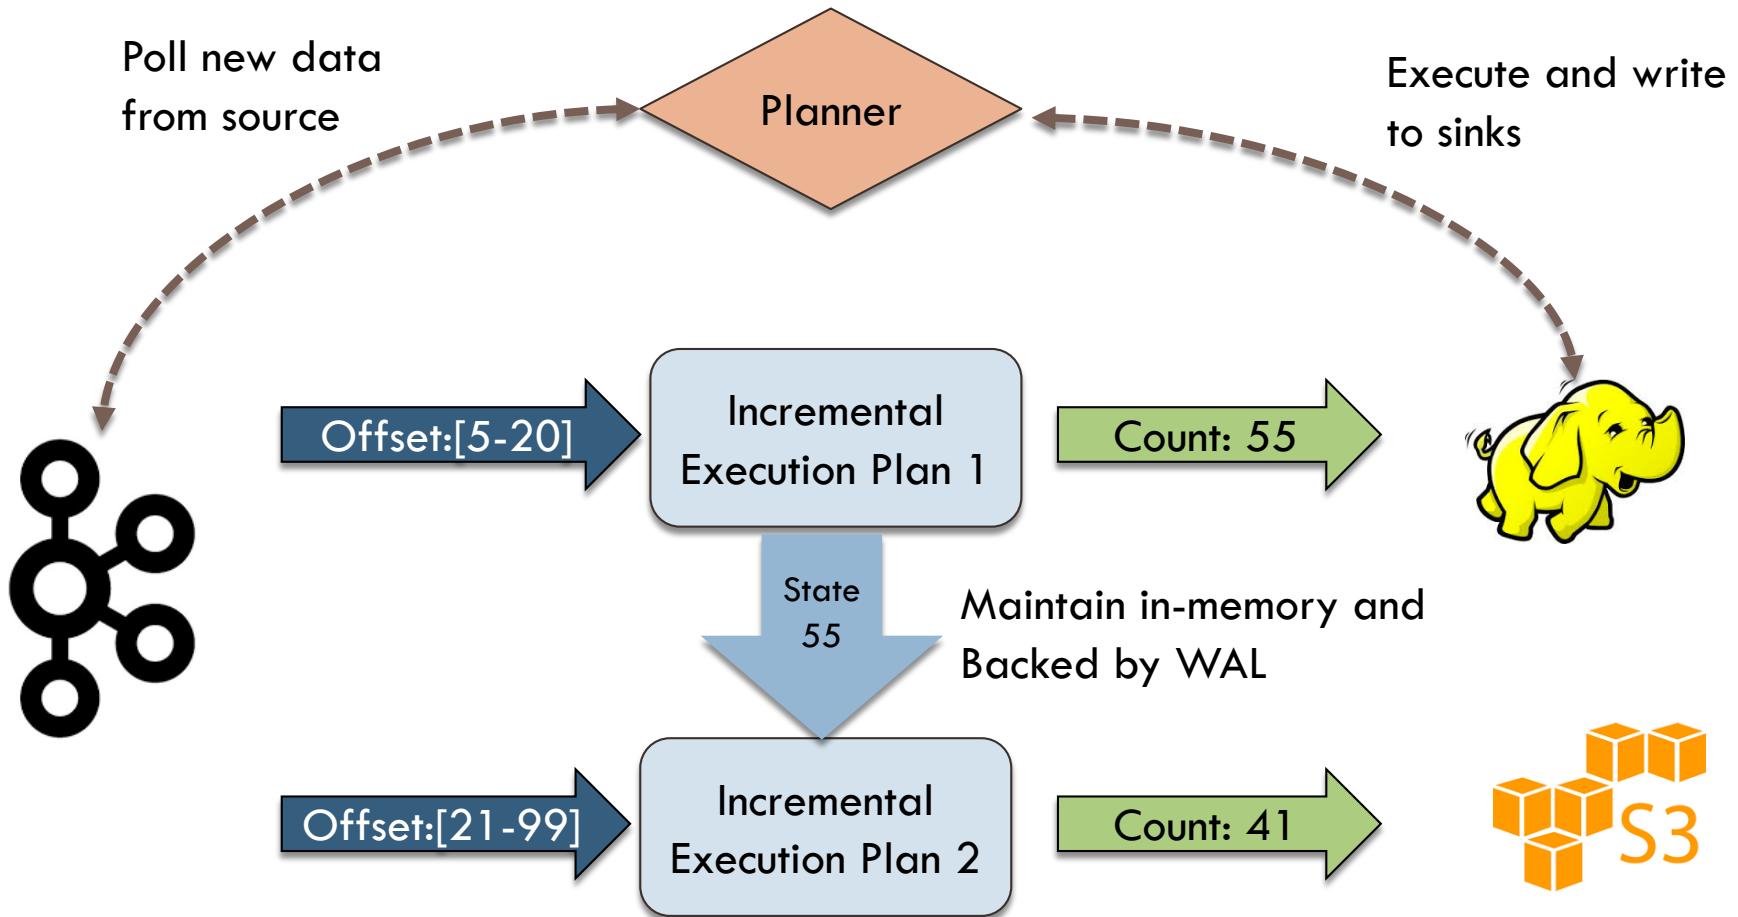
## Continuous & Incremental Execution



# Execution Model

44

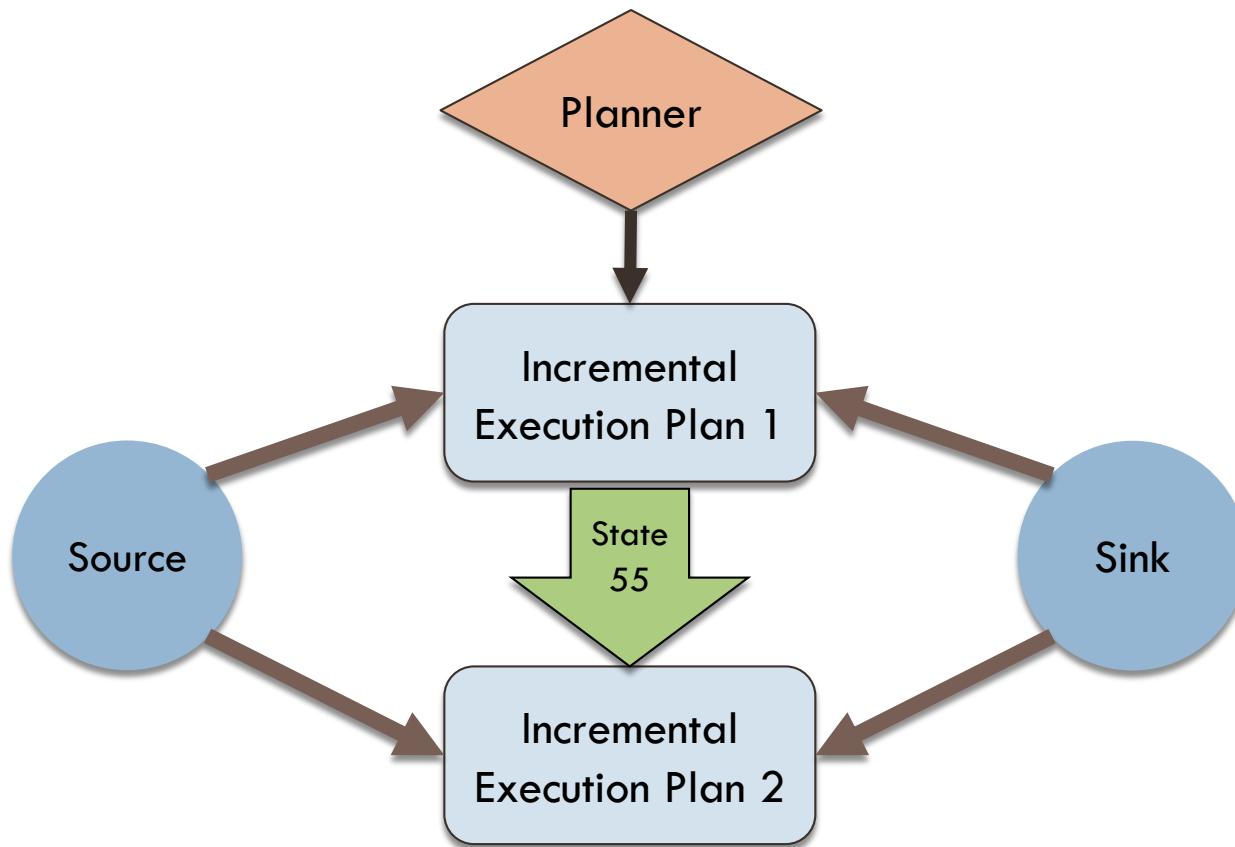
## Continuous & Incremental Execution



# Fault Tolerance

45

## End-to-End Exactly Once Guarantees

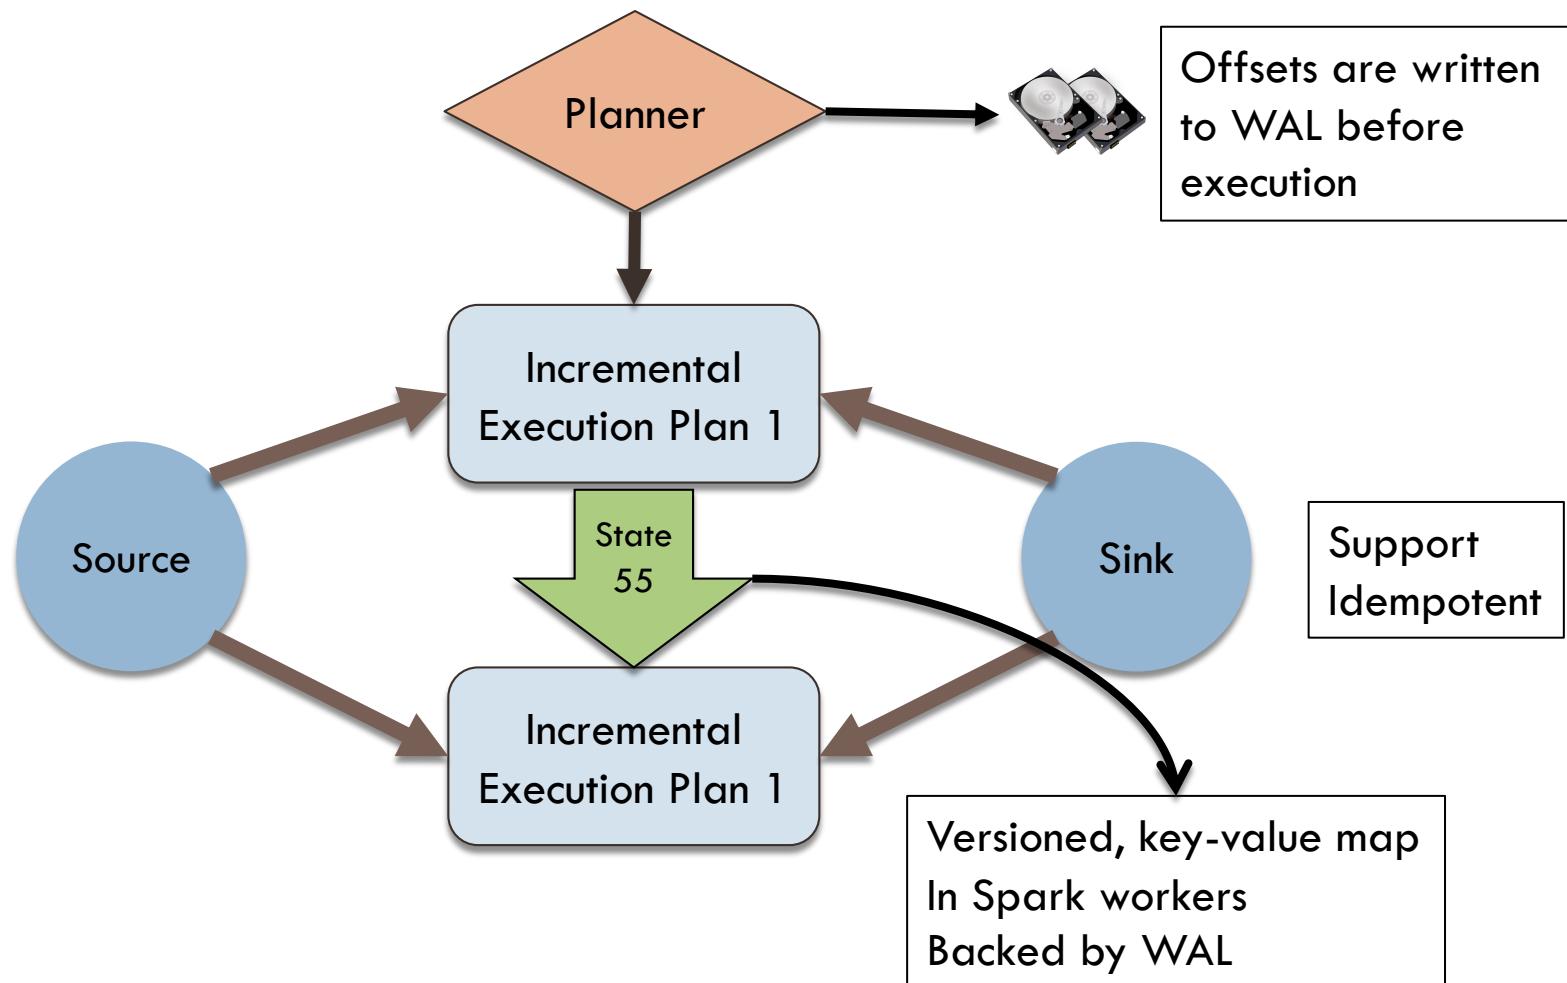


All data & metadata must be recoverable & replayable

# Fault Tolerance

46

## End-to-End Exactly Once Guarantees



# Fault Tolerance

47

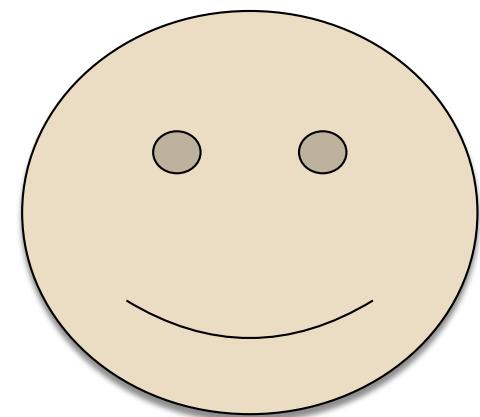
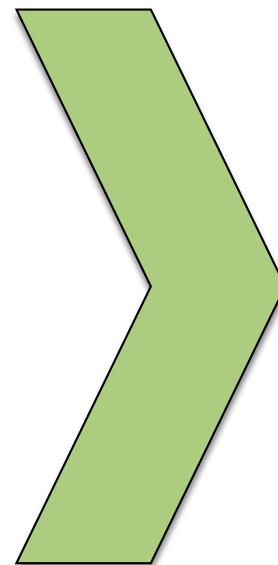
## End-to-End Exactly Once Guarantees

Offset tracking in WAL

State Management

Replayable sources

Idempotent sinks



Happy Customers

# Spark vs Others

48

## Comparing With Other Streaming Engines

Property	Structured Streaming	Spark Streaming	Apache Storm	Apache Flink	Kafka Streams	Google Dataflow
Streaming API	incrementalize batch queries	integrates with batch	separate from batch	separate from batch	separate from batch	integrates with batch
Prefix Integrity Guarantee	✓	✓	✗	✗	✗	✗
Internal Processing	exactly once	exactly once	at least once	exactly once	at least once	exactly once
Transactional Sources/Sinks	✓	some	some	some	✗	✗
Interactive Queries	✓	✓	✗	✗	✗	✗
Joins with Static Data	✓	✓	✗	✗	✗	✗

# Resources

49

- Programming guide
  - <http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>
- Blogs
  - <https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html>
  - <https://databricks.com/blog/2017/01/19/real-time-streaming-etl-structured-streaming-apache-spark-2-1.html>