



Kafka Overview

Hien Luu

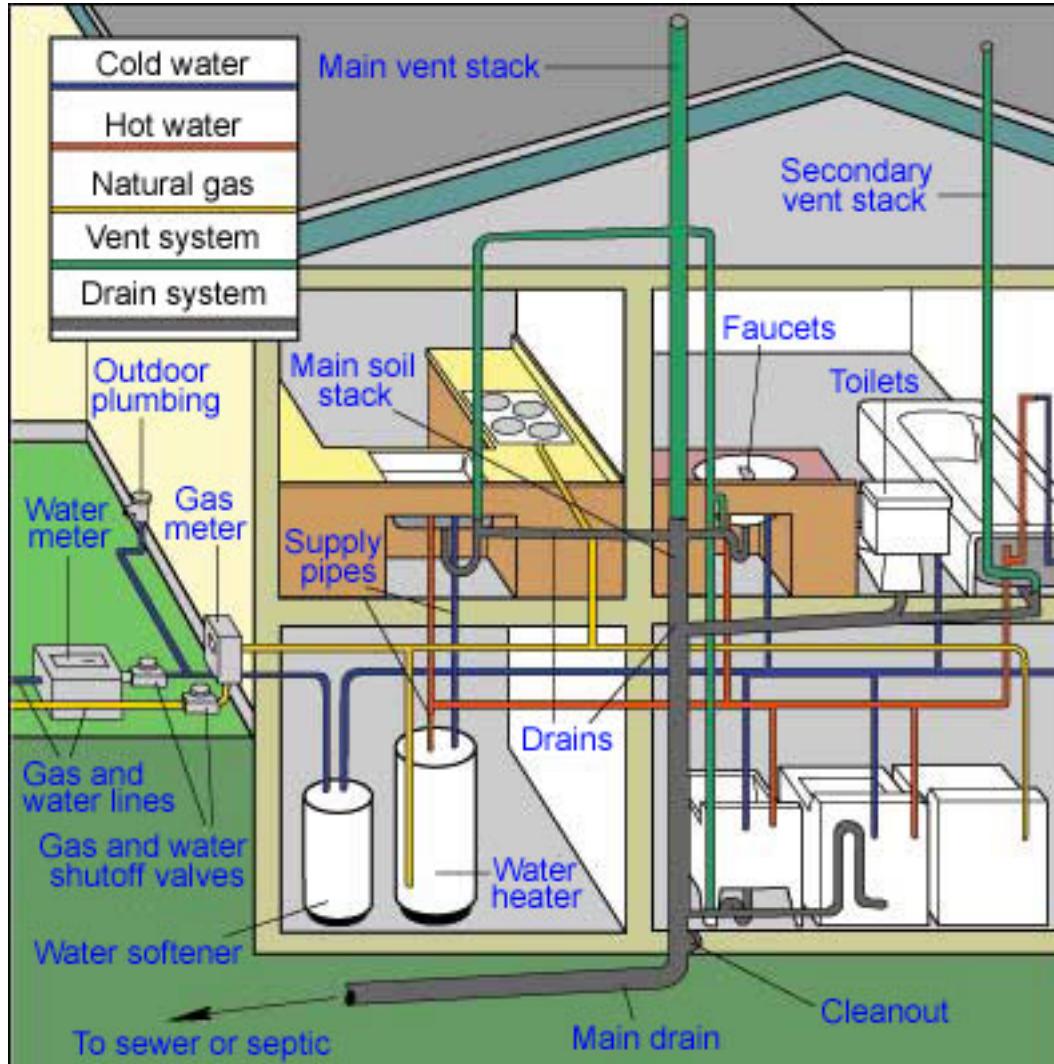
Agenda

2

- Kafka Overview
- Kafka Architecture
- Kafka APIs

Kafka Overview

3



Kafka Overview

4

A unified platform for handling all the real-time data feeds a large company might have

Kafka Overview

5

A distributed messaging system for log processing
that combines the benefits of traditional log
aggregators and messaging systems

Kafka is a distributed, partitioned, replicated commit
log service. It provides the functionality of a
messaging system, but with a unique design

Kafka Overview

6

If data is the lifeblood of high technology, Apache Kafka is the circulatory system in use at LinkedIn

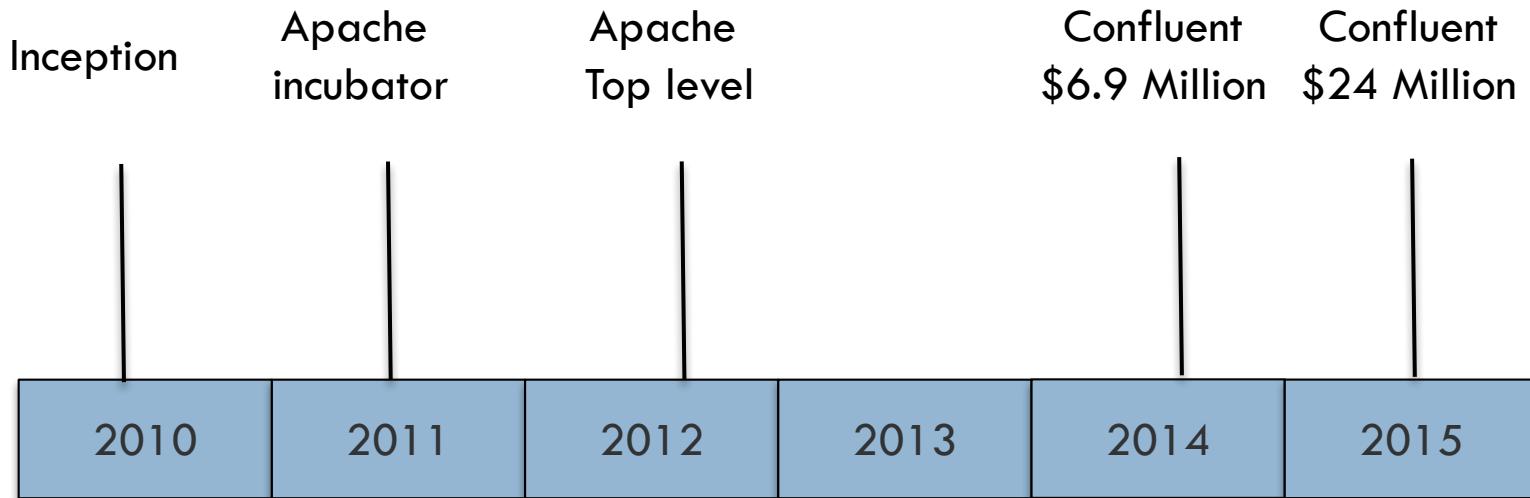
-Todd Palino

Kafka continues to be one of the key pillars in LinkedIn's data infrastructure

-Karthik Paramasivam

Kafka Overview

7



Kafka Overview

8



2,000,000,000,000 messages/day
(4 comma club)

1.5 PB/week

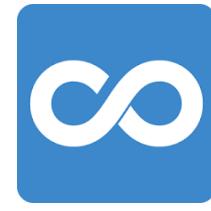
1800 servers

Kafka Overview

9



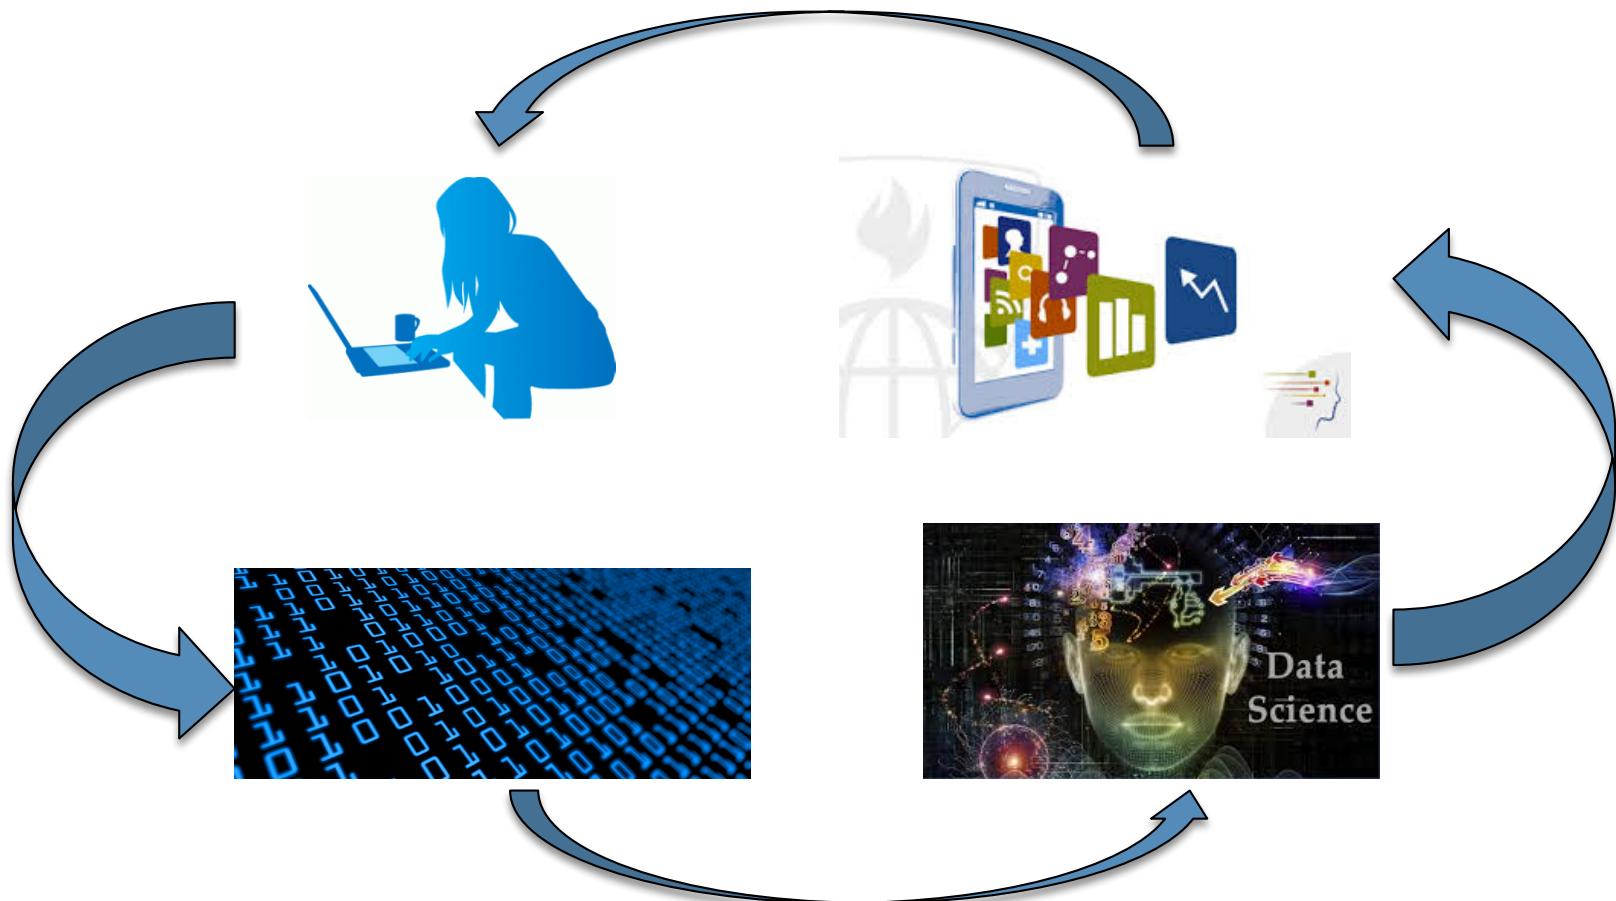
NETFLIX



Kafka Overview

10

Data Patterns in Data-Driven Companies



Kafka Overview

11

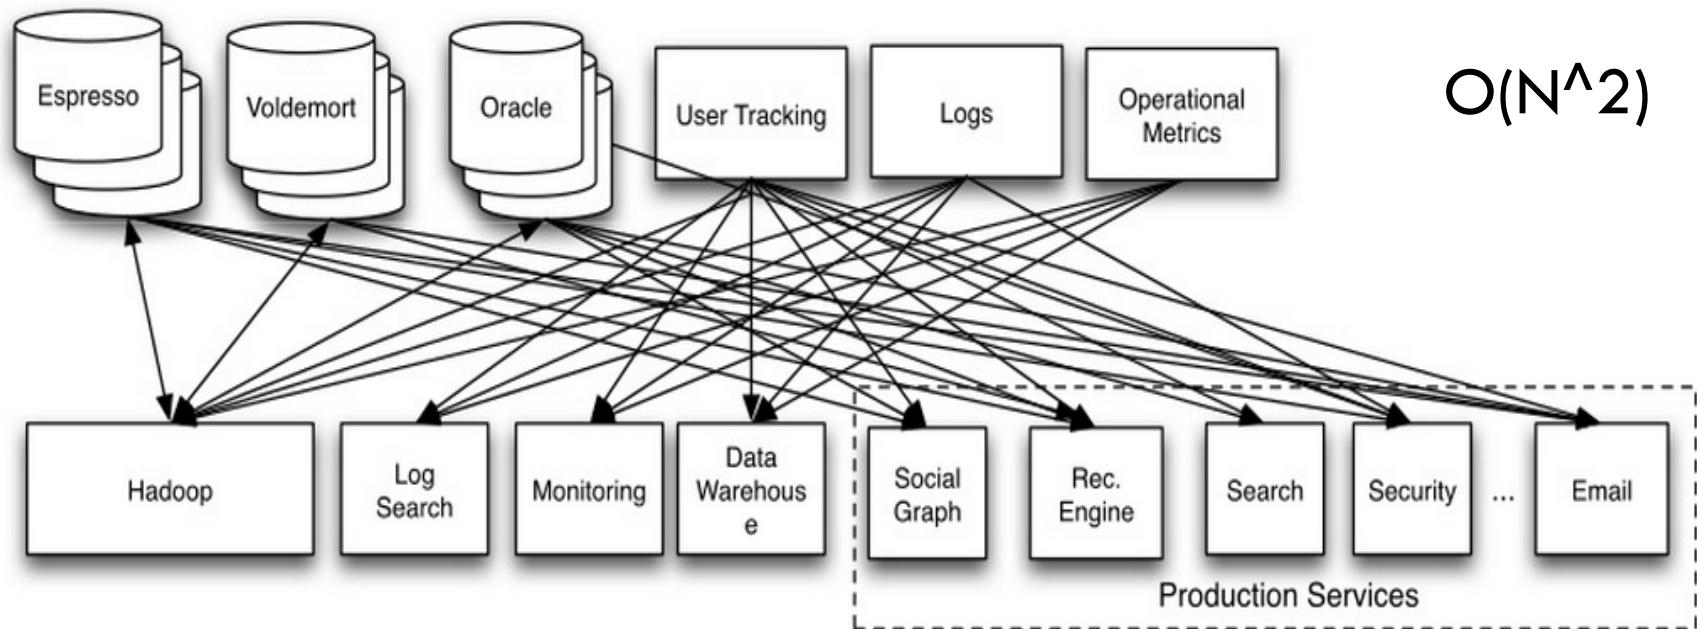
Data Sources

- Database data
 - User, products, orders
- Events
 - Clicks, impressions, pageviews, search queries
- Application metrics
 - CPU usage, request/second
- Application logs
 - Service call, errors

Kafka Overview

12

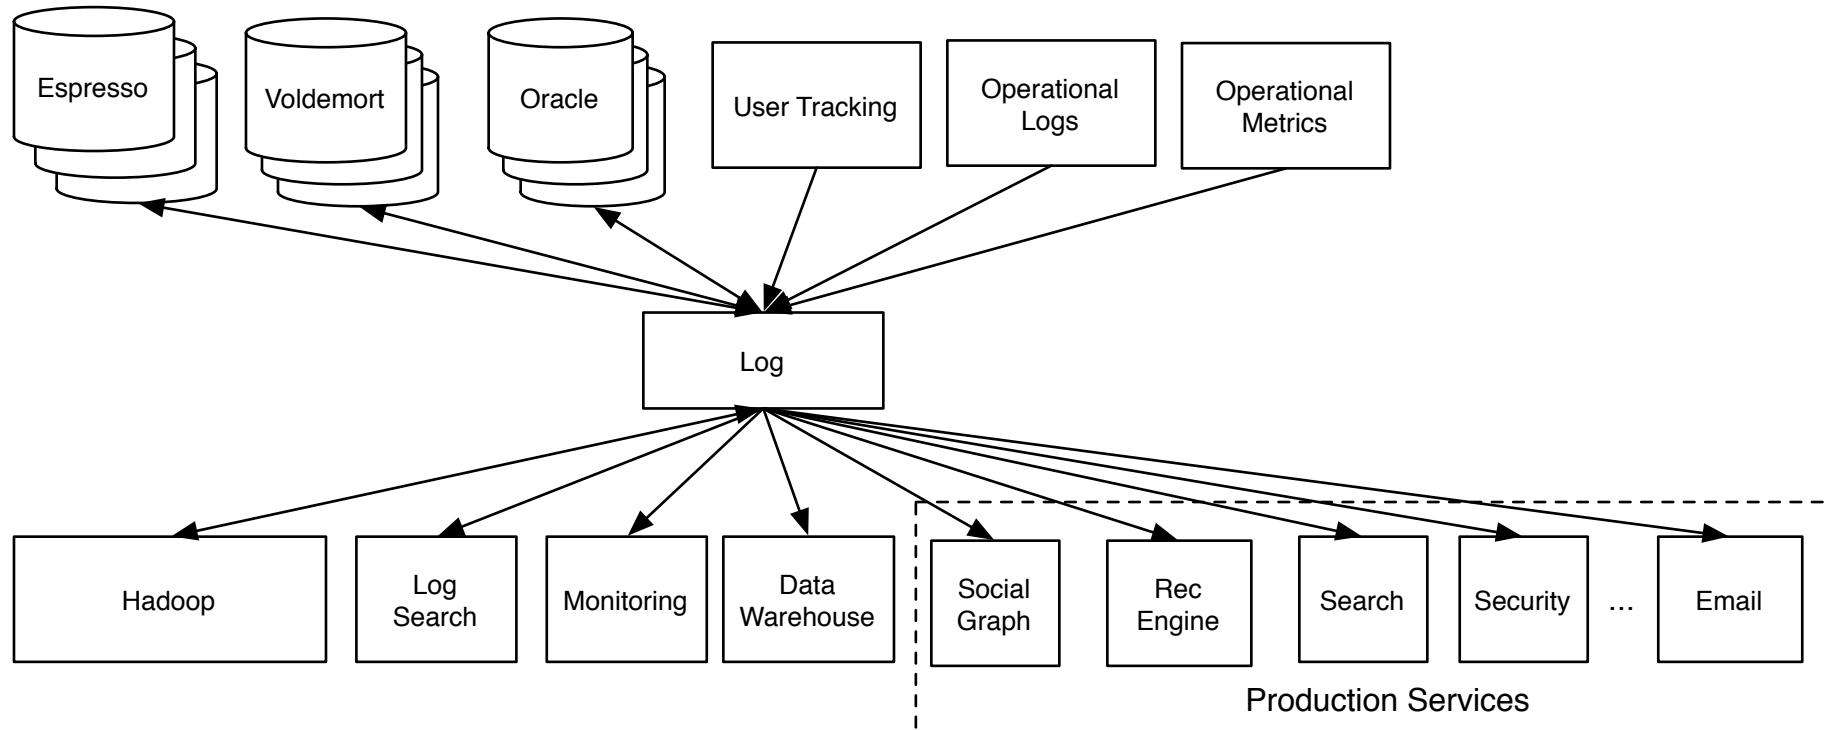
Early Days At LinkedIn



Kafka Overview

13

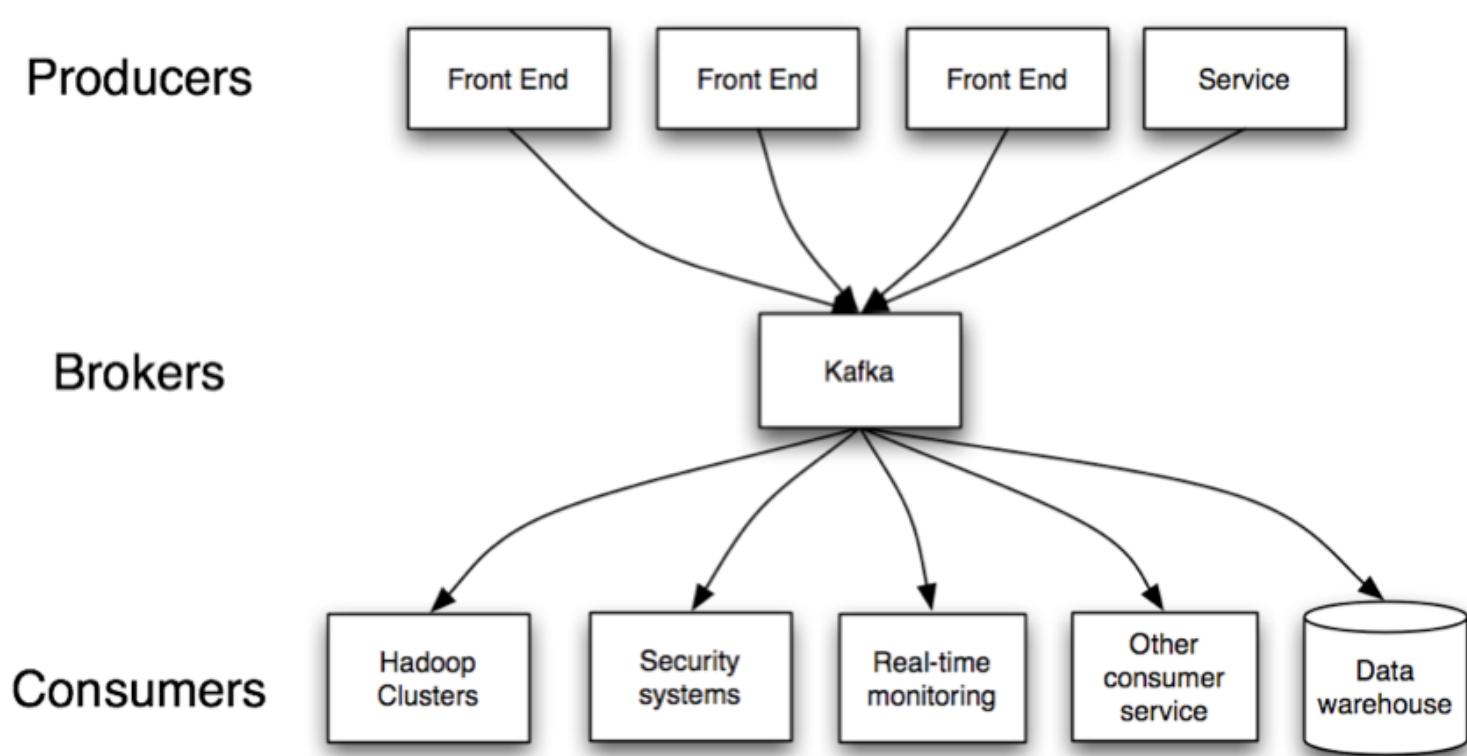
Later Days At LinkedIn



Kafka Overview

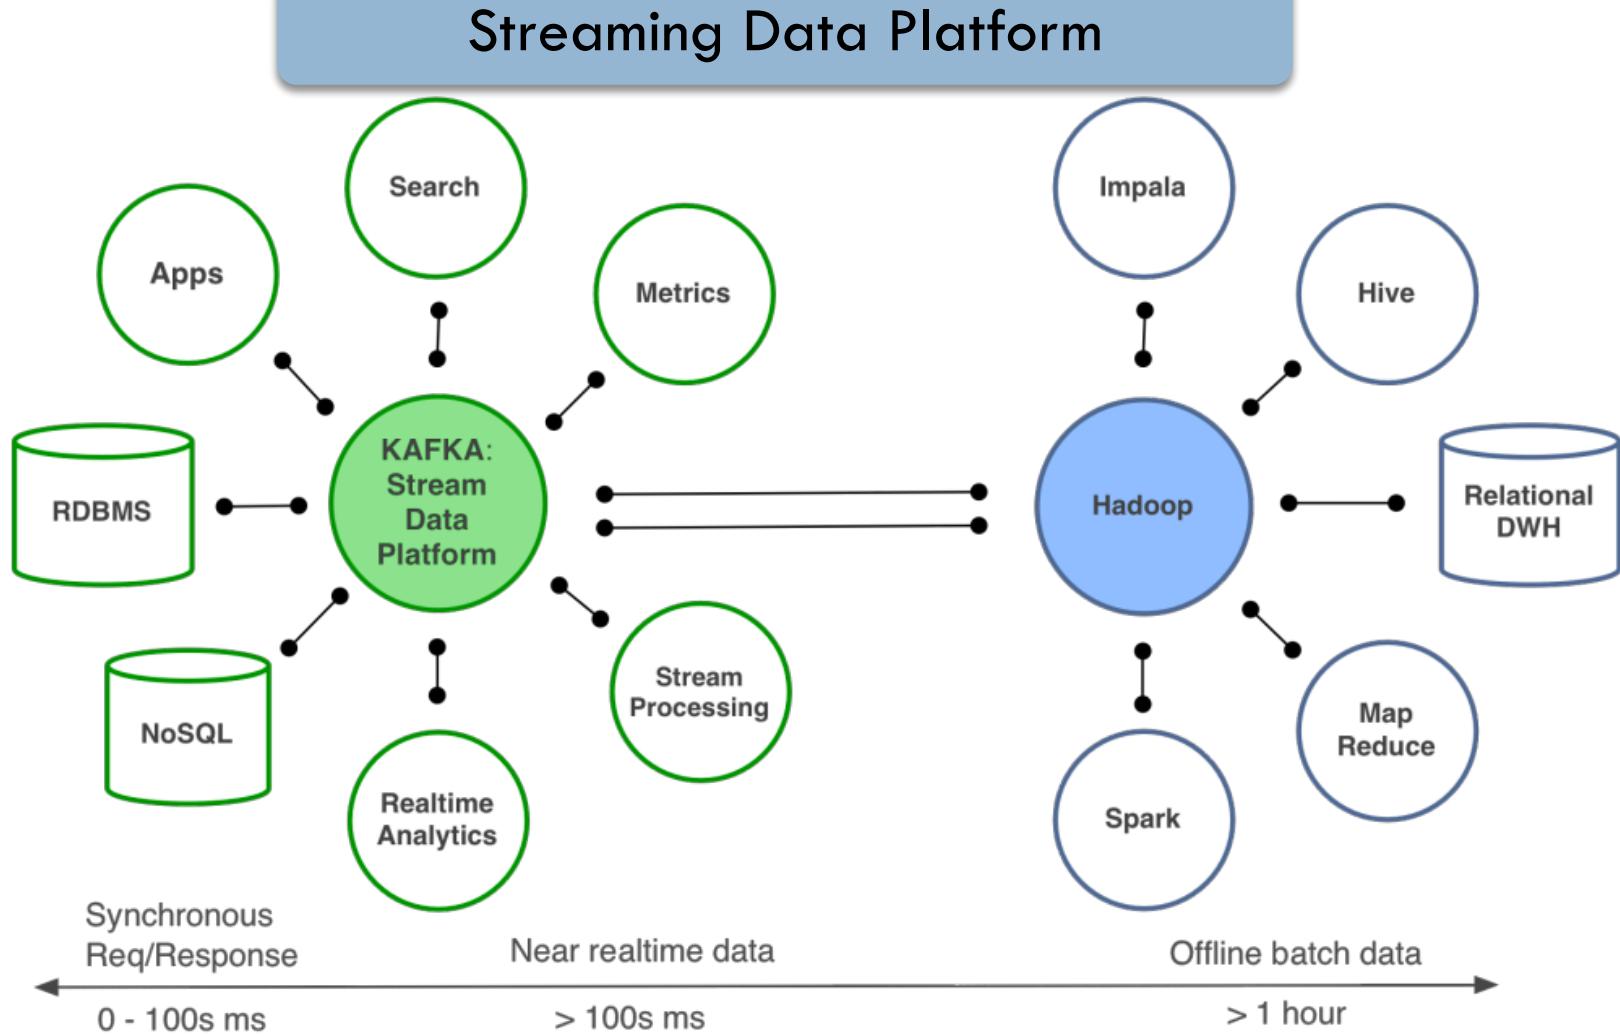
14

De-couples Data Pipelines



Kafka Overview

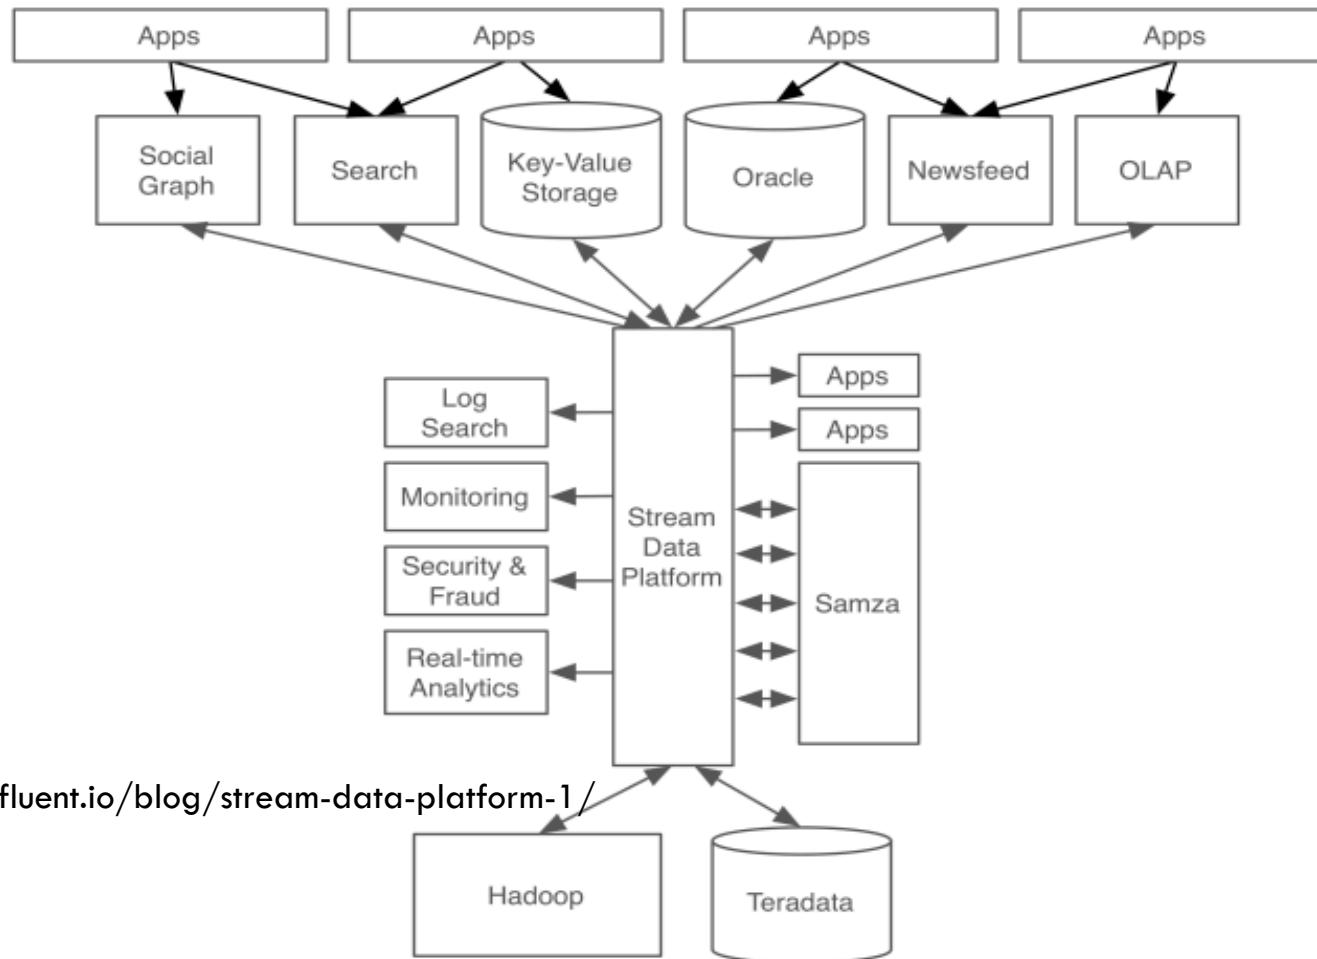
15



Kafka Overview

16

Universal Data Pipeline



Kafka Overview

17

AWS 10/2015



Amazon Kinesis

Amazon Kinesis services make it easy to work with real-time streaming data in the AWS cloud.

A fire hose of streaming data, deliver continuously into AWS

Kafka Architecture

18

Desired Properties

- Scalability
 - Partitioning
- Guarantees of a database
 - Order and persistent
- High throughput, high volume
- Low latency
- Fault tolerant

Kafka Architecture

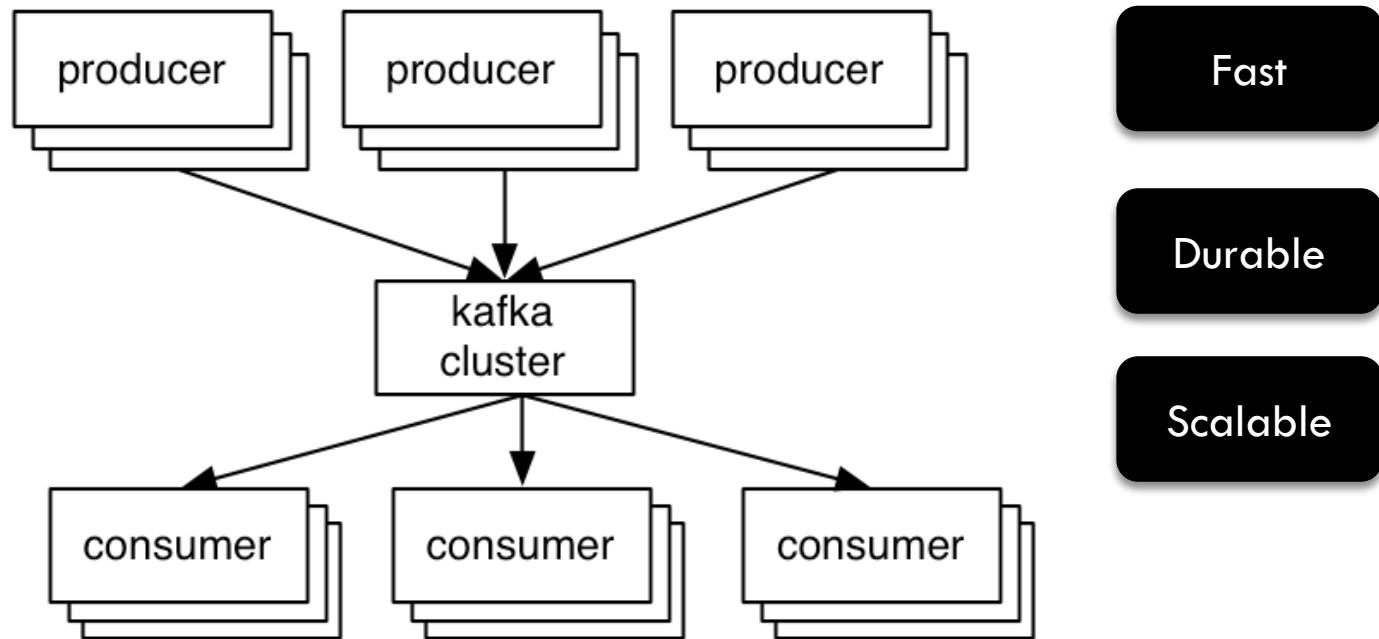
19

- Open source messaging systems
 - ActiveMQ
 - RabbitMQ
- Challenges
 - Low throughput
 - Due to strong delivery guarantees
 - No batching support
 - Weak in distributed support
 - No good persistency model
 - Near immediate message consumption model
 - Push model

Kafka Overview

20

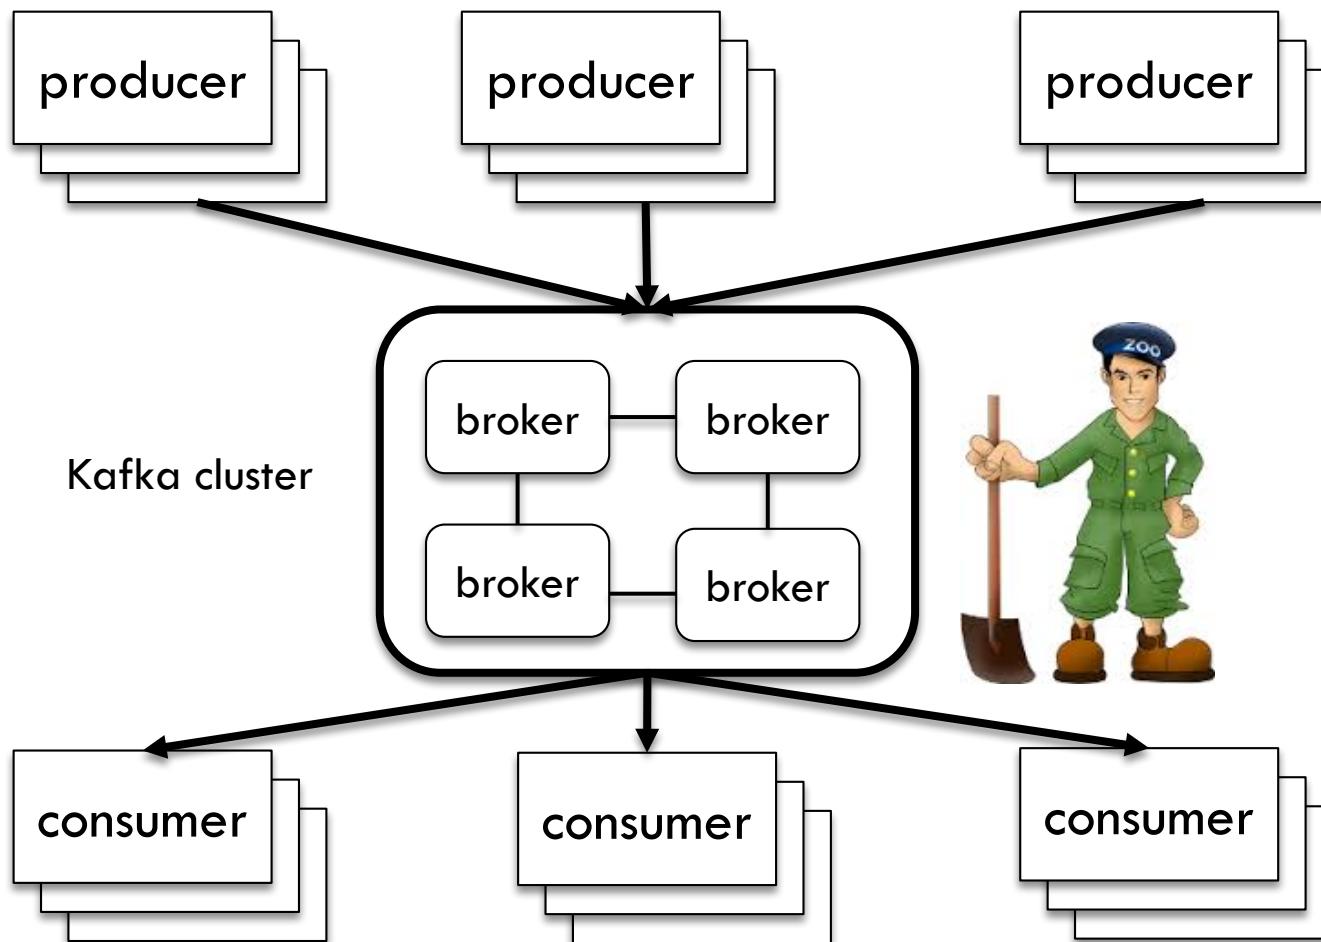
Distributed Publish-Subscribe Messaging System



Kafka Architecture

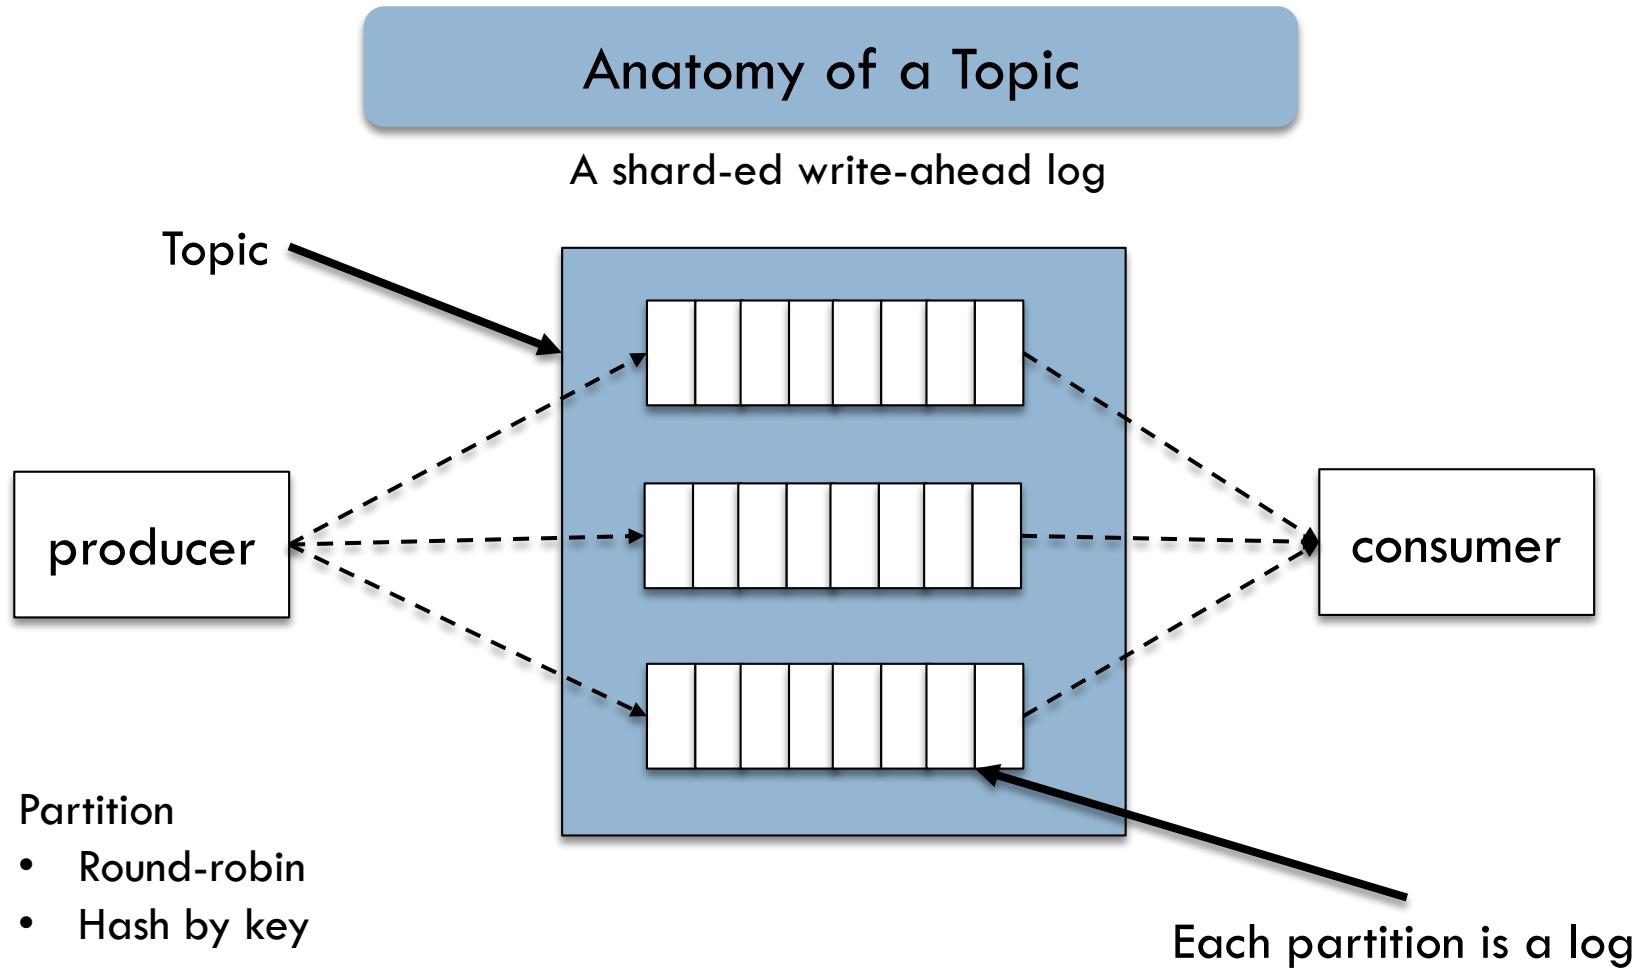
21

Built-in Cluster Management



Kafka Architecture

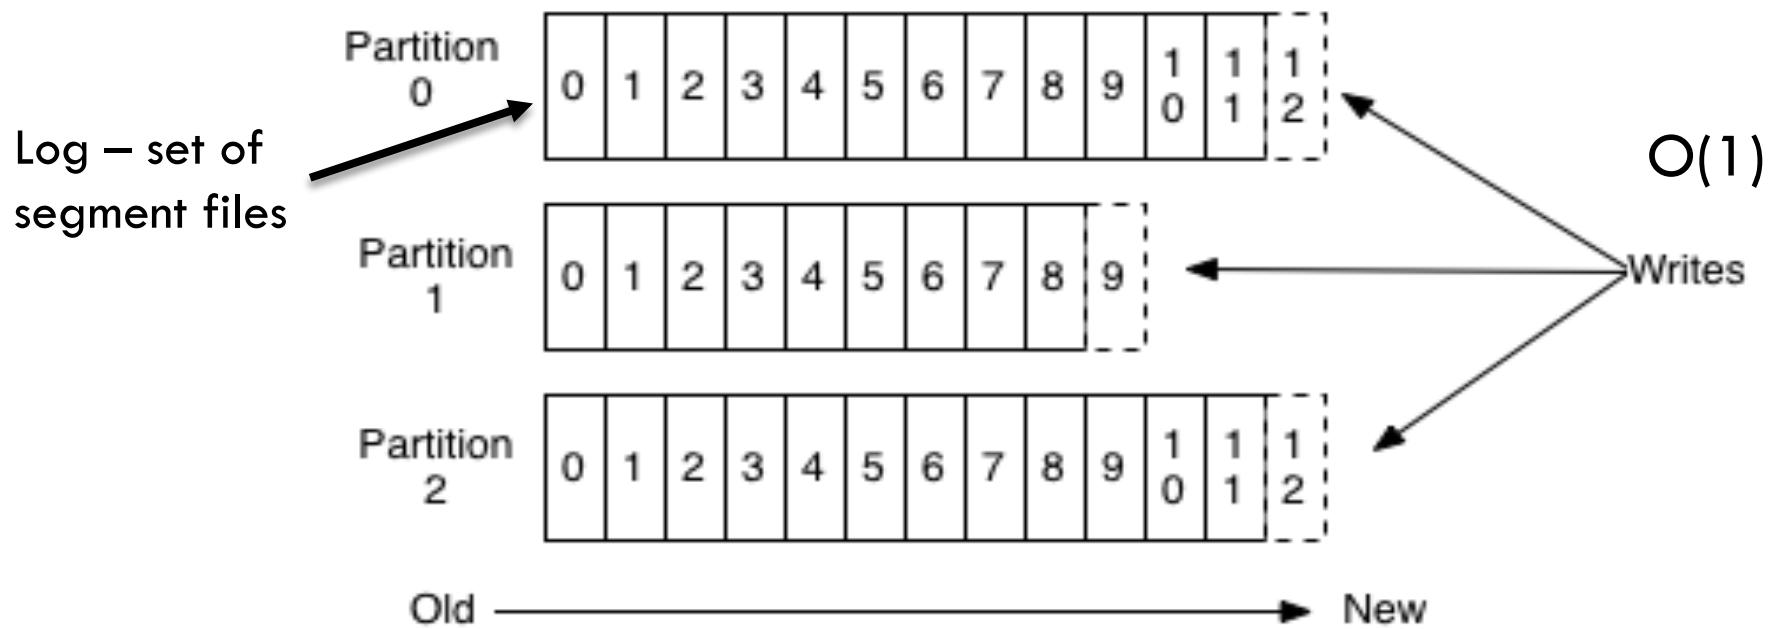
22



Kafka Architecture

23

Partition: ordered + immutable sequence of messages



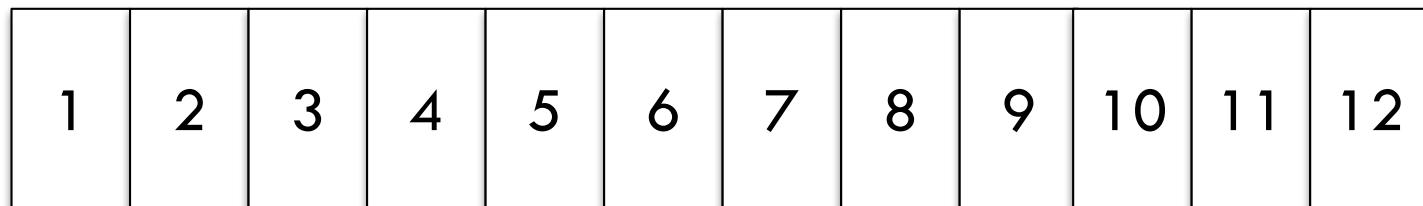
Partition is a unit of parallelism for throughput

Kafka performance is effectively constant with respect to data size.

Kafka Architecture

24

Message Format



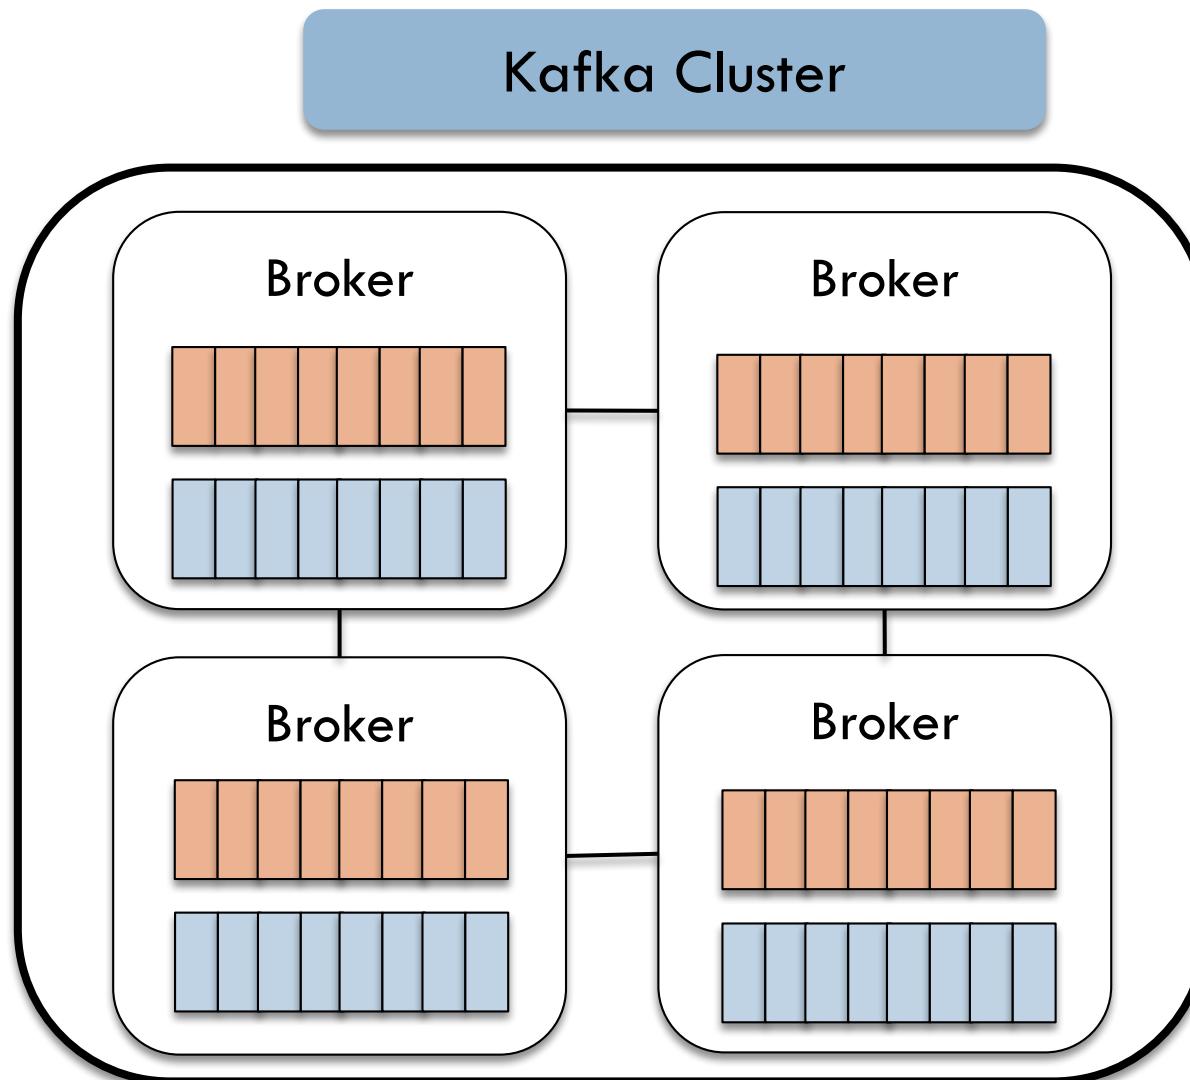
Message Format



Each message is addressed by its logical offset in the log

Kafka Architecture

25



- Two topics
- Four replicas/topic



Kafka Architecture

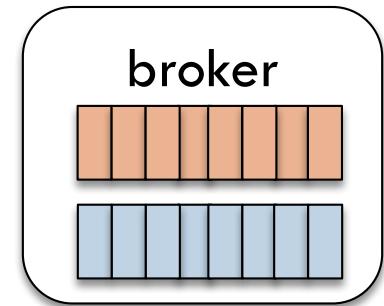
26

□ Partition

- Maintain by a single broker
- Replicate for fault-tolerance
- SLA
 - Time based
 - Size

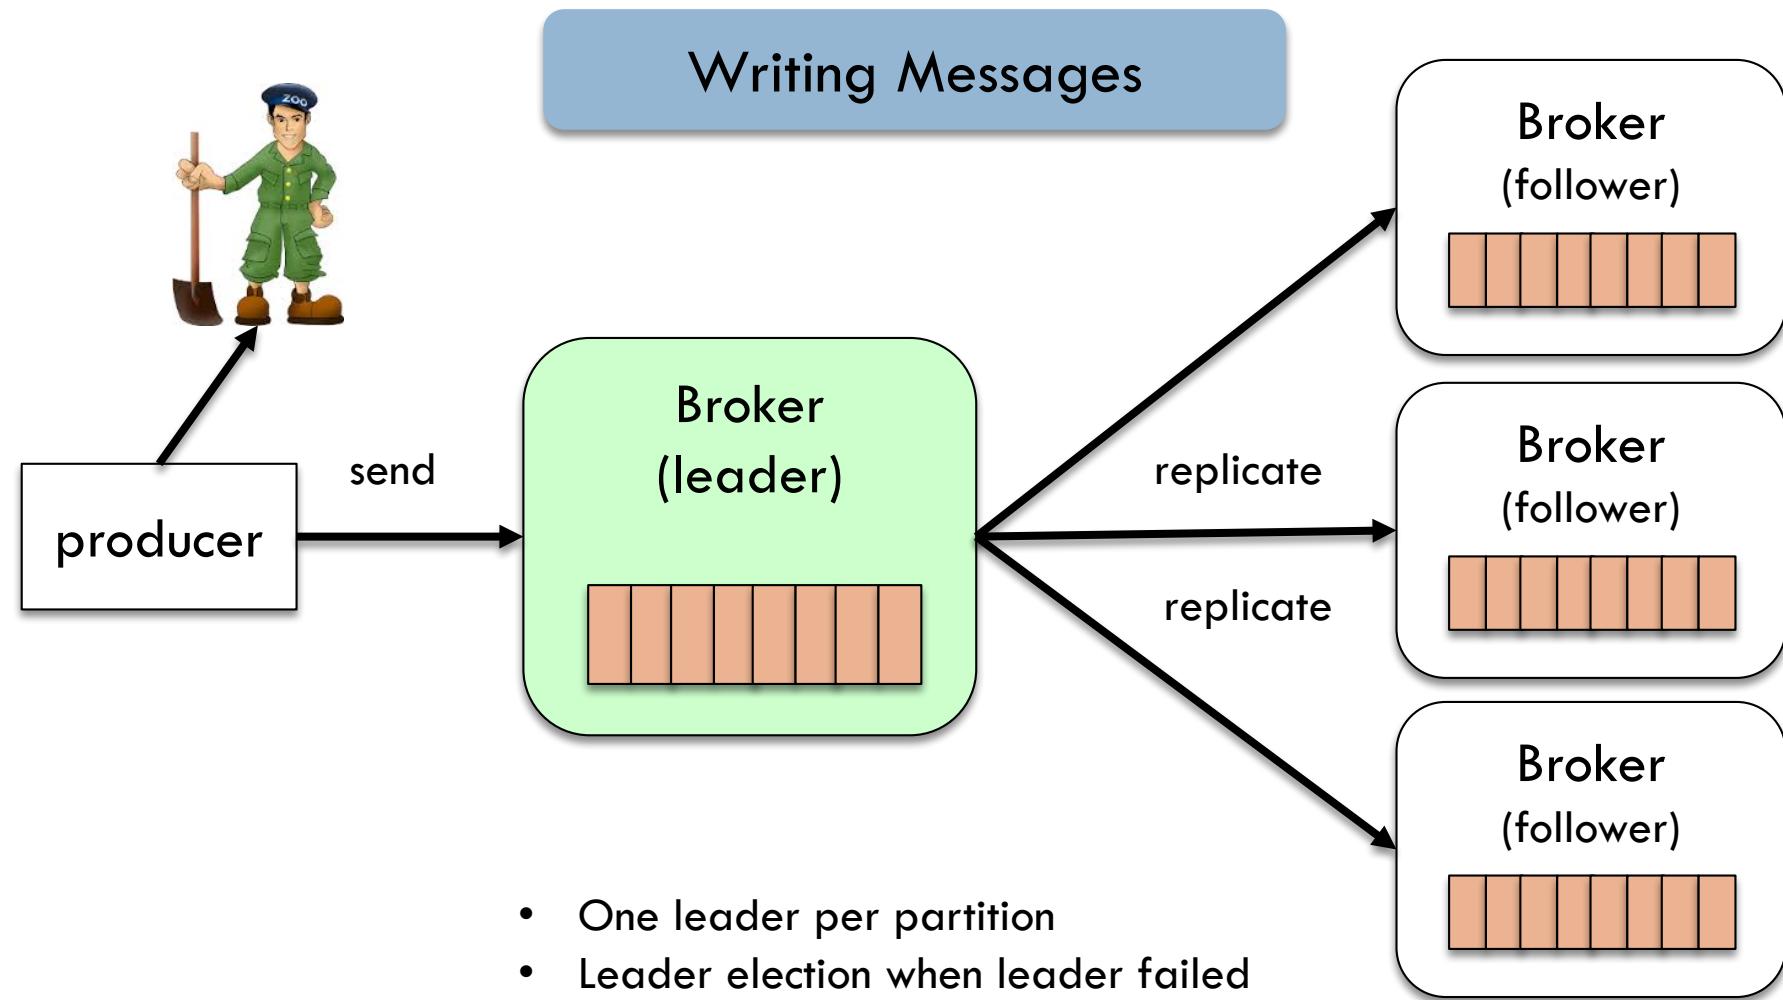
□ Broker

- Stateless
- Responsible for 1 or more partitions
- Handle replication
- Leader for some partitions
- Follower for some partitions



Kafka Architecture

27



Kafka Architecture

28

1 topic with 4 partitions with 3x replicas

Broker #1

Broker #2

Broker #3

Broker #4

topic1-part1

topic1-part1

topic1-part1

topic1-part2

topic1-part2

topic1-part3

topic1-part3

topic1-part3

topic1-part4

topic1-part4

topic1-part4

Leader

Follower



Kafka Architecture

29

- Partition leader responsibility
 - Handle all reads and writes to a partition
 - A message is committed only after successfully replicated
 - Track "in-sync" replica list
 - Detect slow replicas and remove them from list
- Partition follower responsibility
 - Pull messages from the leader
 - Apply messages to its log

Only committed messages are ever given out to the consumer

Kafka Architecture

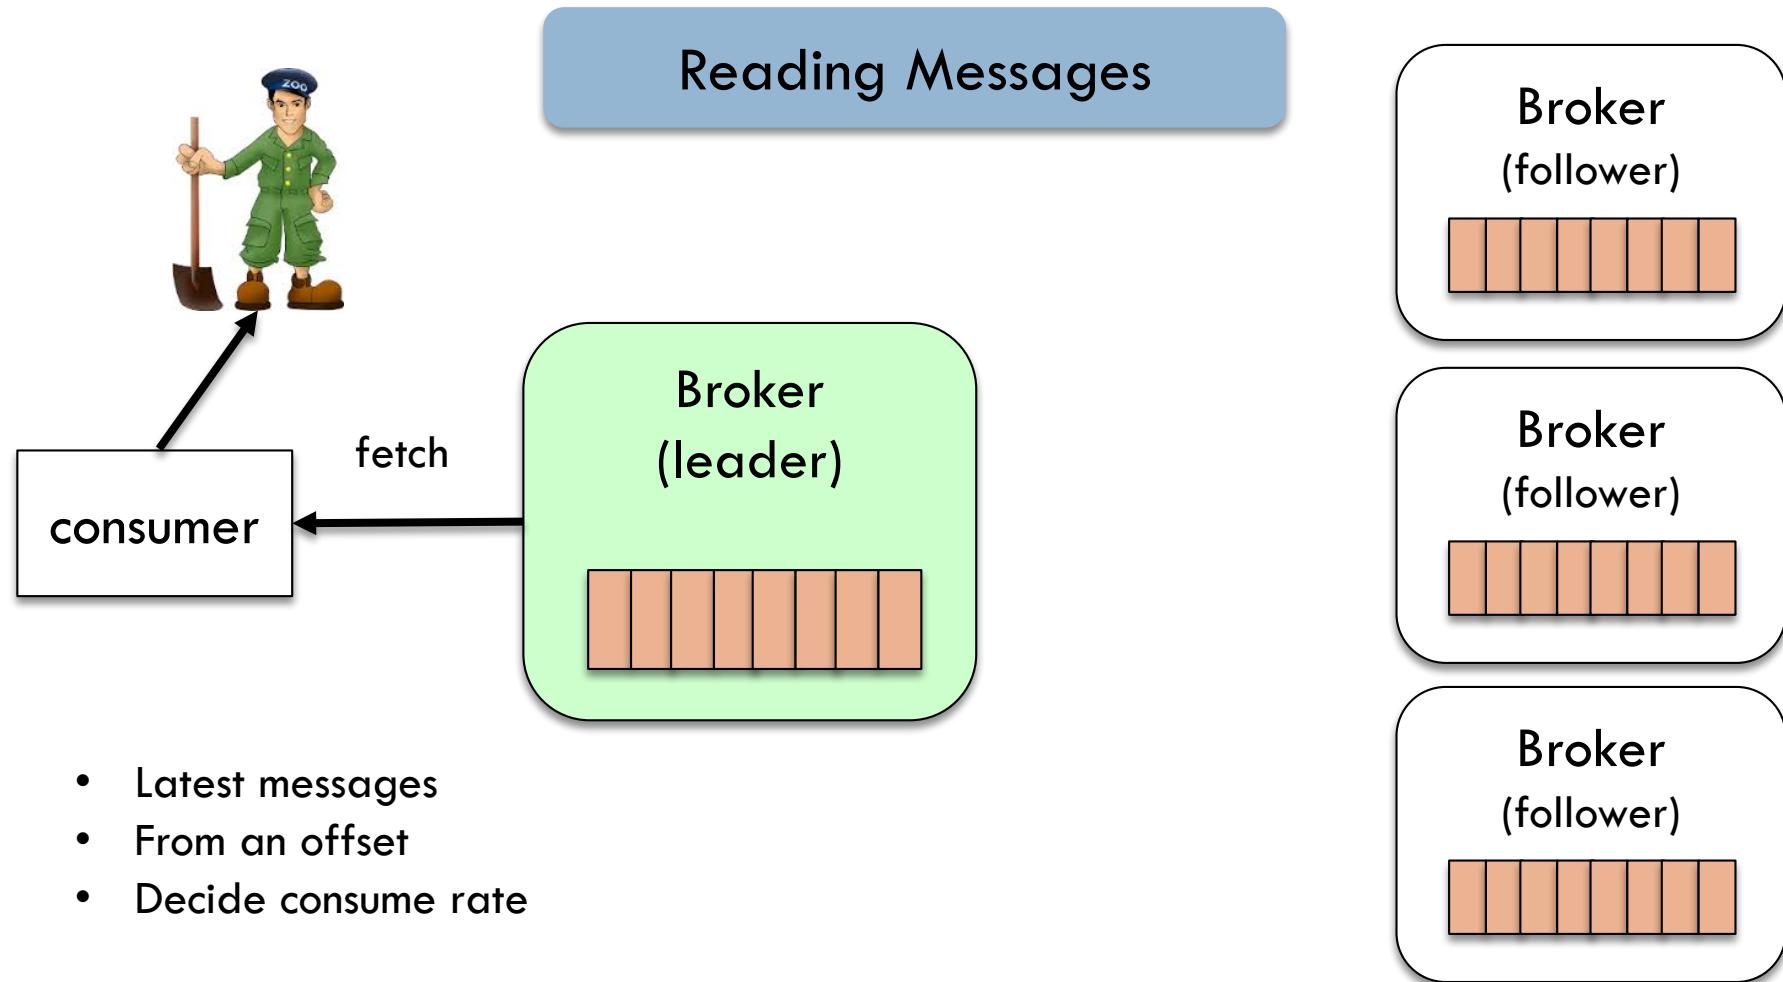
30

A paradigm shift – simplicity through delegation

*"Kafka doesn't keep track of which consumers
consuming which messages"*

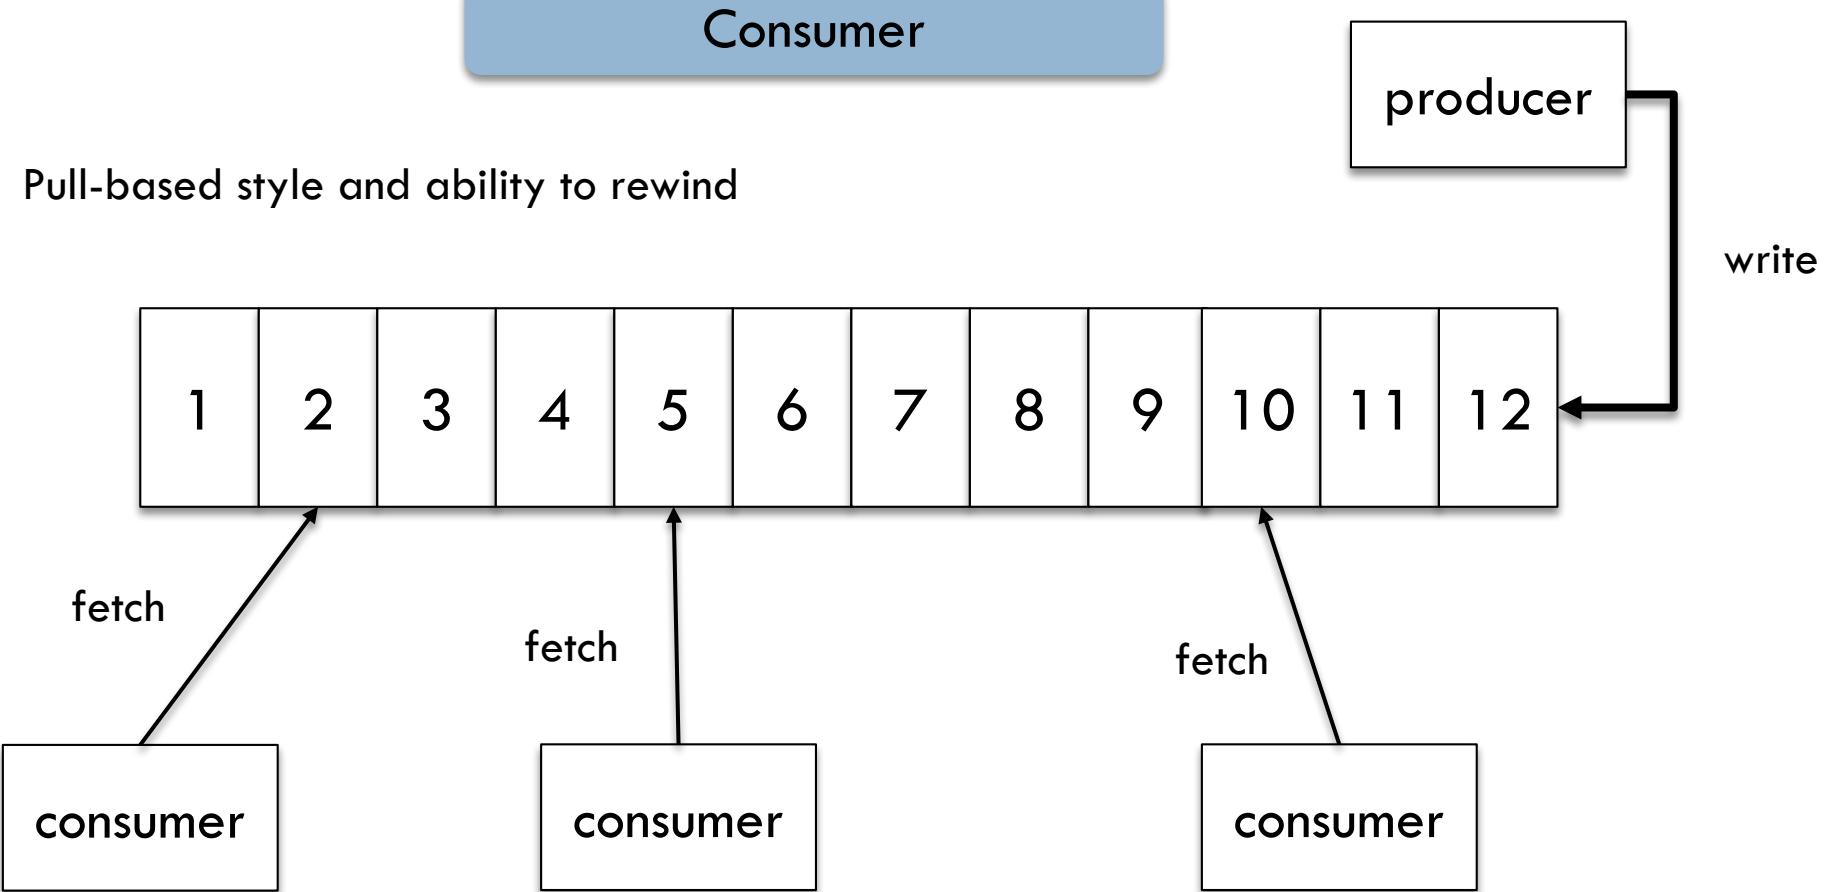
Kafka Architecture

31



Kafka Architecture

32



Leverage Unix sendfile API to deliver bytes from brokers to consumer

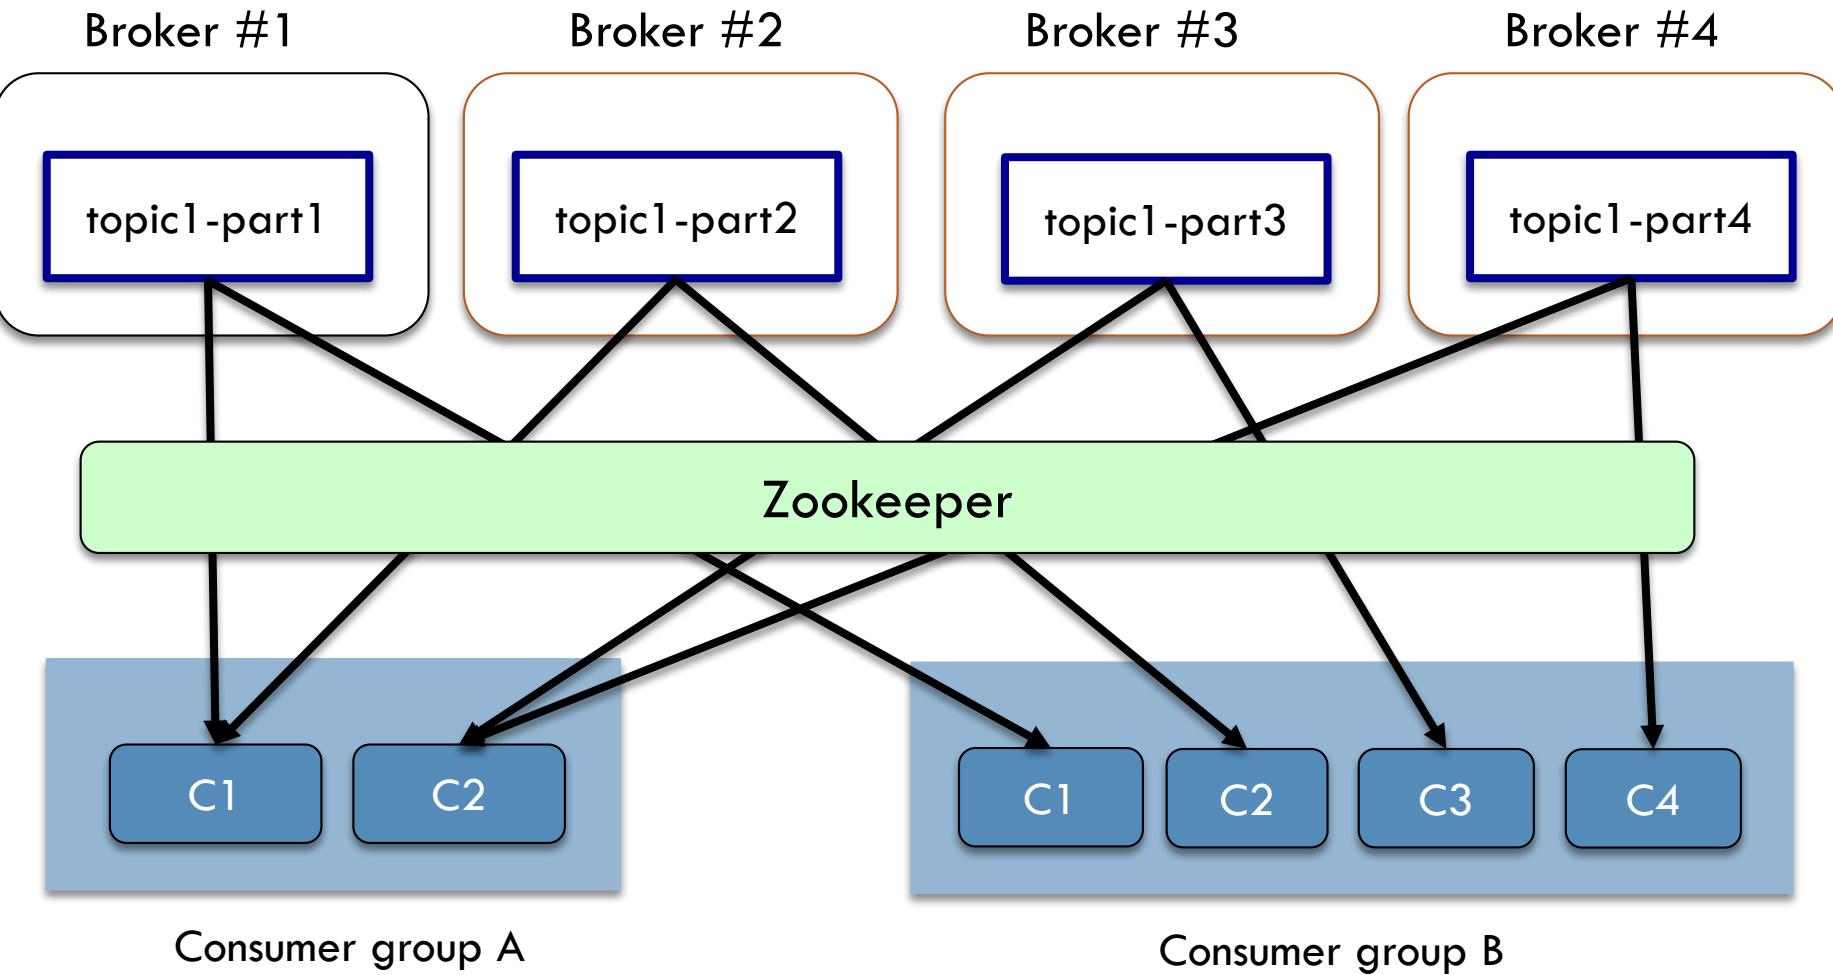
Kafka Architecture

33

- Consumer
 - Responsible for maintaining message offset
 - Combination of offset + partition + topic
 - Leverage Zookeeper for this
 - Consume messages from one partition
 - Flexible for rewinding
- Consumer group
 - A unique name/id within a Kafka cluster
 - A group of consumers consuming messages for a topic
 - Each member consumes messages from one partition

Kafka Architecture

34



Kafka Architecture

35

Broker #1

Broker #2

Broker #3

Broker #4

topic1-part1

topic1-part2

topic1-part3

topic1-part4

Zookeeper

C1

C2

C1

C2

C3

X

Consumer group A

Consumer group B

Kafka Architecture

36

- Message delivery guarantees – consumer in group
 - A message is delivered to only one consumer
 - Messages delivered in the order they were stored
 - Order is not guaranteed across partitions
- To achieve total ordering
 - Consider single partition topic
 - Handle in your own application logic

Kafka Architecture

37

Message Delivery Semantics

Guarantee	Description
At most one	Messages maybe lost, but never redelivered
At least one	Messages are never lost but may be redelivered
Exactly one	Messages are never lost and deliver only once

Producer

Durability	Latency	Description
No Ack	No delay	Some data loss
Leader Ack	1 network rountrip	A few data loss
Full Ack	2 network rountrips	No data loss

Consumer - At least one guarantee delivery

Kafka Architecture

38

□ Performance

- 2 millions writes/second – 3 producers with 3x async replication
- End-to-end latency
 - 2ms median, 14 ms (99.9th percentile)

Kafka APIs

39

Producer

```
public interface Producer<K, V> extends Closeable {  
    public Future<RecordMetadata> send(ProducerRecord<K, V> record);  
    public Future<RecordMetadata> send(ProducerRecord<K, V> record,  
                                       Callback callback);  
  
    public void flush();  
    public List<PartitionInfo> partitionsFor(String topic);  
  
    public Map<MetricName, ? extends Metric> metrics();  
  
    public void close();  
    public void close(long timeout, TimeUnit unit);  
}
```

Kafka APIs

40

Producer

```
public class SimpleProducer {  
    public static void main(String[] args) {  
        String topic = (String) args[0];  
        String count = (String) args[1];  
        String msg = "Message Publishing Time - " + runtime;  
  
        Properties props = new Properties();  
        props.put("metadata.broker.list", "<host>:<port>");  
        props.put("serializer.class", "kafka.serializer.StringEncoder");  
        props.put("request.required.acks", "1");  
        ProducerConfig config = new ProducerConfig(props);  
  
        Producer<String, String> producer = new Producer<String, String>(config);  
        KeyedMessage<String, String> data =  
            new KeyedMessage<String, String>(topic, msg);  
  
        producer.send(data);  
    }  
}
```

Kafka APIs

41

Consumer

```
public interface Consumer<K, V> extends Closeable {  
    public void subscribe(Collection<String> topics);  
    public ConsumerRecords<K, V> poll(long timeout);  
  
    public void close();  
    public void close(long timeout, TimeUnit unit);  
}
```

Kafka APIs

42

Consumer

```
public class SimpleConsumer {  
    public static void main(String[] args) {  
        String topic = (String) args[0];  
  
        Properties props = new Properties();  
        props.put("zookeeper.connect", "<host>:<port>");  
        props.put("group.id", "mygroup");  
        ConsumerConfig config = new ConsumerConfig(props);  
  
        KafkaConsumer consumer = new KafkaConsumer(config);  
  
        consumer.subscribe(Collections.singletonList(topic));  
    }  
}
```

Kafka APIs

43

Consumer

```
// continue from previous slide
try {
    while (true) {
        ConsumerRecords<String, String> records = consumer.poll(100);
        for (ConsumerRecord<String, String> record : records) {
            String debugMsg = String.format("topic: %s, partition: %s,
                offset: %d, message: %s",
                record.topic(), record.partition(), record.offset(),
                record.value());
            System.out.println("message: " + debugMsg);
        }
    }
} finally {
    consumer.close();
}
```

Kafka Summary

44

□ Key takeaways

- Distributed, partitioned and replicated
- Simple storage
- Message offset
- Efficient data transfer
- Stateless broker
- Pull-based consumption model
- Consumer group
- Guarantee at-least-once delivery

Resources

45

- <http://kafka.apache.org/documentation.htm>
- <http://www.confluent.io/blog>
- <https://engineering.linkedin.com/kafka/benchmarking-apache-kafka-2-million-writes-second-three-cheap-machines>
- <https://engineering.linkedin.com/kafka/running-kafka-scale>
- https://engineering.linkedin.com/apache-kafka/how-we_re-improving-and-advancing-kafka-linkedin